



A massive Optimus Prime robot stands in a city street amidst fire and smoke. He is positioned in the center, facing forward. His body is primarily blue and red, with intricate mechanical details. The background shows a city skyline with several buildings engulfed in flames and smoke. In the sky, there are other robotic figures and flying vehicles.

Attention!
The Transformers are coming!


Tim Niklas Witte



Computer Vision



Deep Reinforcement Learning



Attention!
The Transformers are coming!

Tim Niklas Witte



Natural Language Processing



GANs

Image sources:

https://machinelearningmastery.com/wp-content/uploads/2021/08/attention_research_1-727x1024.png

https://images.ctfassets.net/3viuren4us1n/1Ghw96A2tcYRfRezOwtmjx/e646778f3f53e50ea3e857e9cdb23120/Computer_vision.jpg

<https://image.stern.de/7931130/t/NJ/v2/w1440/r1/-/roboter.jpg>

https://www.studienscout-nl.de/fileadmin/_processed/_csm_Sprache-und-Sprachkurse_08f53ae680.jpg

https://studip.uni-osnabrueck.de/sendfile.php?type=0&file_id=8a554d27d543b1ce06eb1a68a34e22ae&file_name=vqgan_4_i350.png

(call dates: 20.07.22)

Structure

1. Motivation
2. Intuitive understanding of self-attention
3. Intuitive understanding of transformer architecture
4. Natural Language Processing
5. Computer Vision
6. GANs
7. Deep Reinforcement Learning
8. Modifications of attention: Synthesizer, Linformer, Reformer
9. Conclusions

Discussion: Candy-wise attention



Image sources:
https://www.flaticon.com/free-icon/candy-jar_1075135?related_id=1075135&origin=tag
<https://emojipedia.org/de/toss-face/march-2022/gesicht-mit-umarmenden-h%C3%A4nden/>
(call dates: 20.07.22)

1. Motivation




Figure 1: Performance (BLEU score) differences between RNN-based models for neural machine translation [1].

1. Motivation

Bahdanau
attention

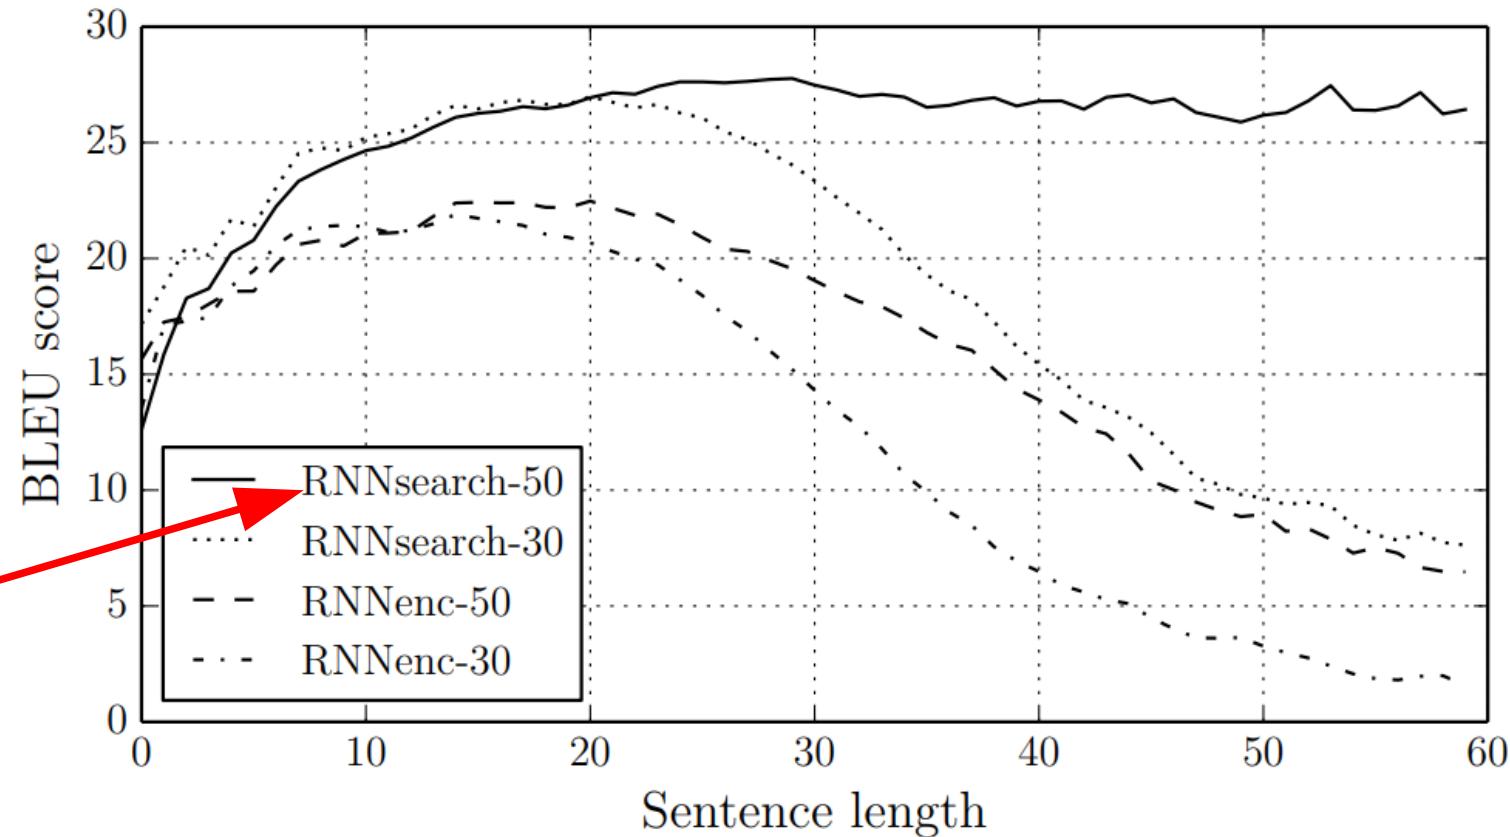


Figure 1: Performance (BLEU score) differences between RNN-based models for neural machine translation [1].



2. Intuitive understanding of self-attention



Image sources:
<https://www.campussafetymagazine.com/wp-content/uploads/2019/03/Bull-horn-announcement-PA-iStock-1000x500.jpg>
https://upload.wikimedia.org/wikipedia/commons/thumb/7/7b/Attention_Sign.svg/2302px-Attention_Sign.svg.png
https://upload.wikimedia.org/wikipedia/commons/thumb/9/9a/Simple_Attention.svg/1024px-Simple_Attention.svg.png
(call dates: 20.07.22)

2. Intuitive understanding of self-attention

The server ...

2. Intuitive understanding of self-attention



Figure 2: A Waiter.



Figure 3: Server in a data center.

2. Intuitive understanding of self-attention



Figure 2: A Waiter.



Figure 3: Server in a data center.

2. Intuitive understanding of self-attention



Figure 2: A Waiter.

The server is crashed.




Figure 3: Server in a data center.

Image sources:

<https://content.gastivo.de/wp-content/uploads/2019/09/Der-Kellner-Knigge-Perfektes-Servieren-1024x576.jpg>

<https://i0.wp.com/blog.fyipe.com/wp-content/uploads/2020/01/servers-e1605085513709.jpeg>

(call dates: 20.07.22)

2. Intuitive understanding of self-attention



Figure 2: A Waiter.

The server is crashed.

A green rectangular box with a black arrow points downwards to the text "The server is crashed.". A larger black arrow points from this text towards Figure 3, which shows a server room.



Figure 3: Server in a data center.

Image sources:

<https://content.gastivo.de/wp-content/uploads/2019/09/Der-Kellner-Knigge-Perfektes-Servieren-1024x576.jpg>

<https://i0.wp.com/blog.fyipe.com/wp-content/uploads/2020/01/servers-e1605085513709.jpeg>

(call dates: 20.07.22)

2. Intuitive understanding of self-attention

English

This is the first book I did.

Portuguese

Este e o primerio livro que eu fiz.

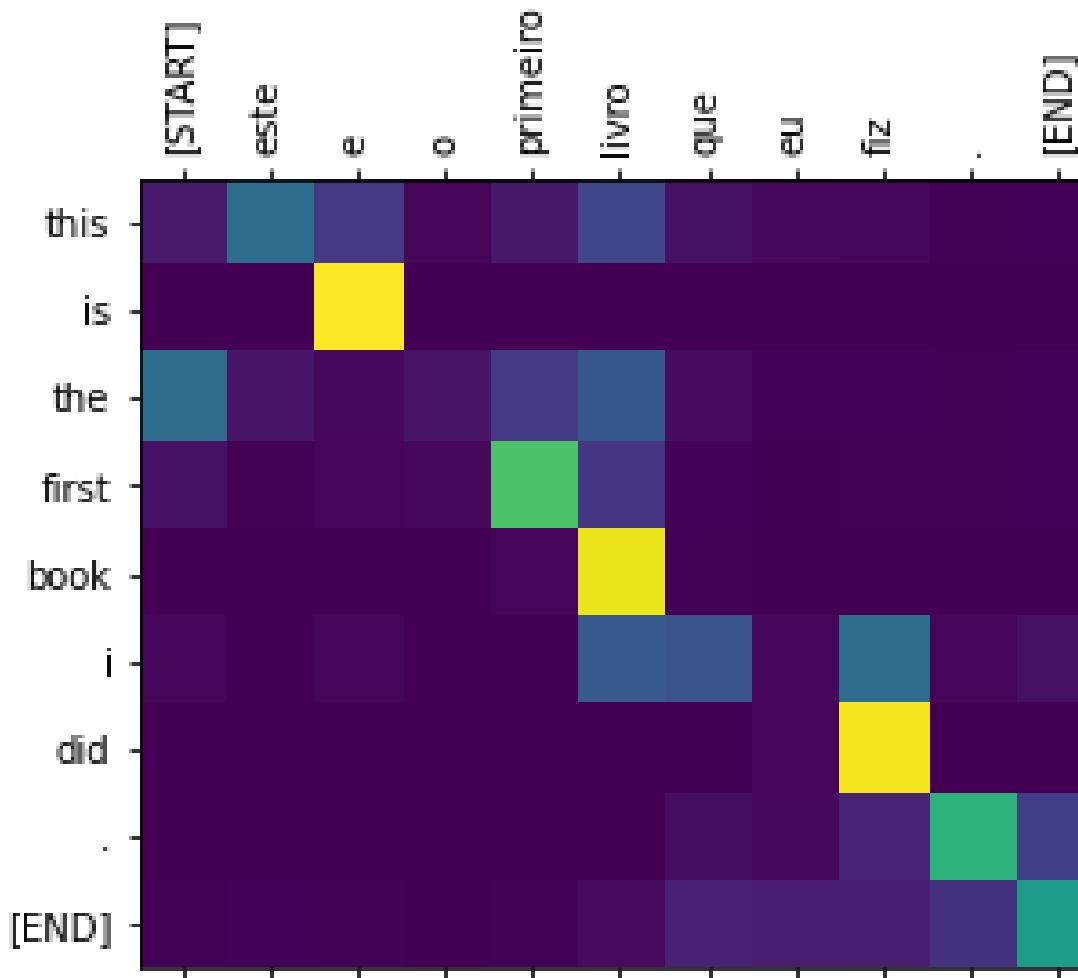


Figure 4: Self-attention heatmap.

2. Intuitive understanding of self-attention

English

This is the first book I did.

Portuguese

Este **e** o **primerio** **livro** que **eu** **fiz**.

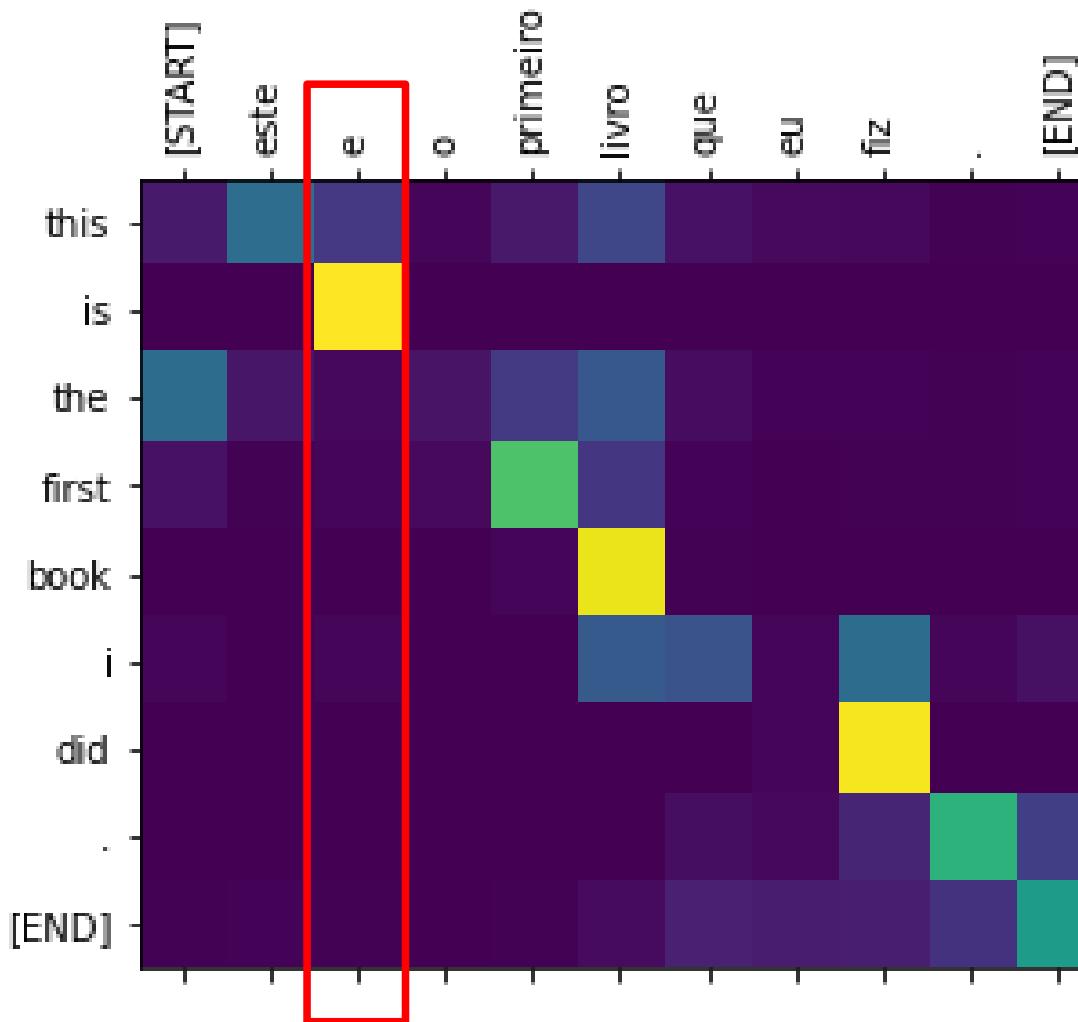


Figure 4: Self-attention heatmap.

2. Intuitive understanding of self-attention

English

This is the first book I did.

Portuguese

Este **e** o **primerio** **livro** que eu **fiz**.


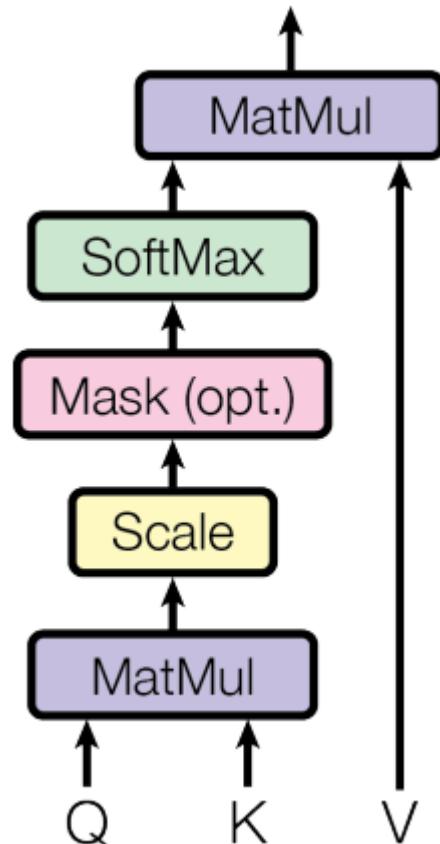


Figure 4: Self-attention heatmap.

2. Intuitive understanding of self-attention



$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Q: Query matrix shape: $d_k \times D_H$
K: Key matrix shape: $d_k \times D_H$ Hidden size
V: Value matrix shape: $d_v \times D_H$

Figure 5: Self-attention [2].

2. Intuitive understanding of self-attention

Database		
Video title	Description	Video

Figure 6: Self-attention example based on a YouTube search query.

2. Intuitive understanding of self-attention

Database		
Video title	Description	Video
Key		Value

Figure 6: Self-attention example based on a YouTube search query.

2. Intuitive understanding of self-attention

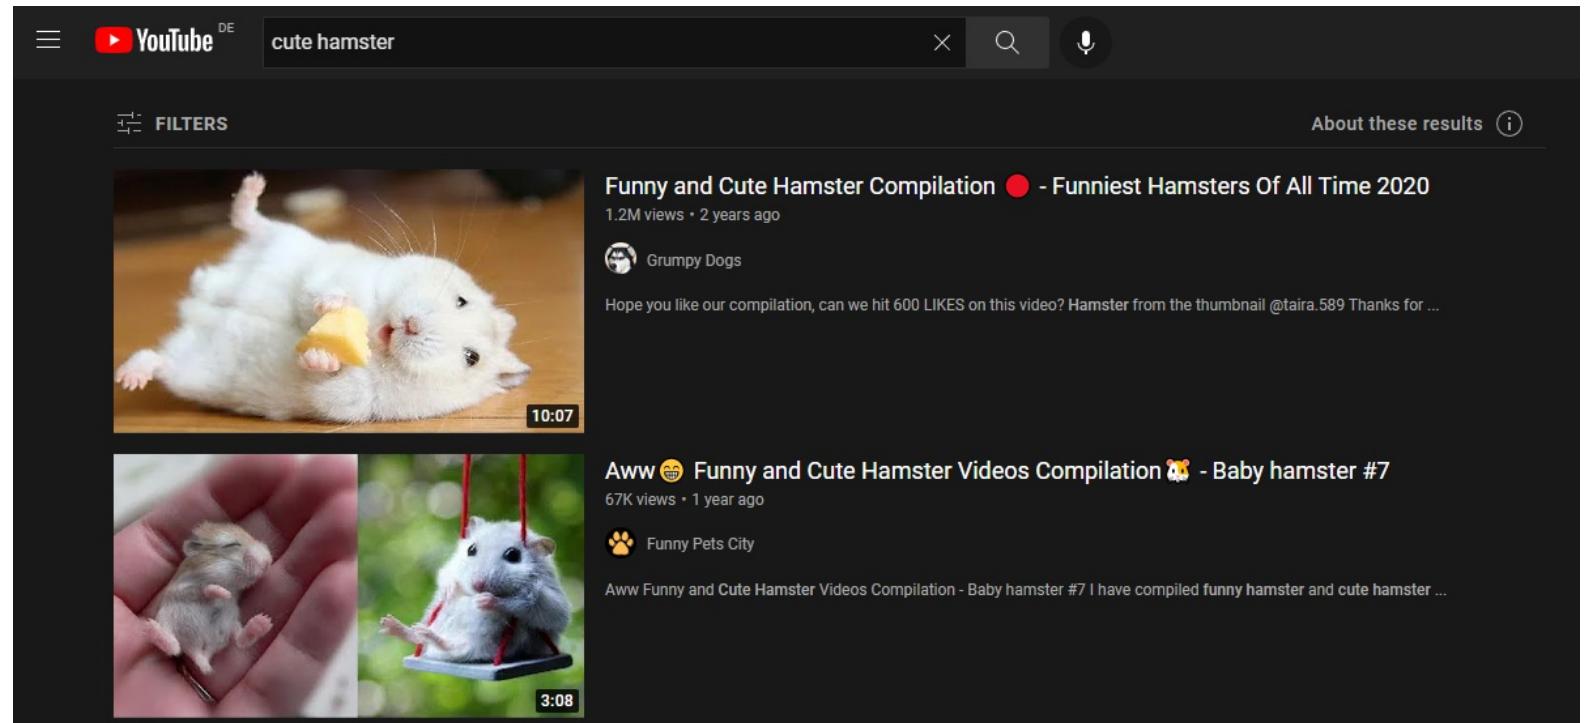
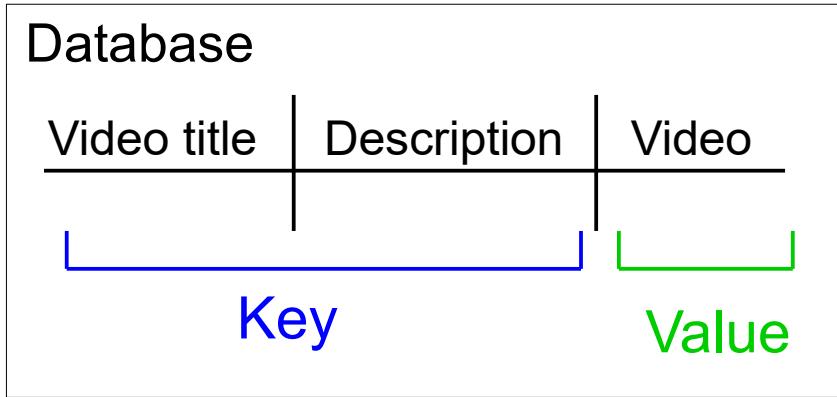


Figure 6: Self-attention example based on a YouTube search query.

2. Intuitive understanding of self-attention

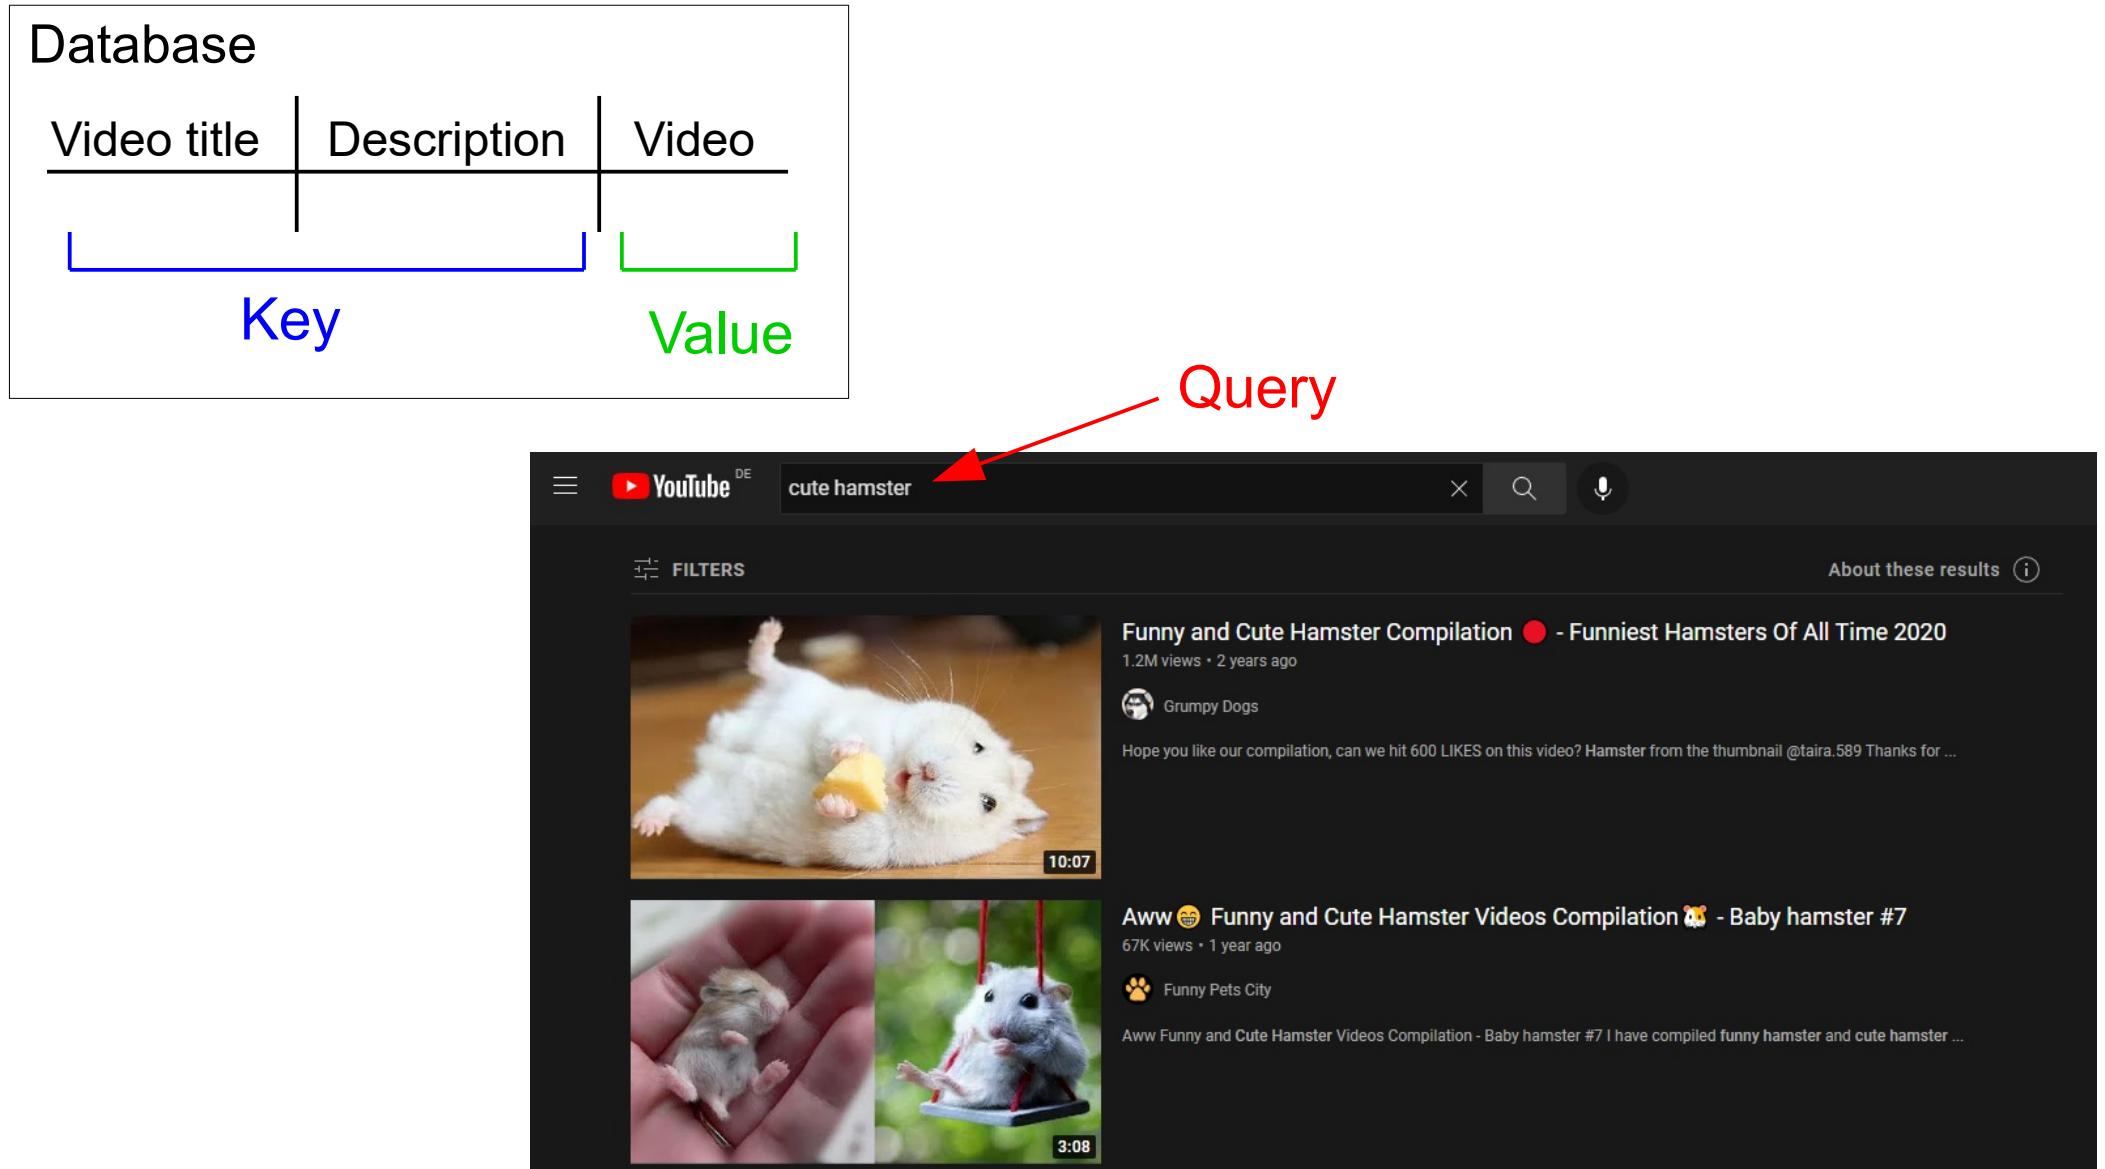


Figure 6: Self-attention example based on a YouTube search query.

2. Intuitive understanding of self-attention

Video title	Description	Video

Key

Value

Query

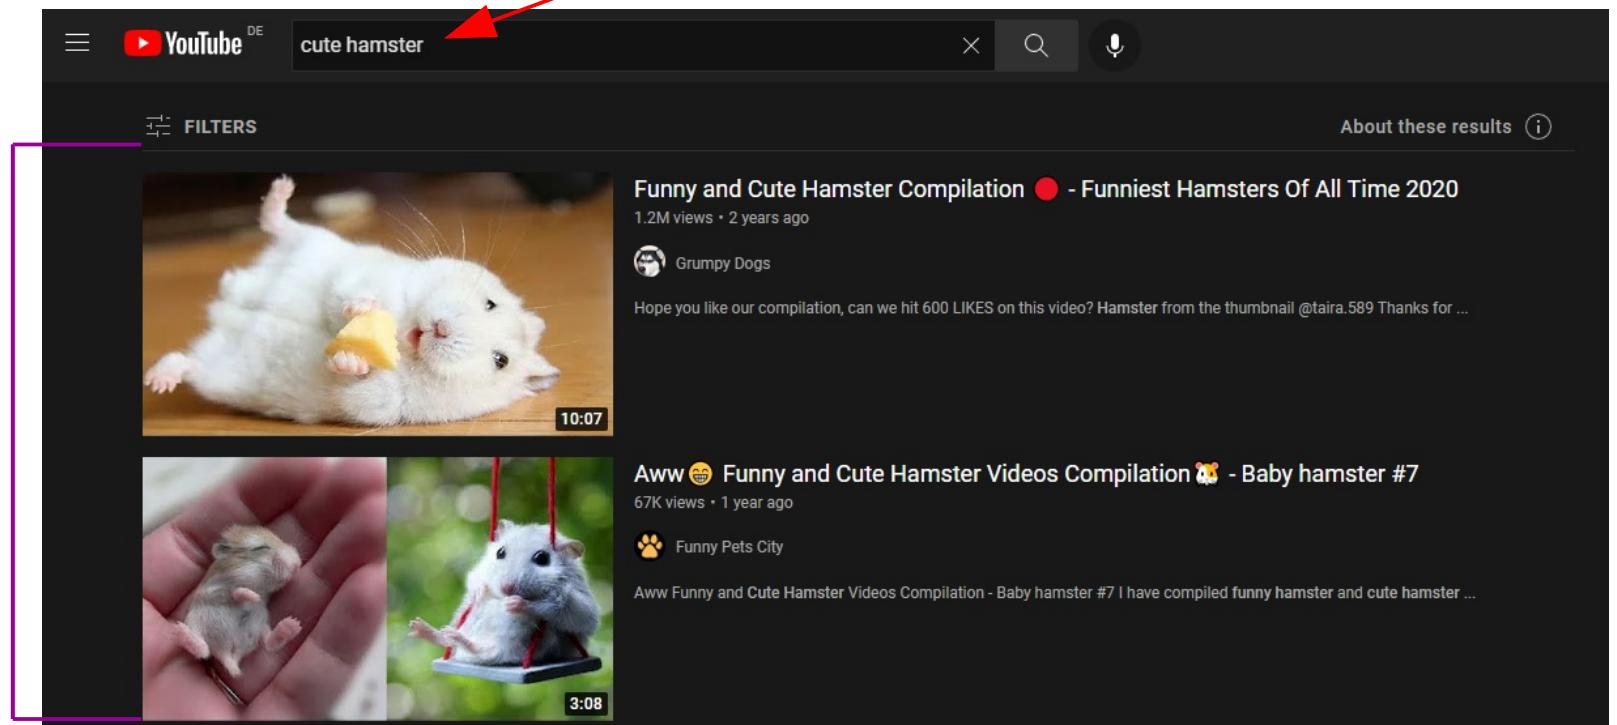


Figure 6: Self-attention example based on a YouTube search query.

2. Intuitive understanding of self-attention

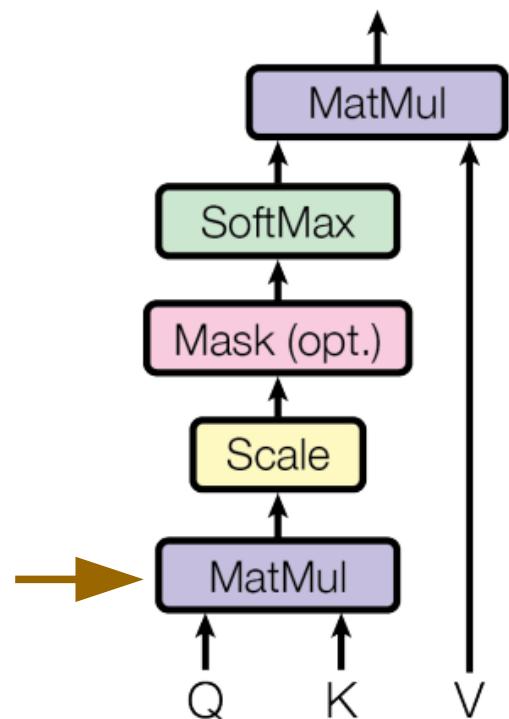


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

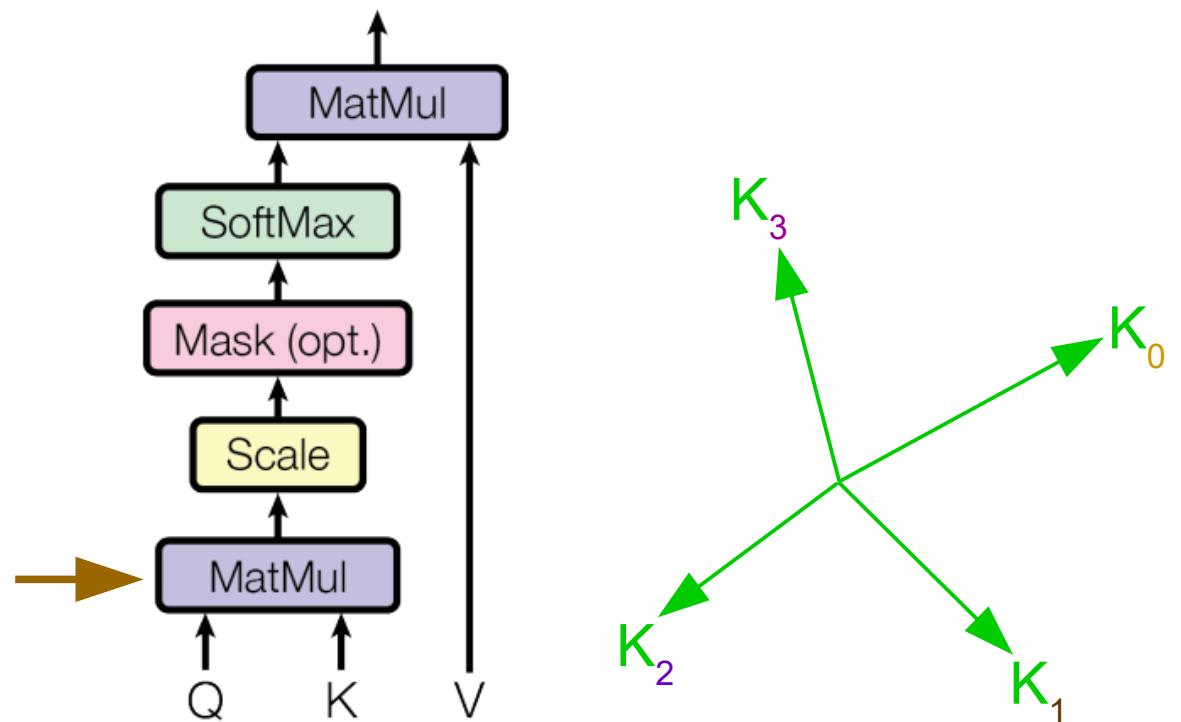


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

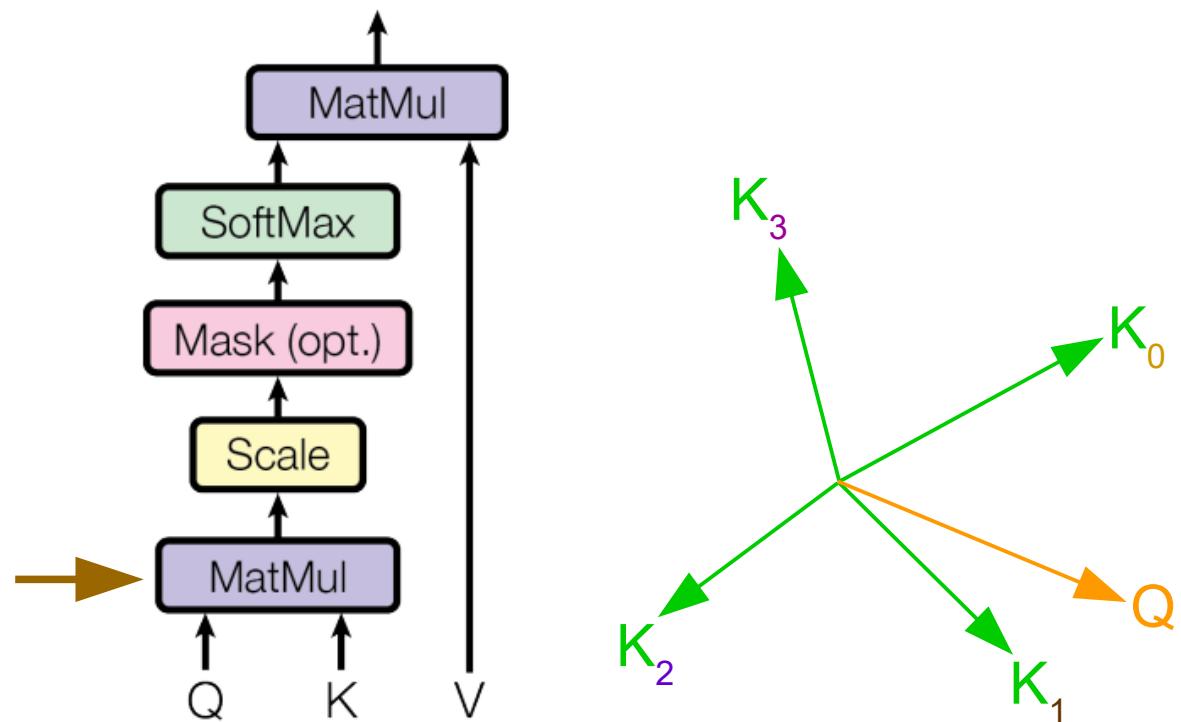


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

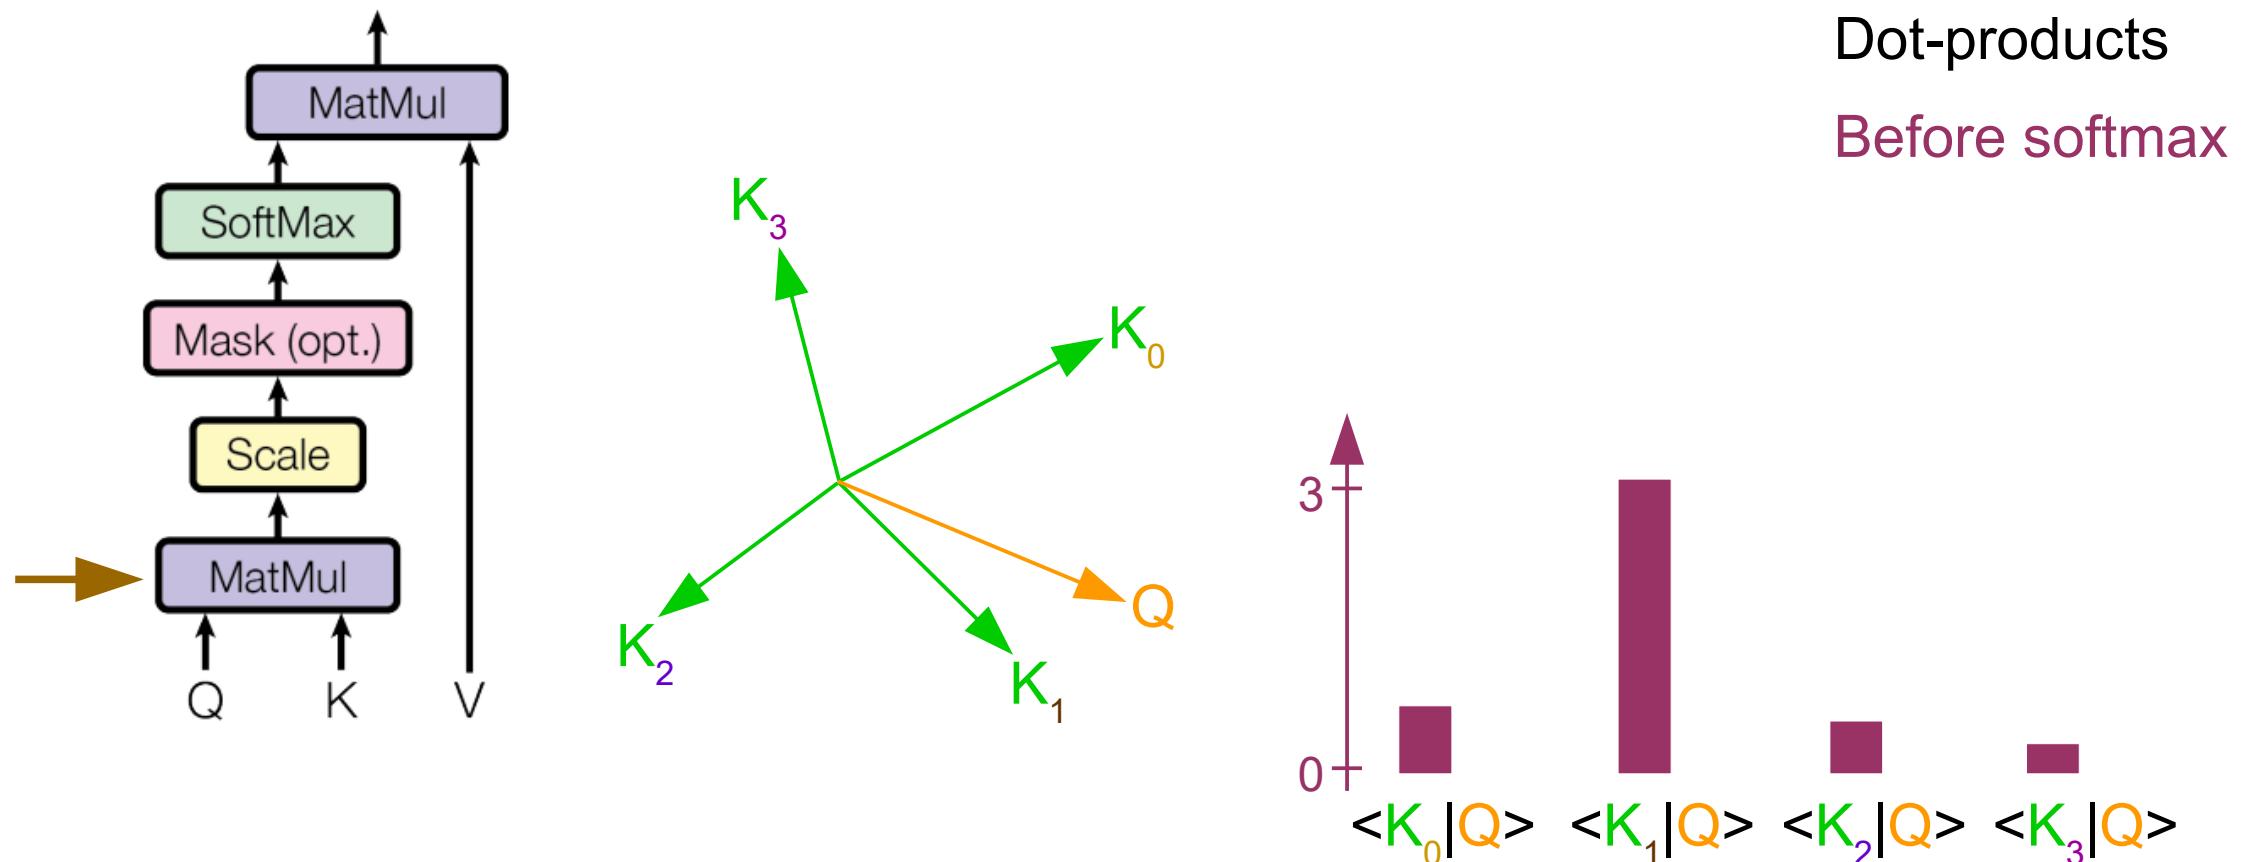


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention




Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

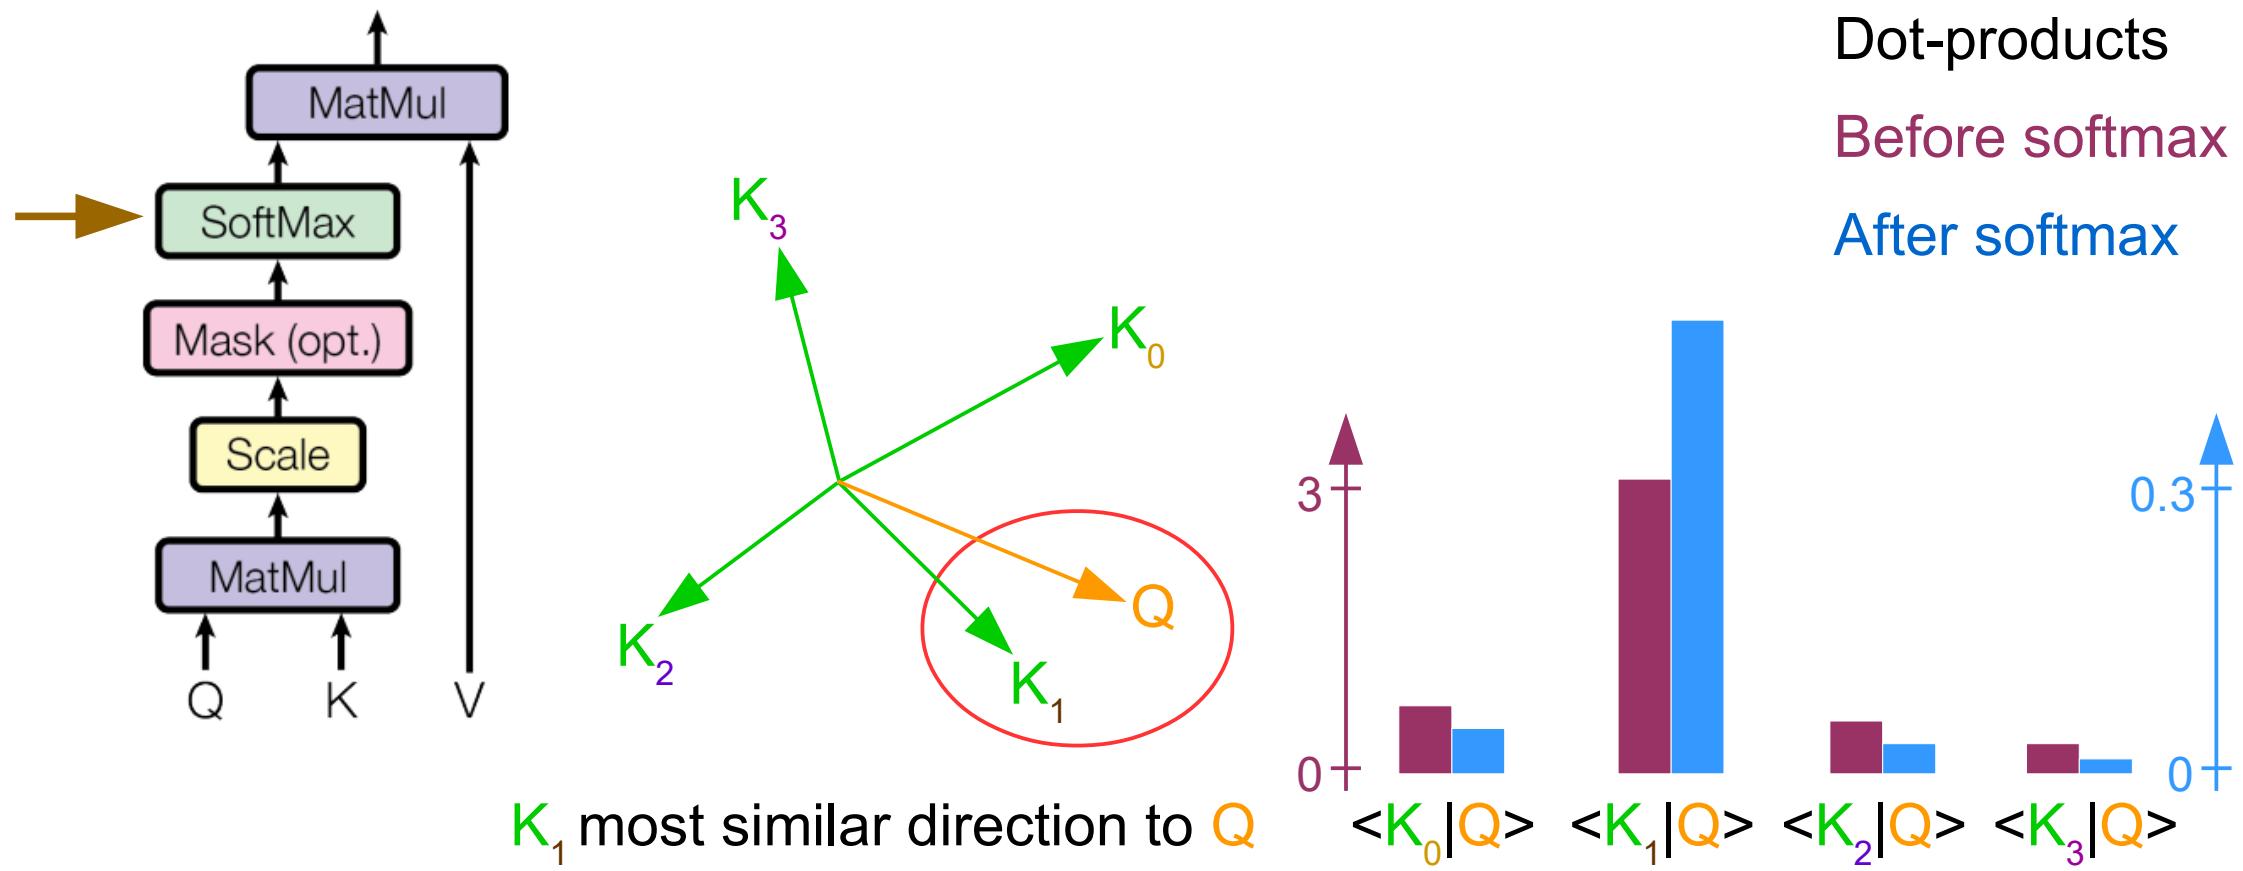


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

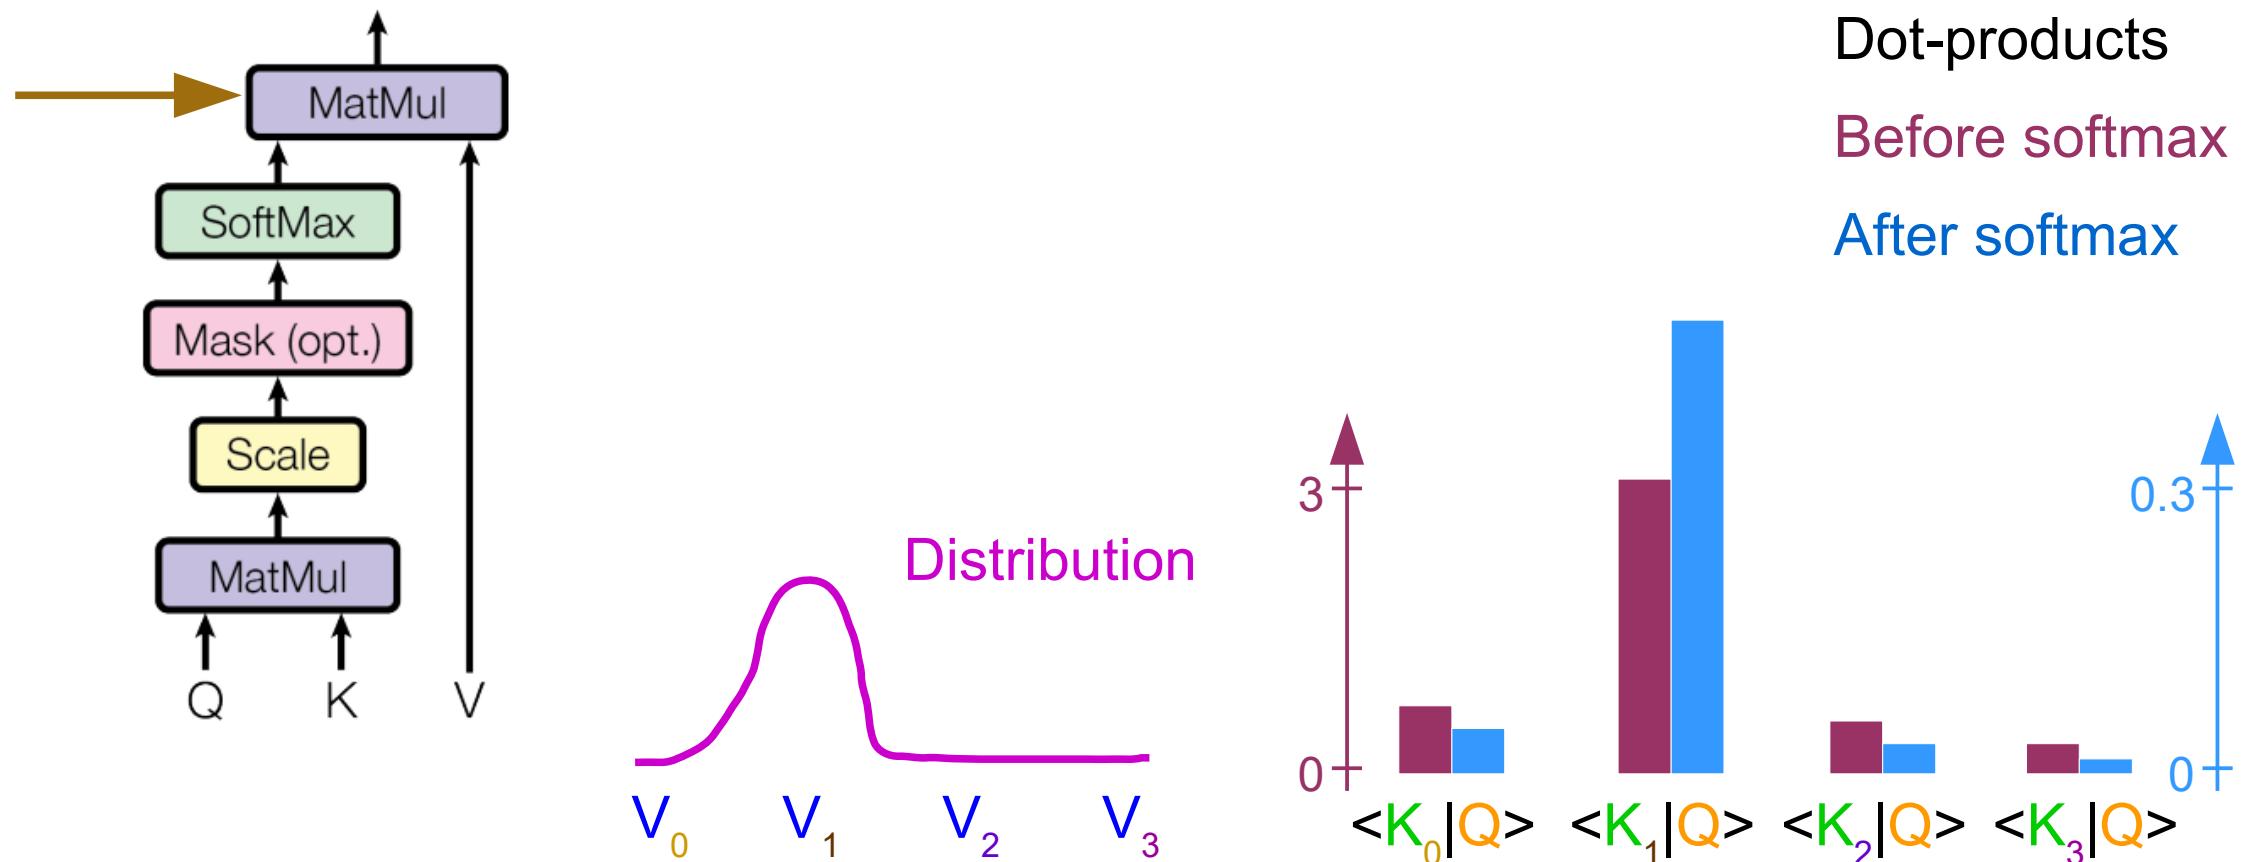


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

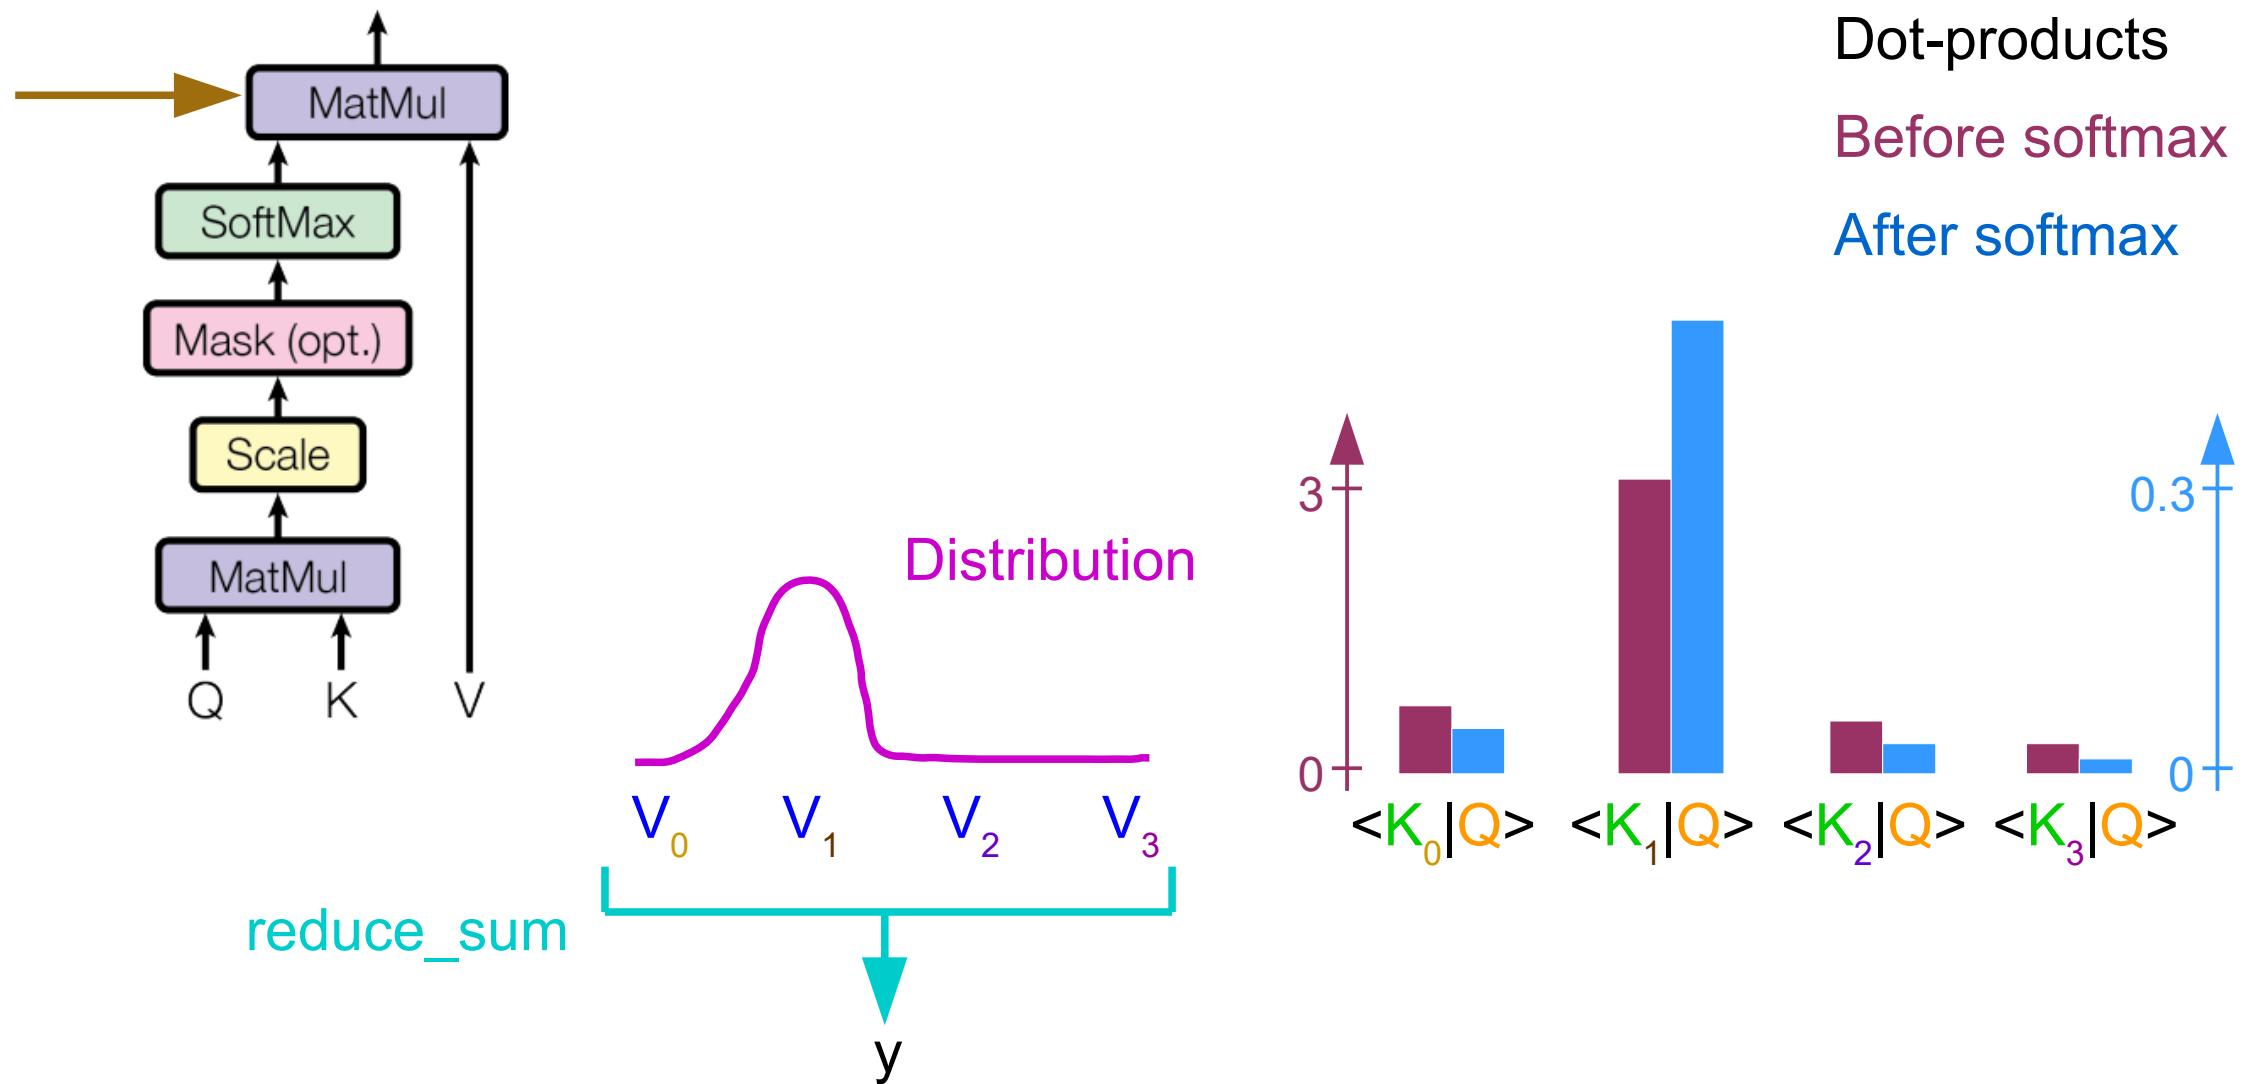


Figure 7: Self-attention visualised [2].

2. Intuitive understanding of self-attention

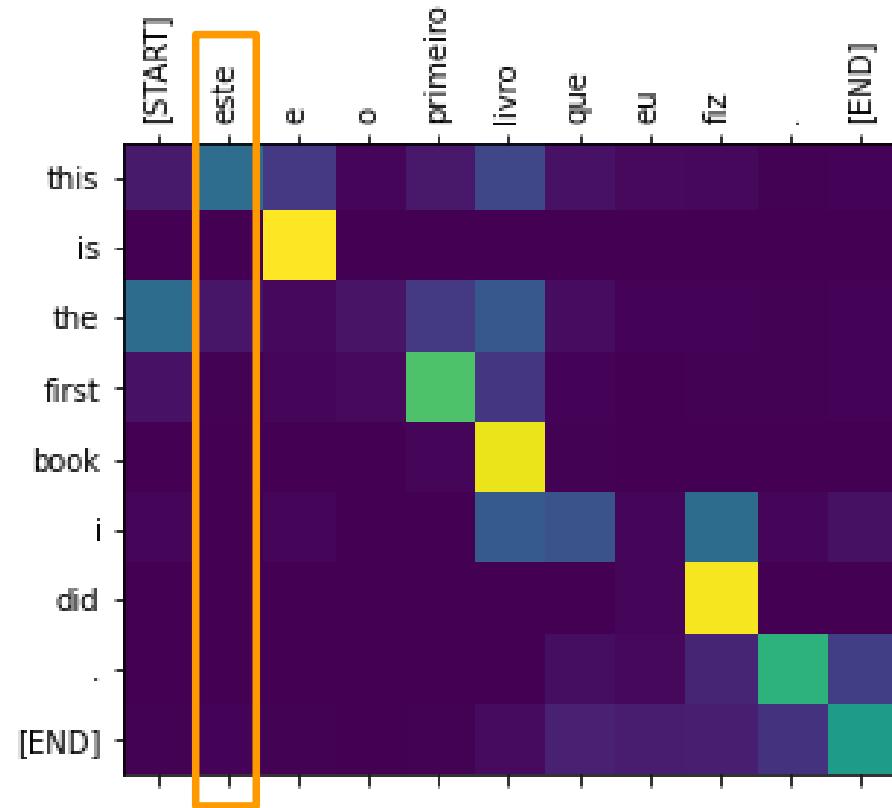



Figure 8: Self-attention – origin of the heatmap [2].

2. Intuitive understanding of self-attention

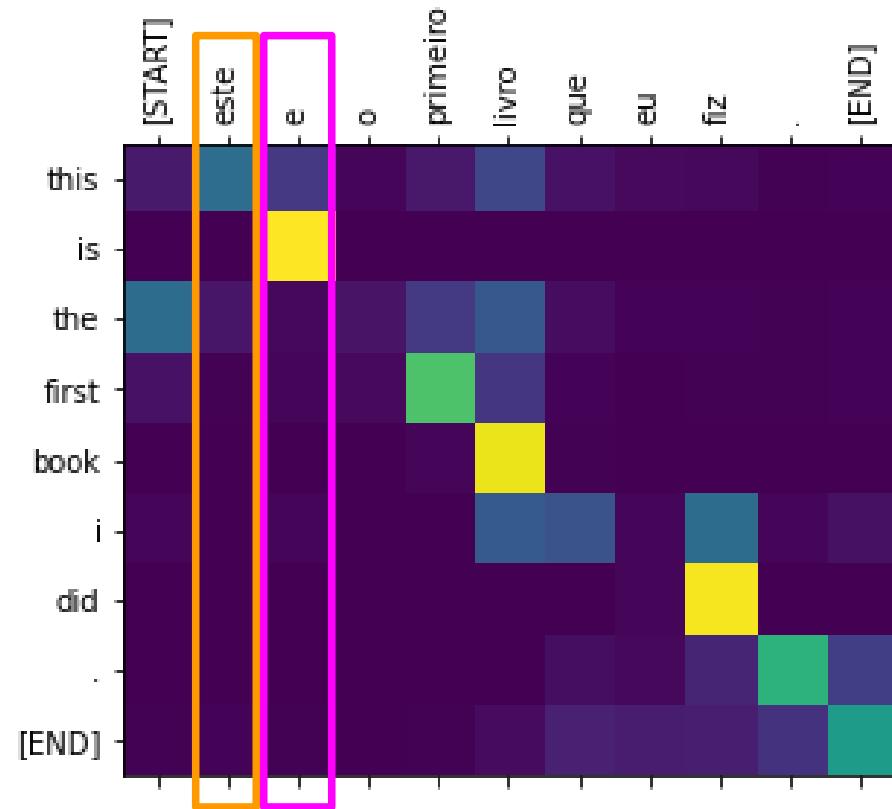
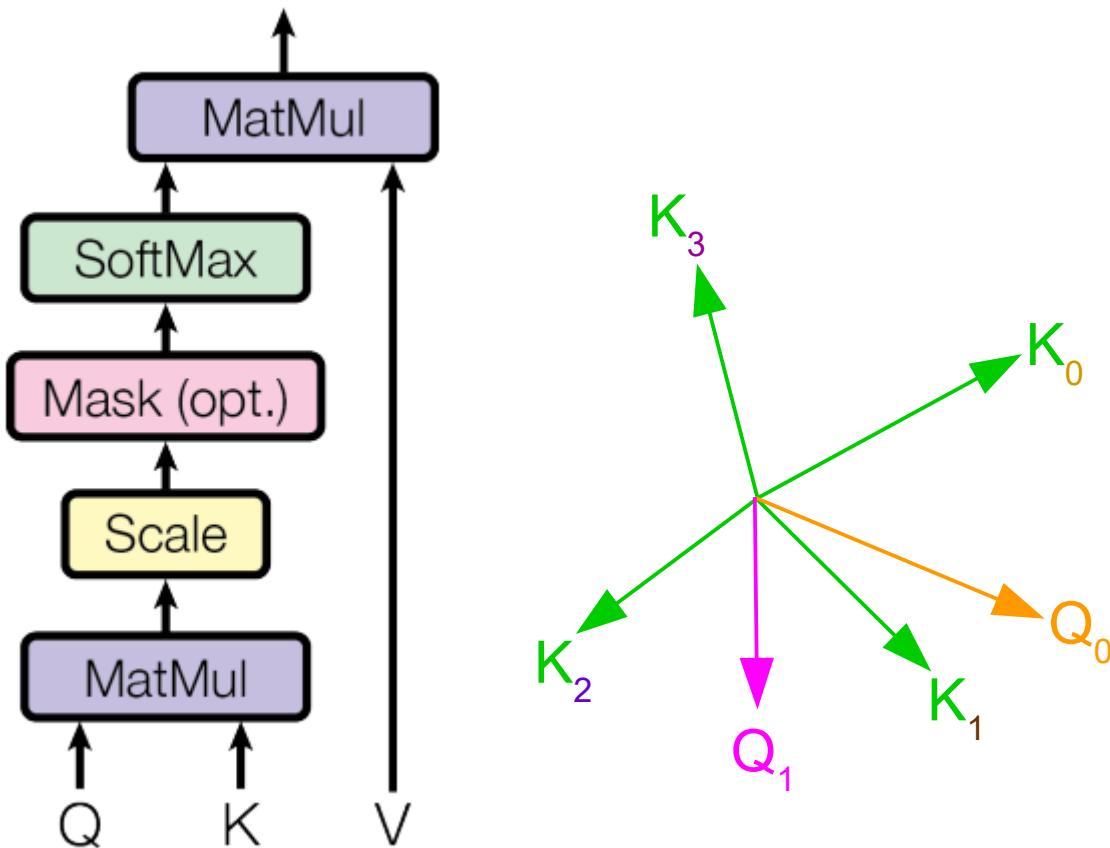


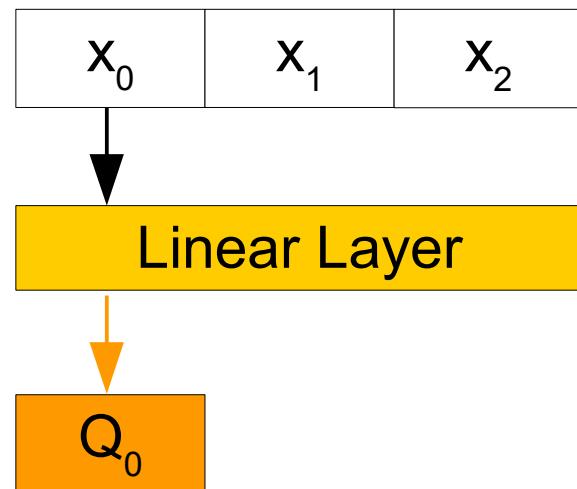
Figure 8: Self-attention – origin of the heatmap [2].

2. Intuitive understanding of self-attention


x_0	x_1	x_2
-------	-------	-------

Linear Layer


2. Intuitive understanding of self-attention



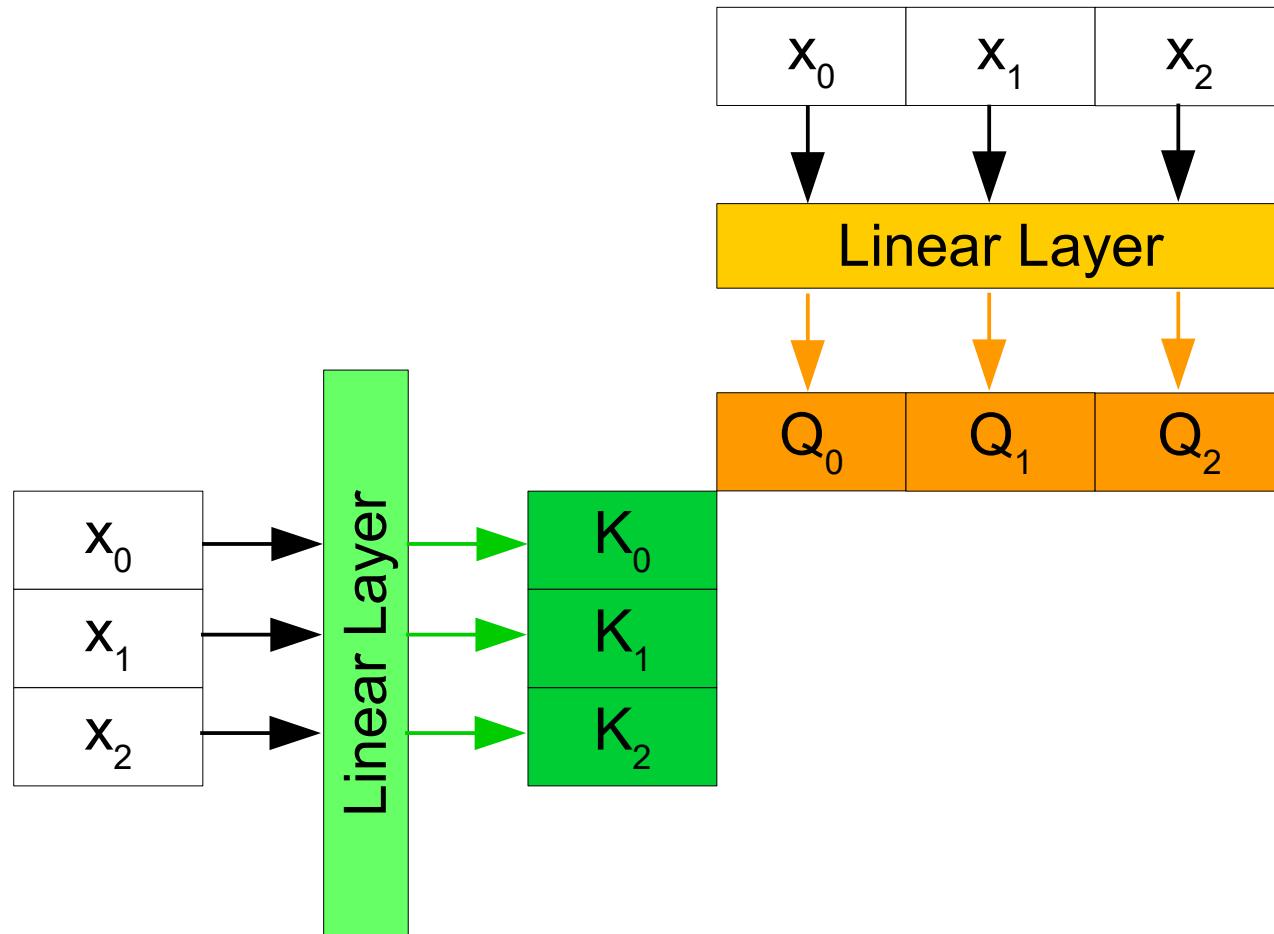
2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention



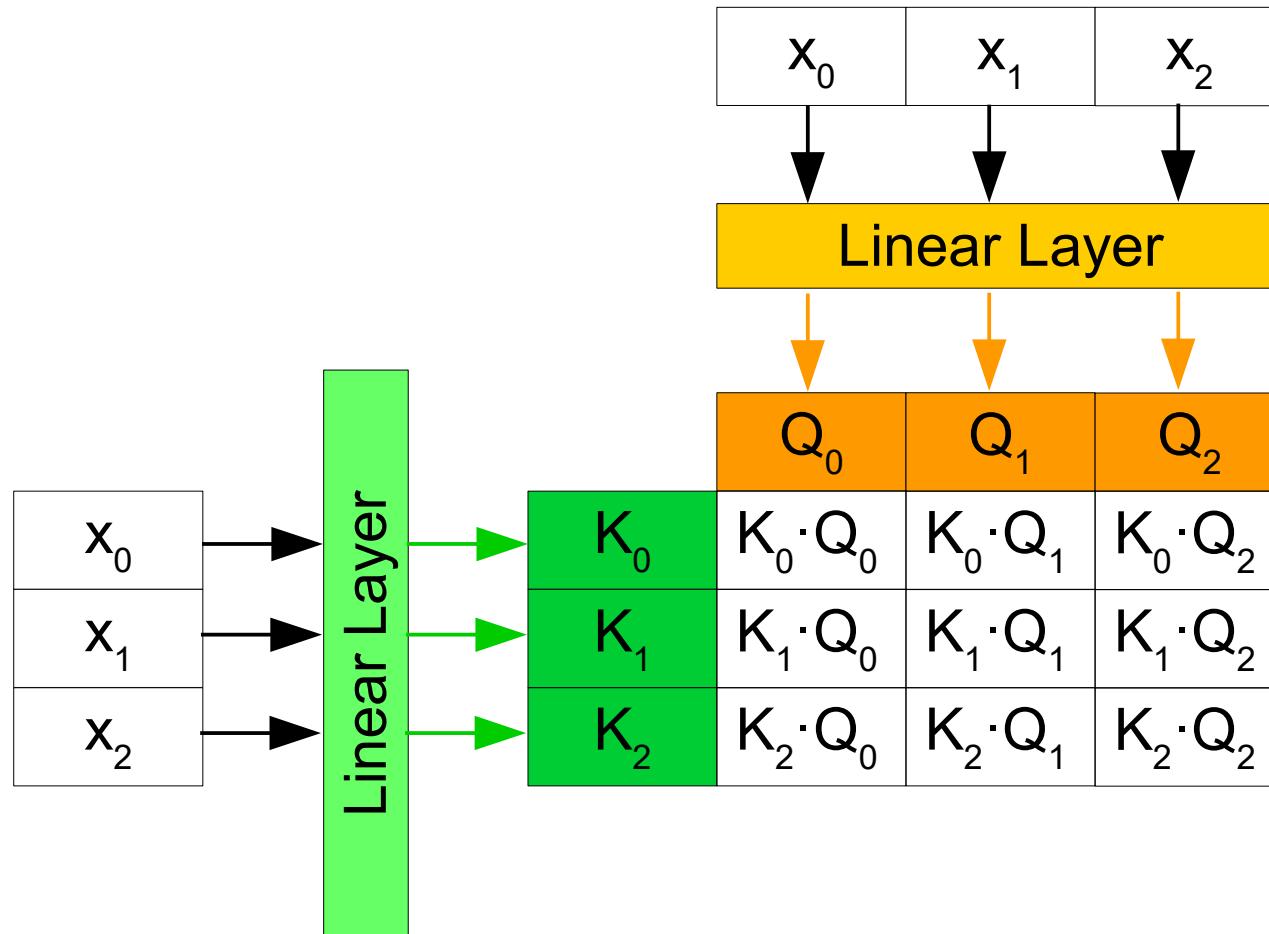
2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention



2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention




$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

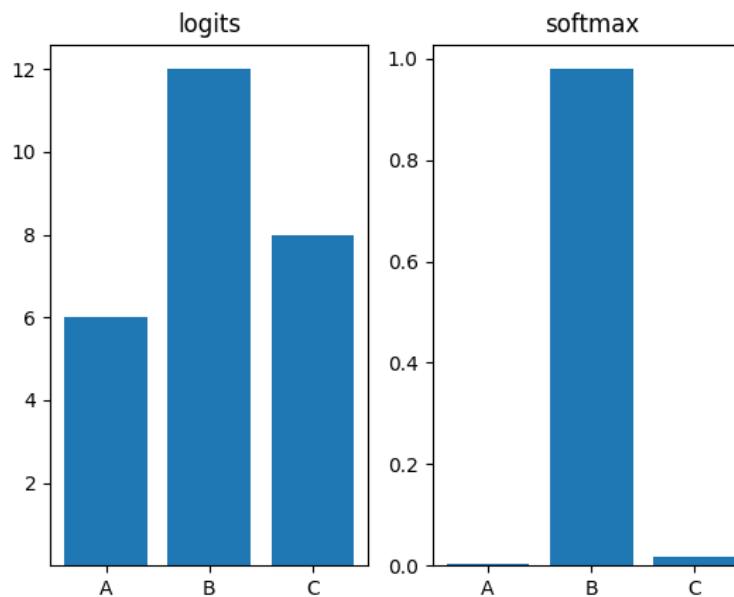
Side note: softmax is not scale invariant



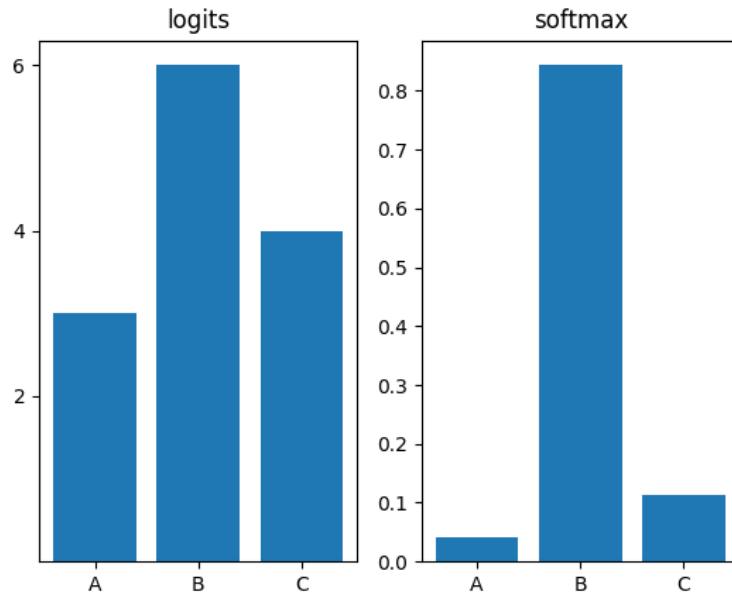
Side note: softmax is not scale invariant




$2x$




Side note: softmax is not scale invariant



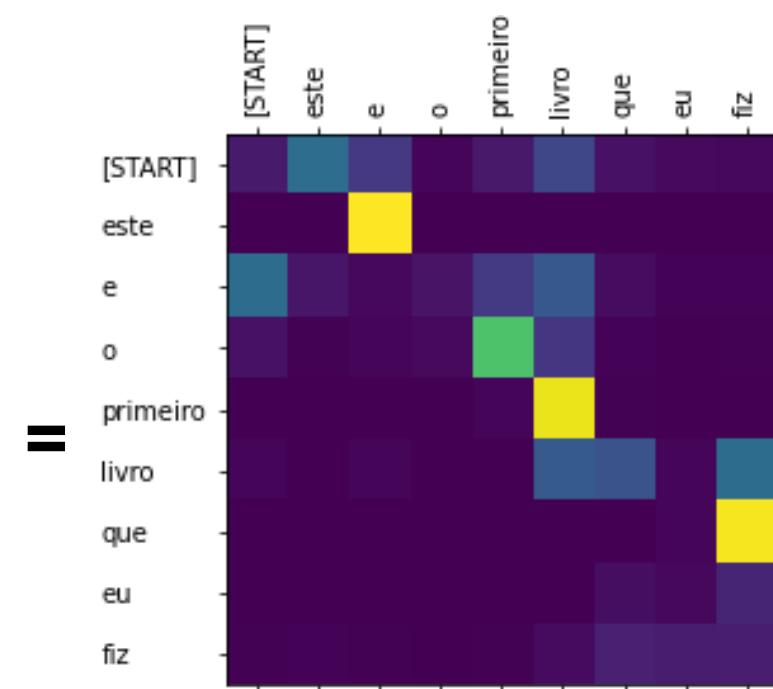
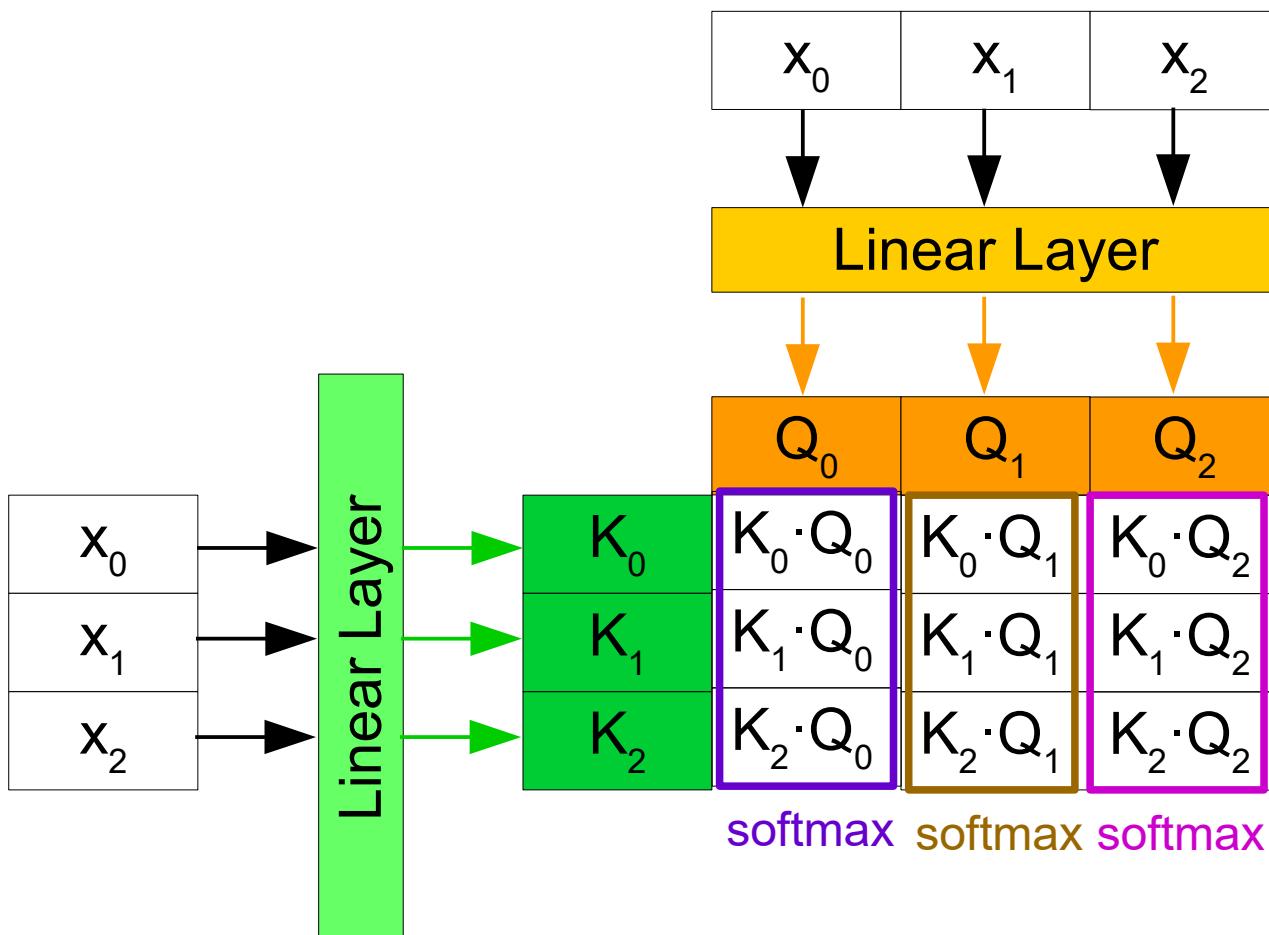
$2x$



2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention

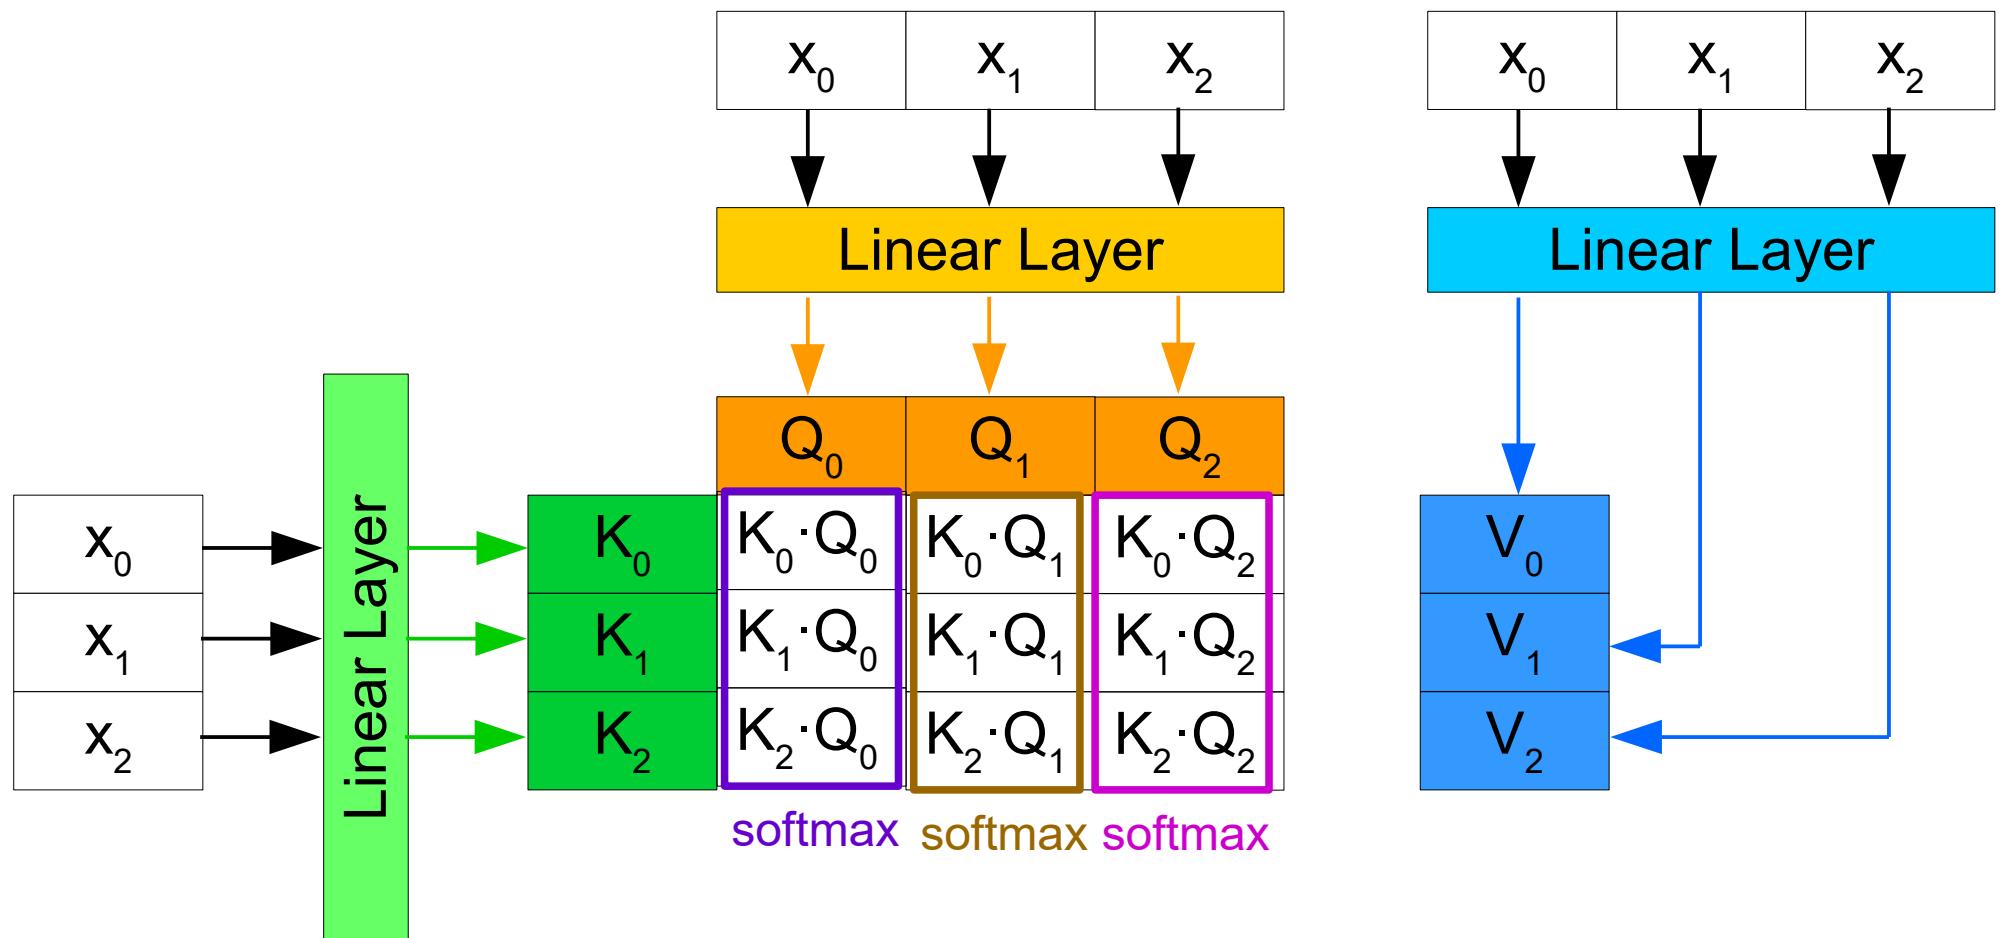


Here: 3 words instead of 9 ;)


2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention




2. Intuitive understanding of self-attention



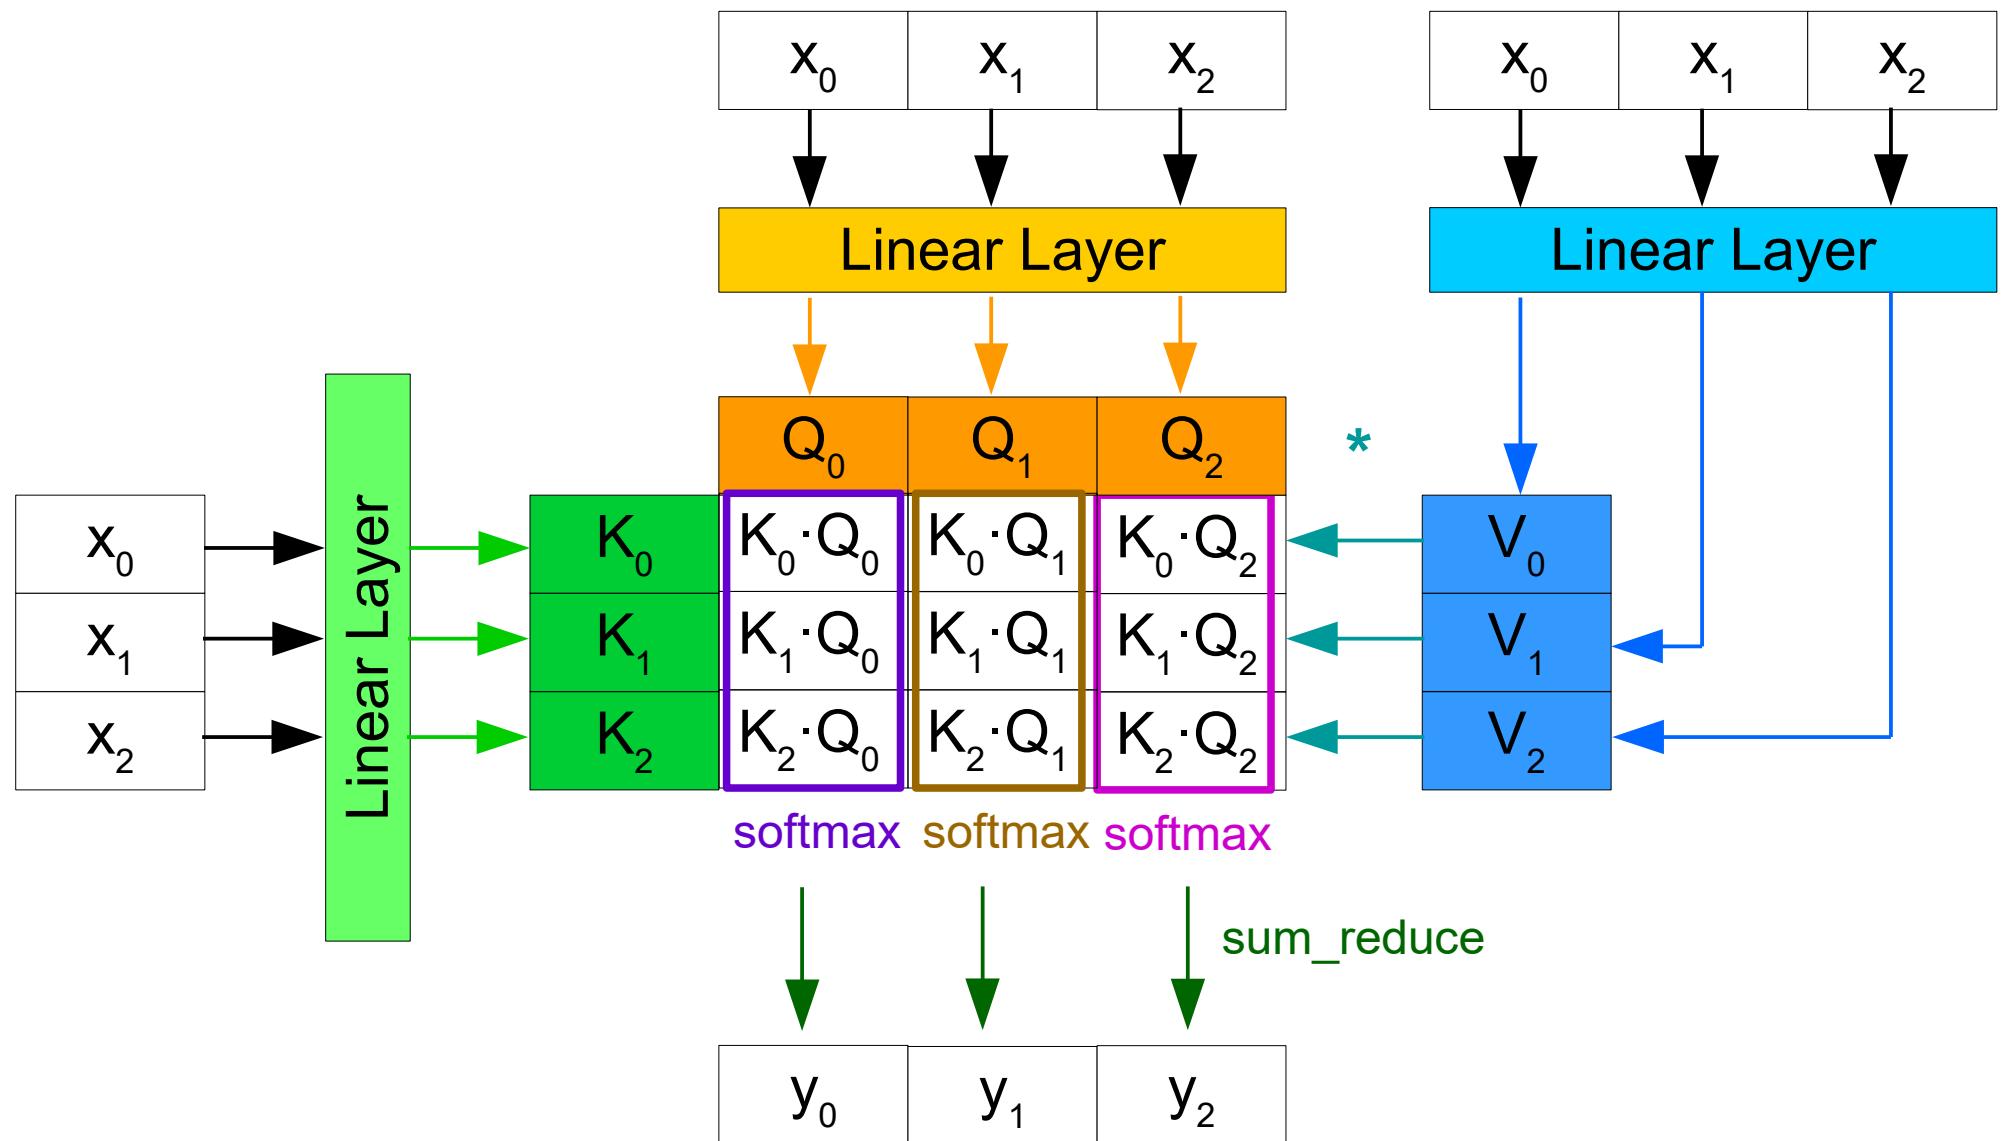
2. Intuitive understanding of self-attention



2. Intuitive understanding of self-attention



2. Intuitive understanding of self-attention



2.1 Multi-head attention

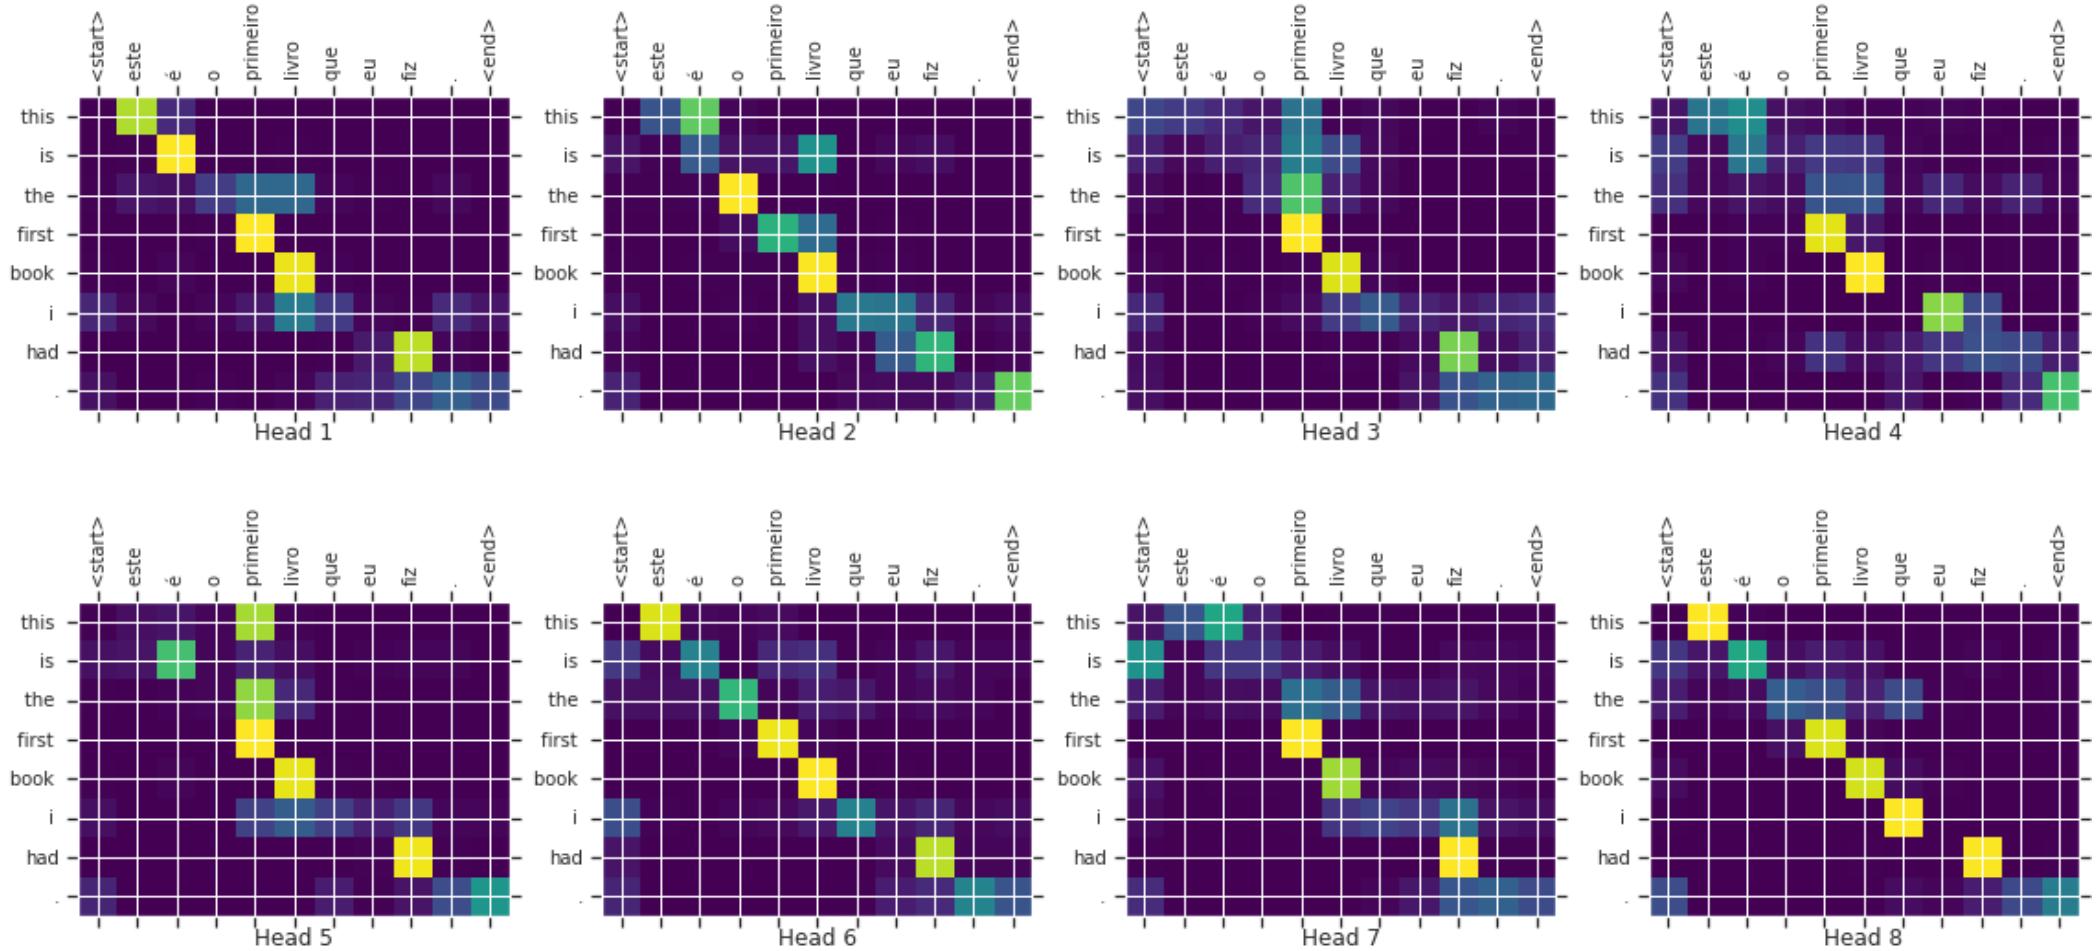


Figure 9: Heatmaps from Multi-head attention [2].

2.1 Multi-head attention


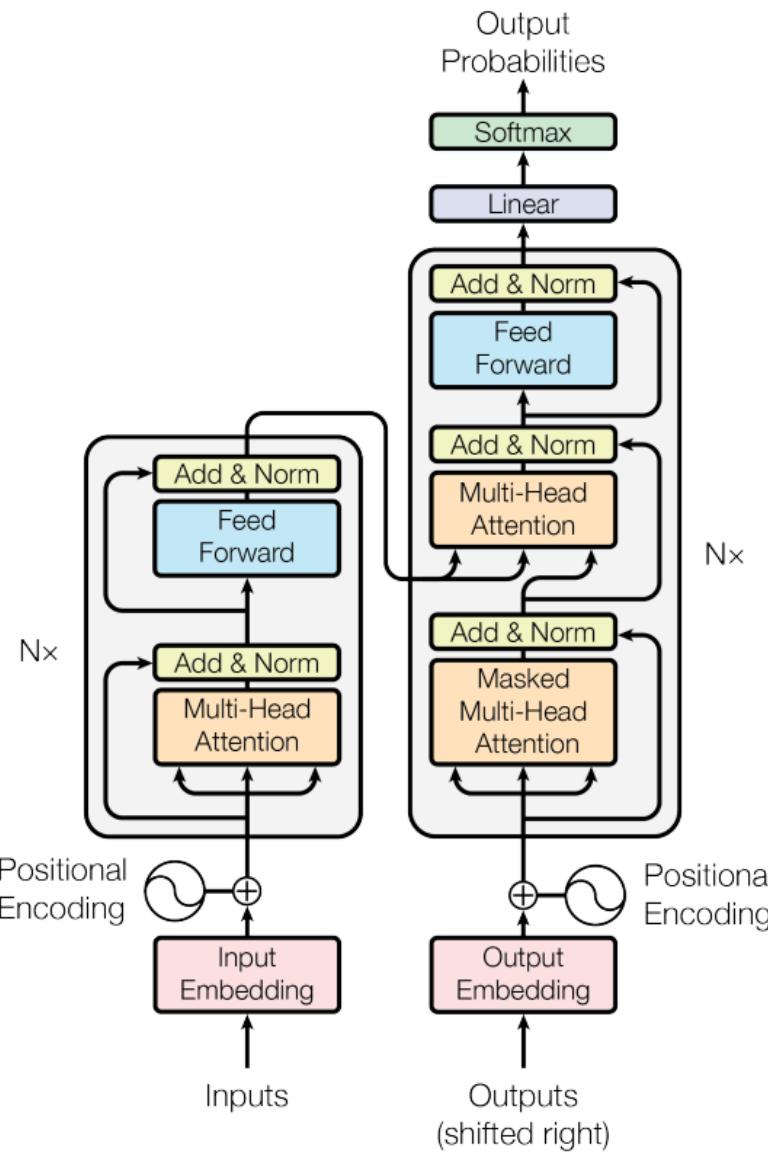


Figure 10: Overview of multi-head attention [2].



3. Intuitive understanding of the transformer architecture

Image sources:

https://machinelearningmastery.com/wp-content/uploads/2021/08/attention_research_1-727x1024.png
(call date: 20.07.22)

3.1 Let's build a transformer

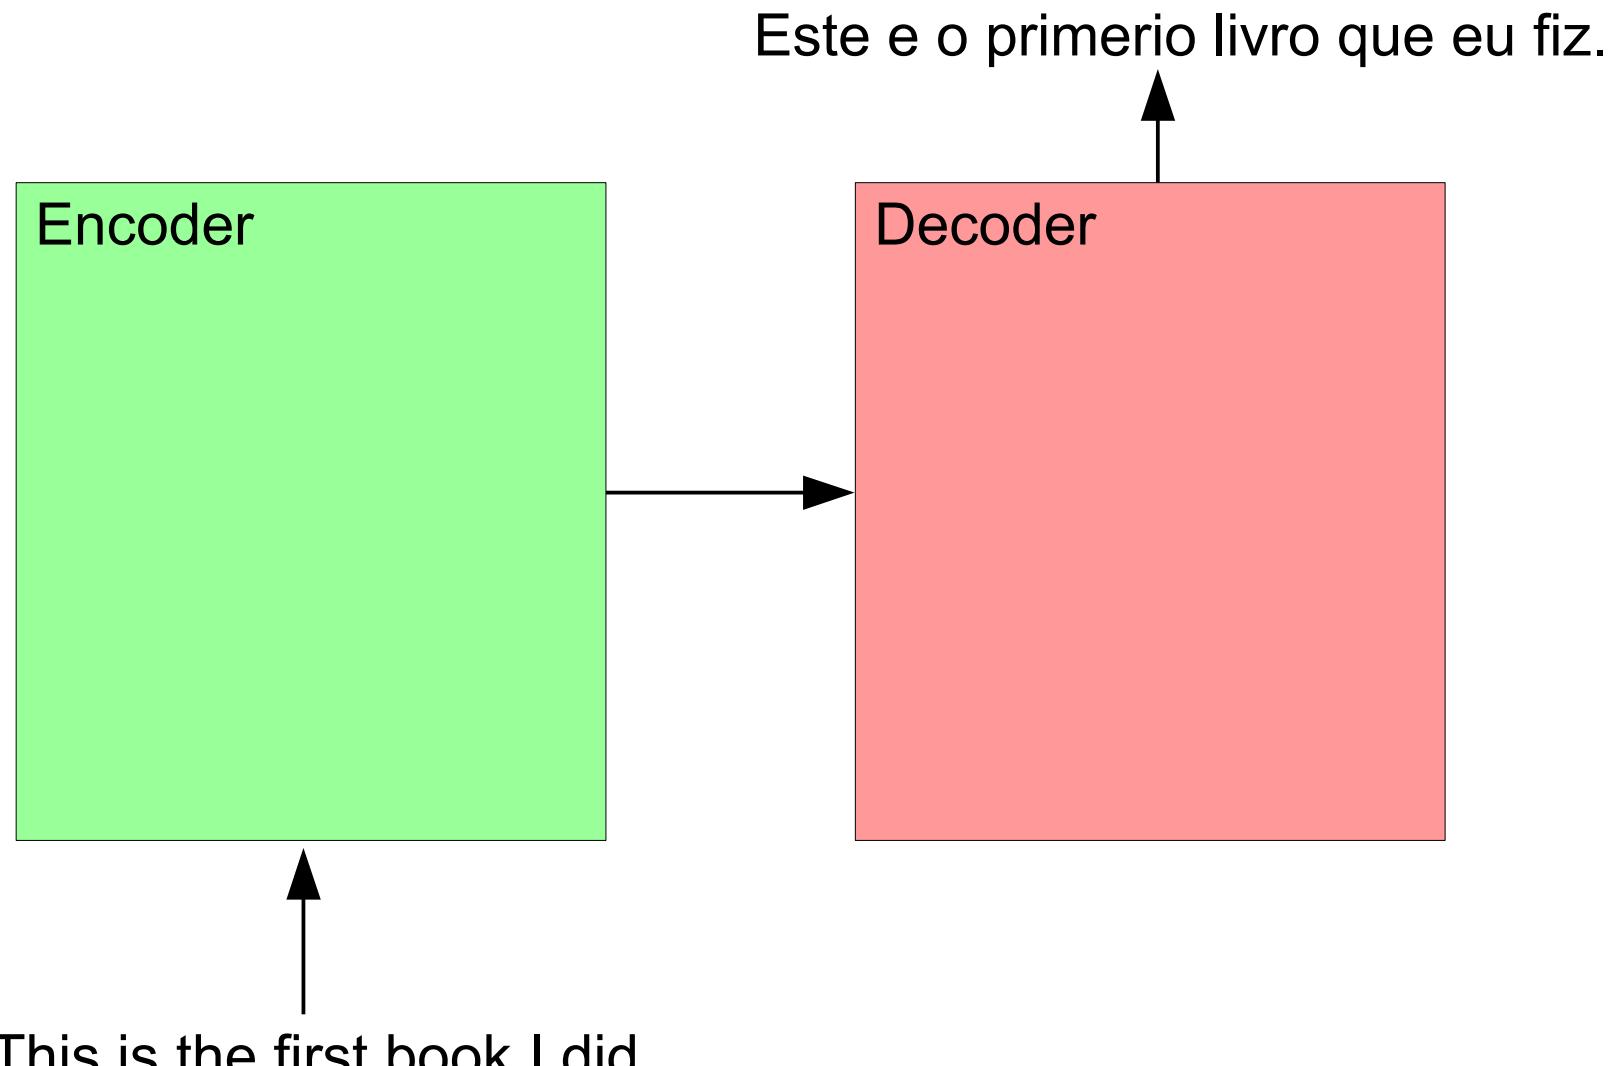


Figure 11: Neural machine translation based on a transformer [2].

3.1 Let's build a transformer

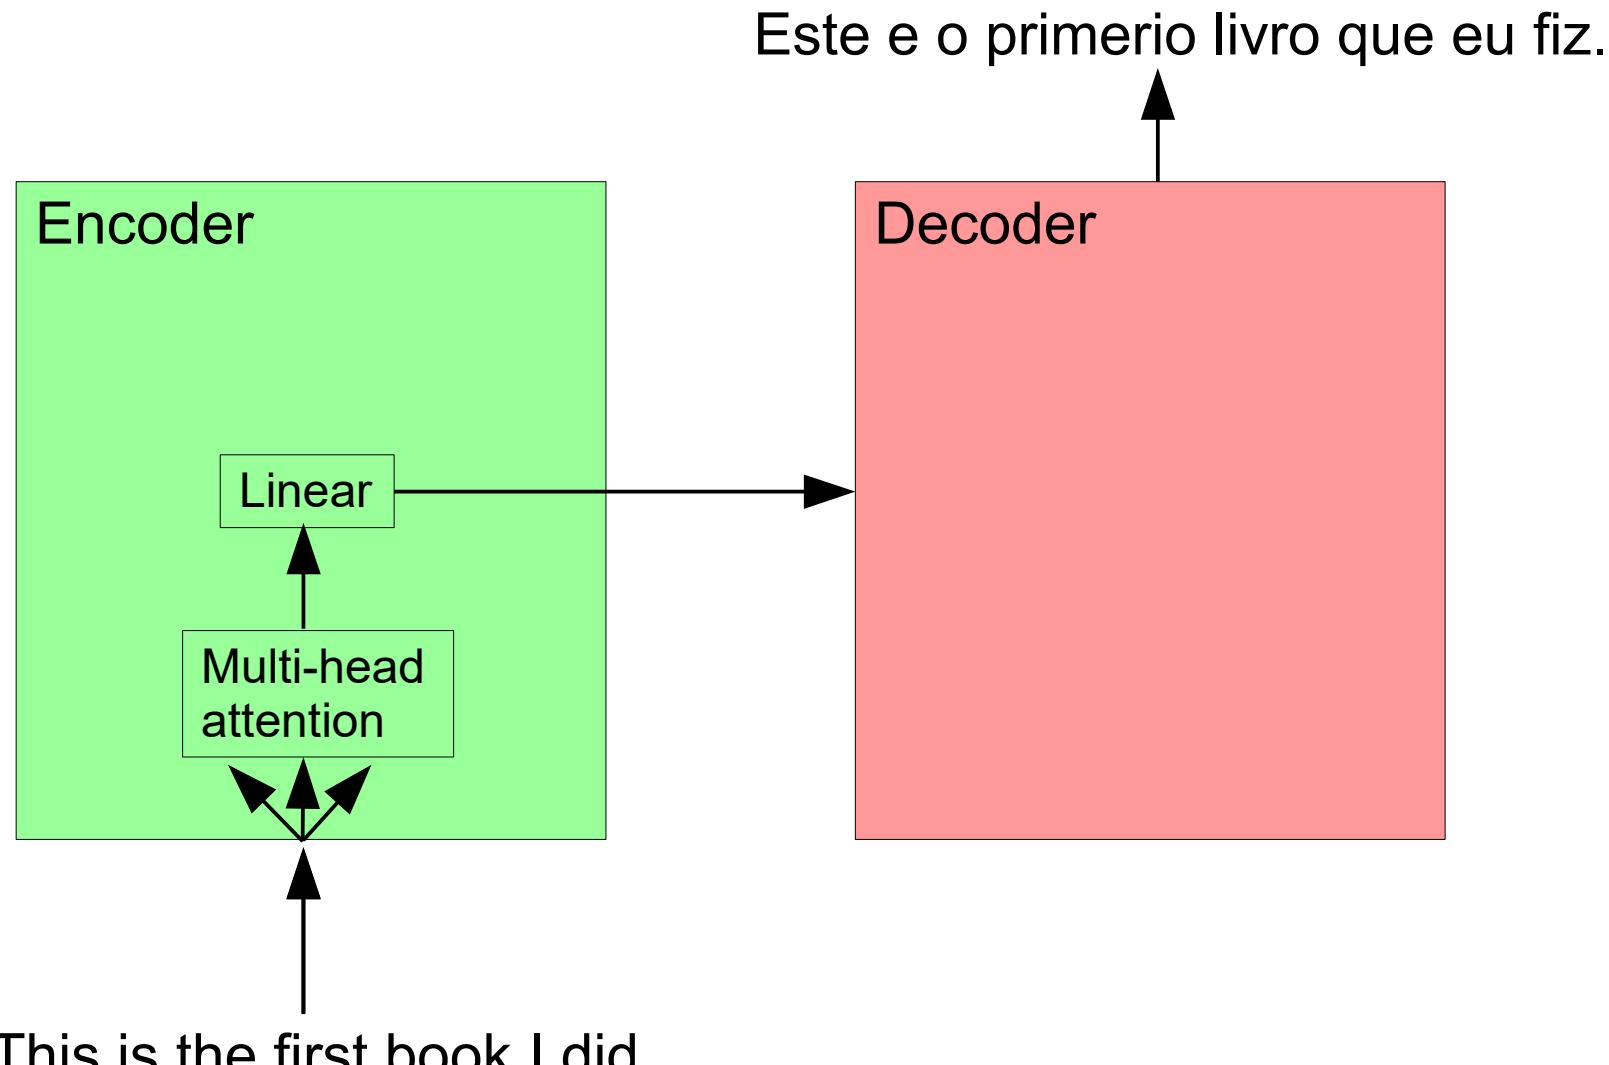


Figure 11: Neural machine translation based on a transformer [2].

3.1 Let's build a transformer

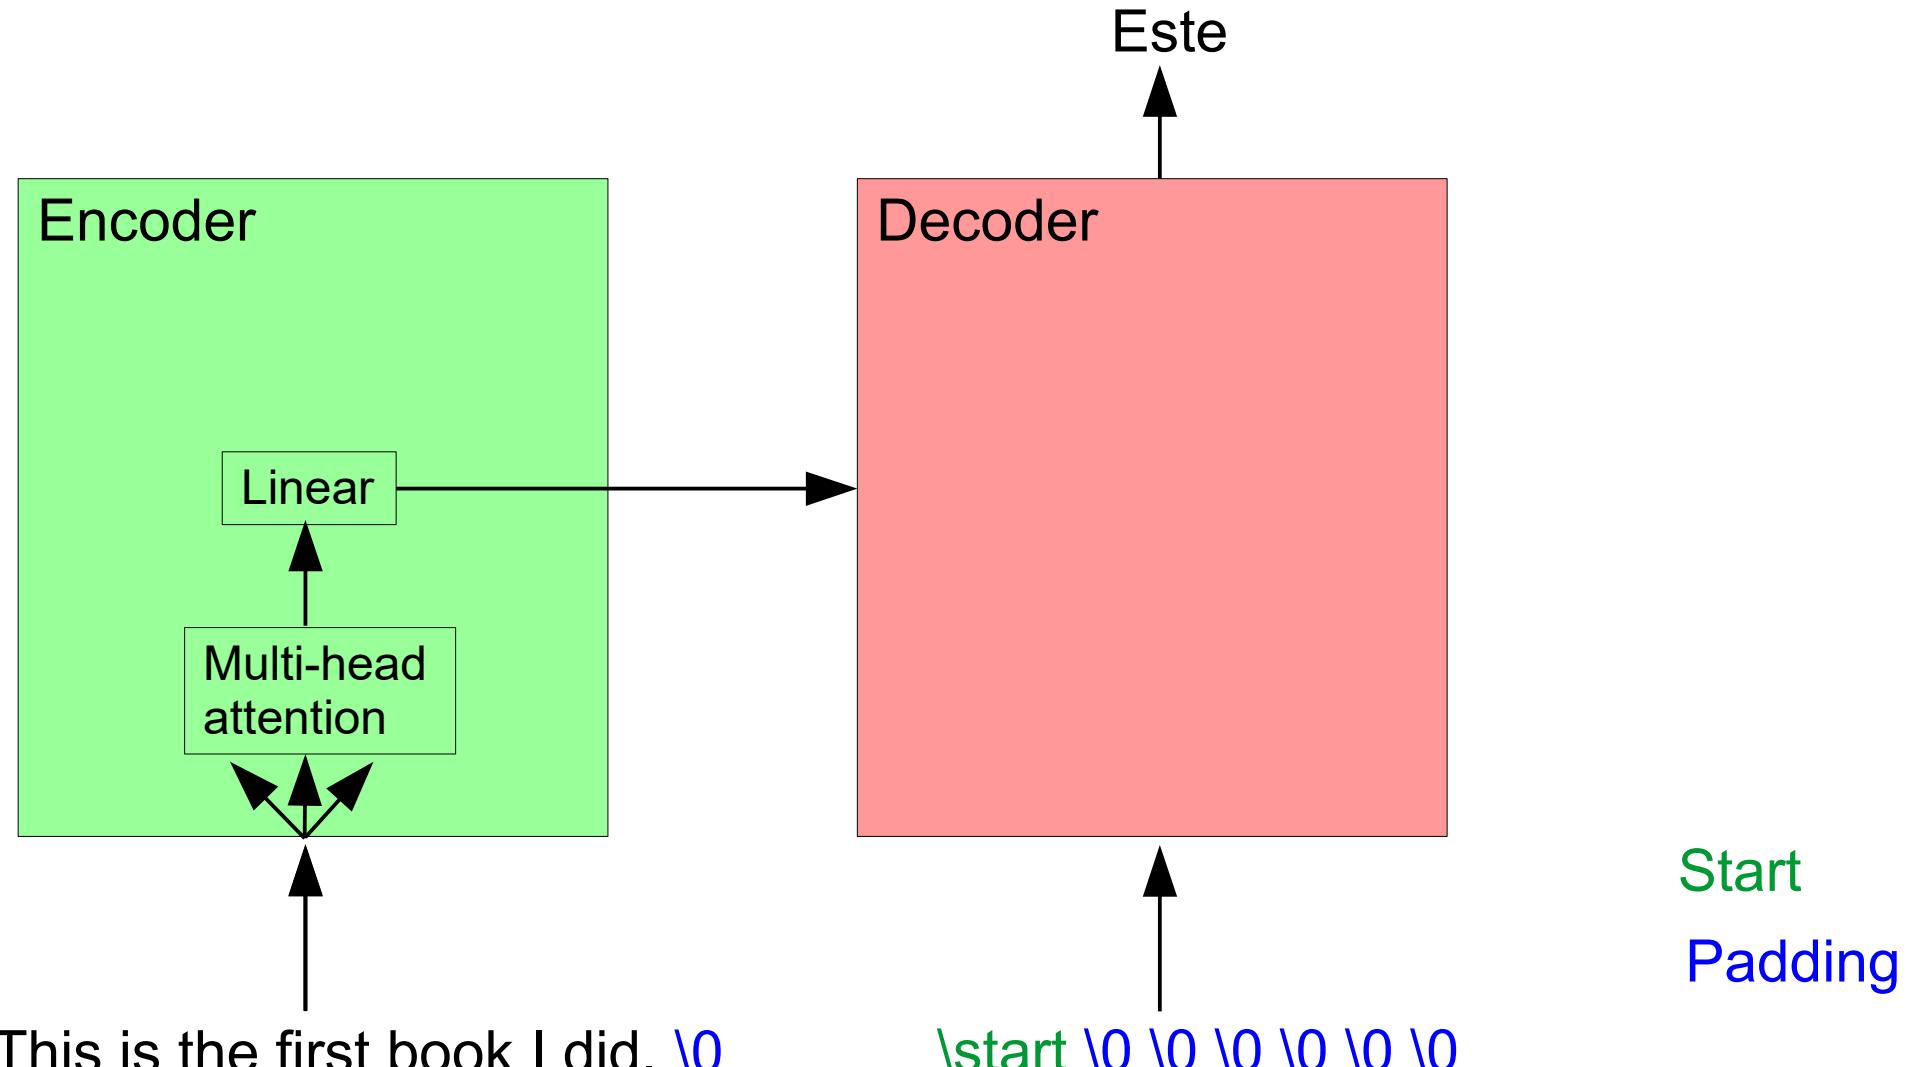


Figure 11: Neural machine translation based on a transformer [2].

3.1 Let's build a transformer

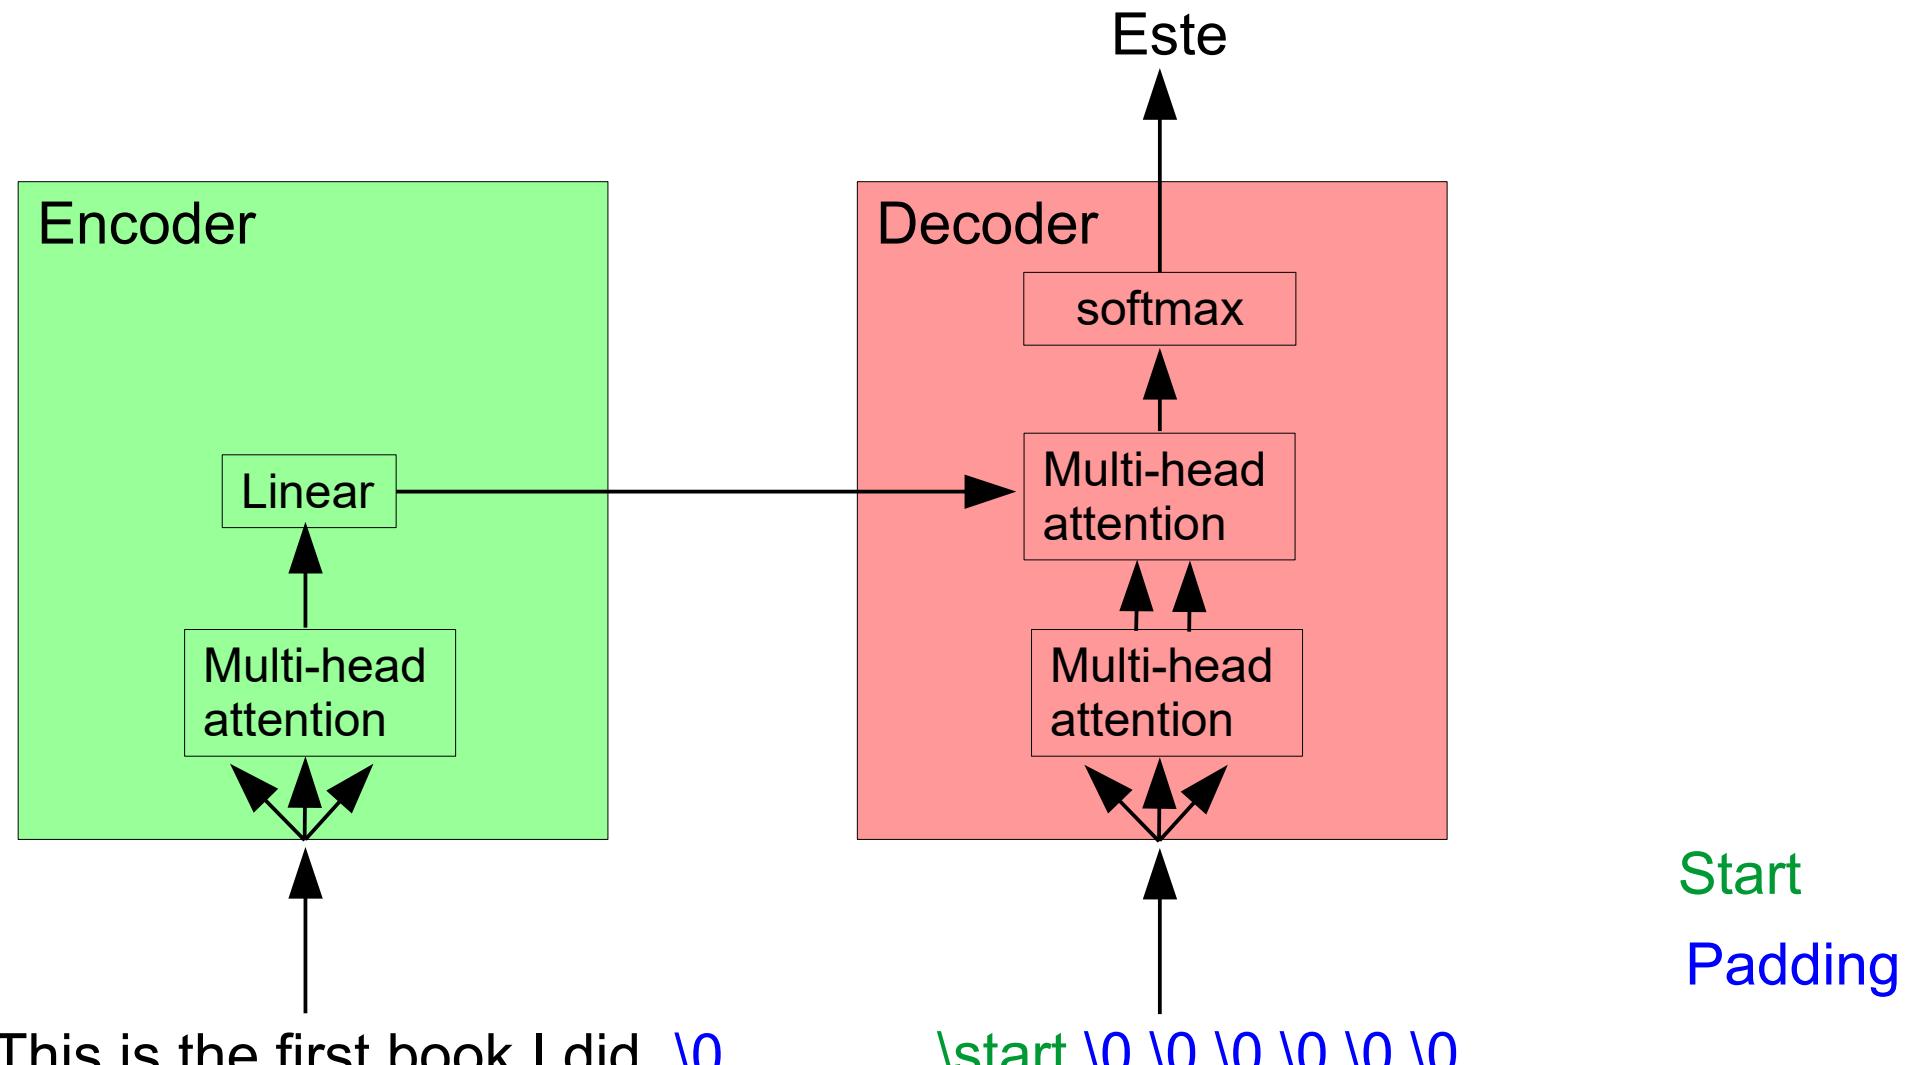


Figure 11: Neural machine translation based on a transformer [2].

3.1 Let's build a transformer

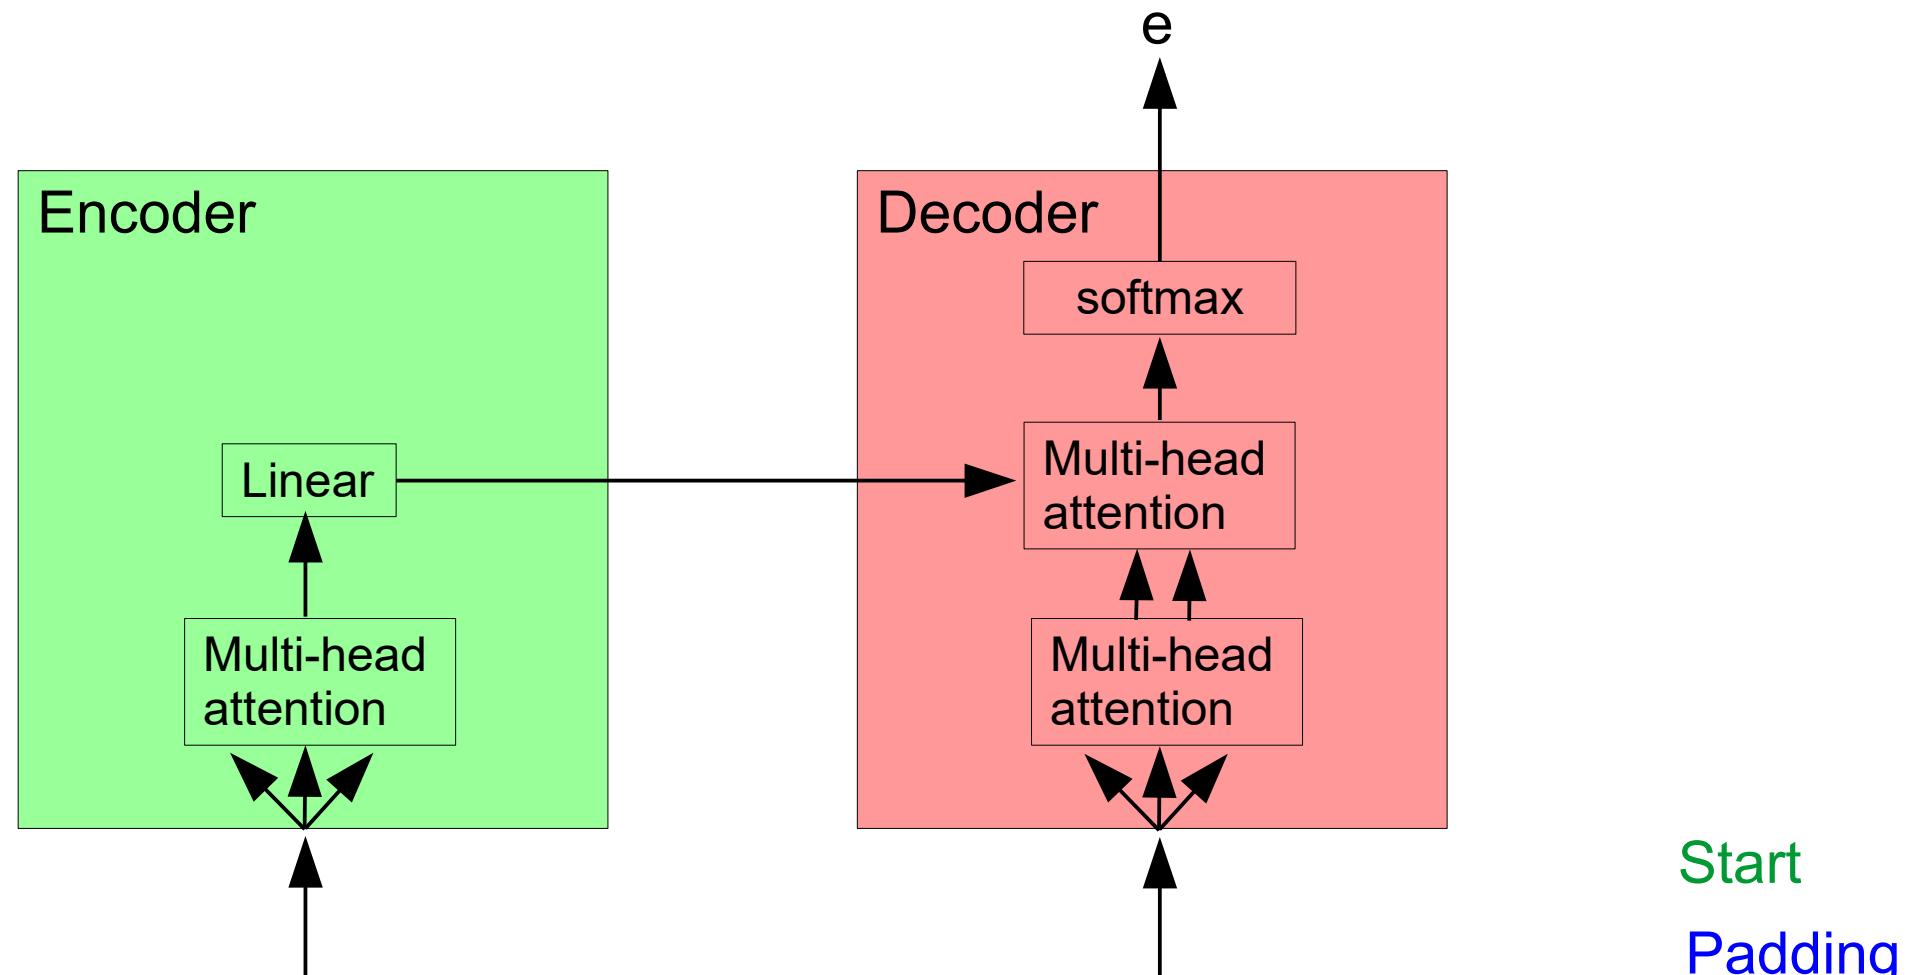


Figure 11: Neural machine translation based on a transformer [2].

3.1 Let's build a transformer

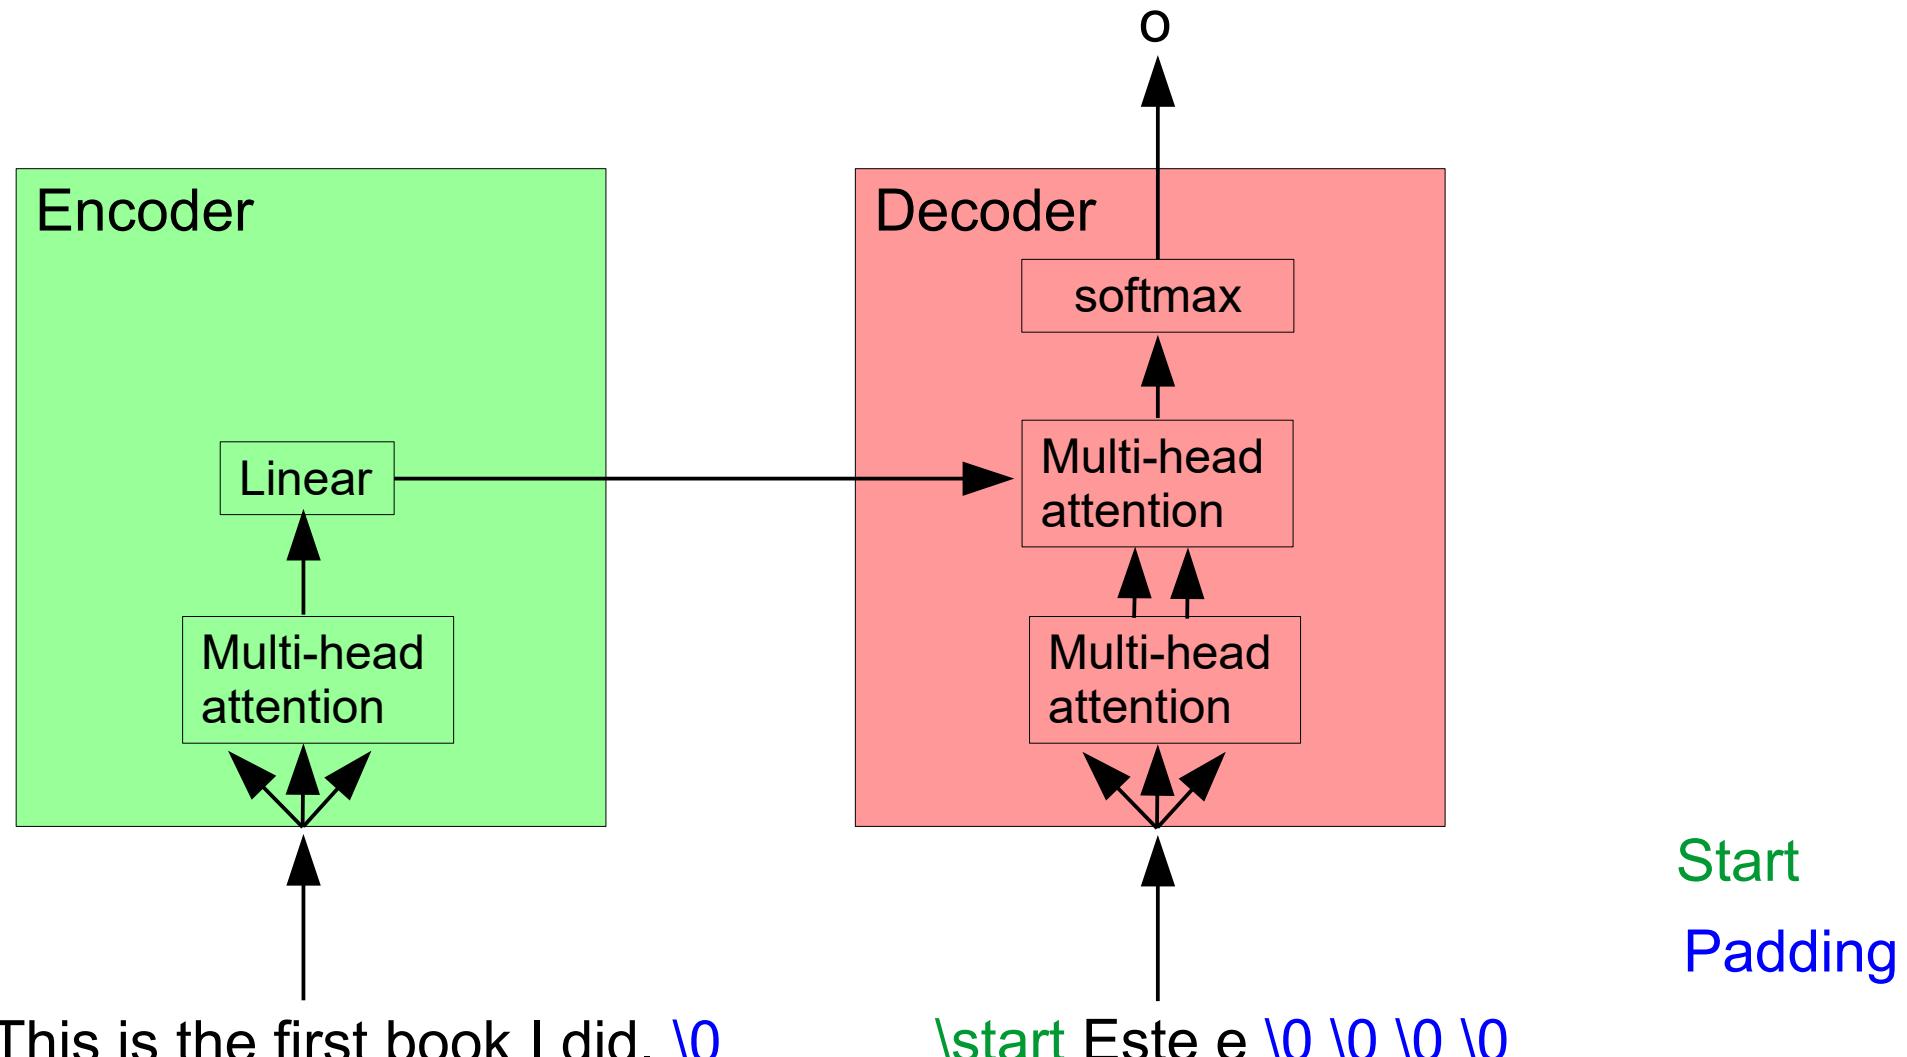


Figure 11: Neural machine translation based on a transformer [2].

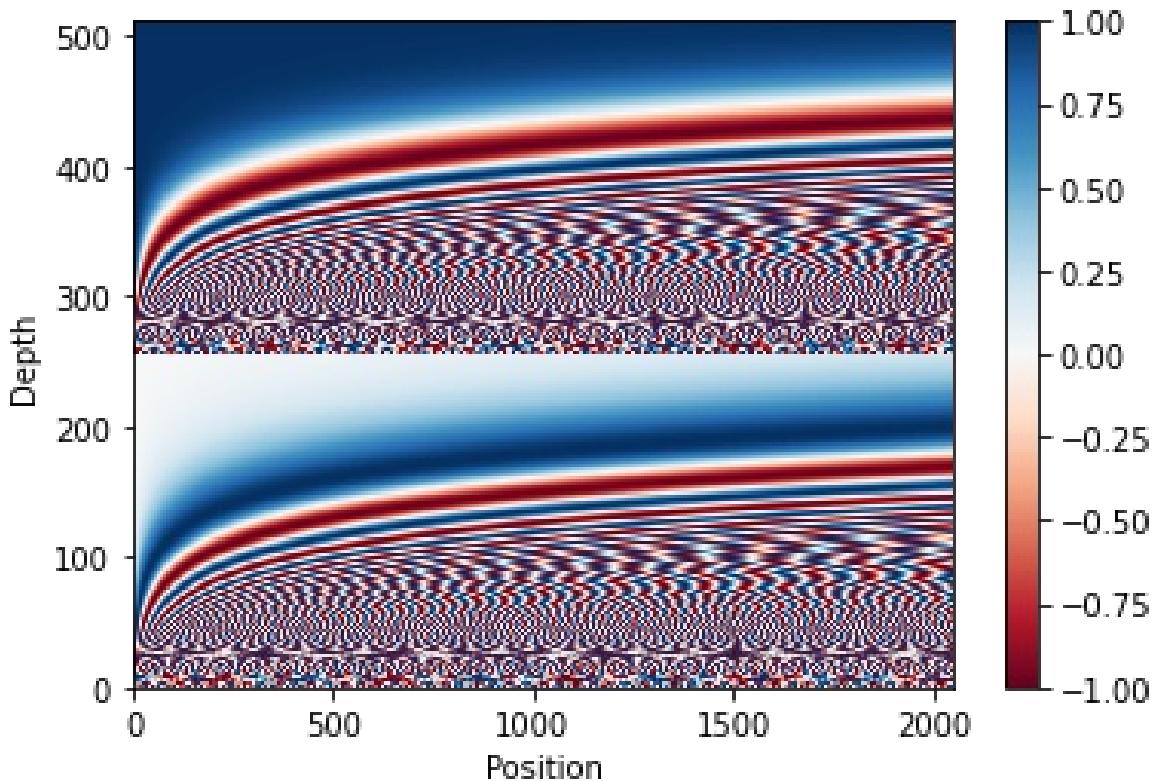
3.2 Positional encoding

$$PE_{pos, 2i} = \sin(pos / 10000^{2i/d_{model}})$$

$$PE_{pos, 2i+1} = \cos(pos / 10000^{2i/d_{model}})$$

depth/dimension index

3.2 Positional encoding



$$PE_{pos, 2i} = \sin(pos / 10000^{2i/d_{model}})$$

$$PE_{pos, 2i+1} = \cos(pos / 10000^{2i/d_{model}})$$

depth/dimension index

Figure 12: Positional encoding applied in a transformer [13].

3.3 Masking – Encoder and Decoder Padding

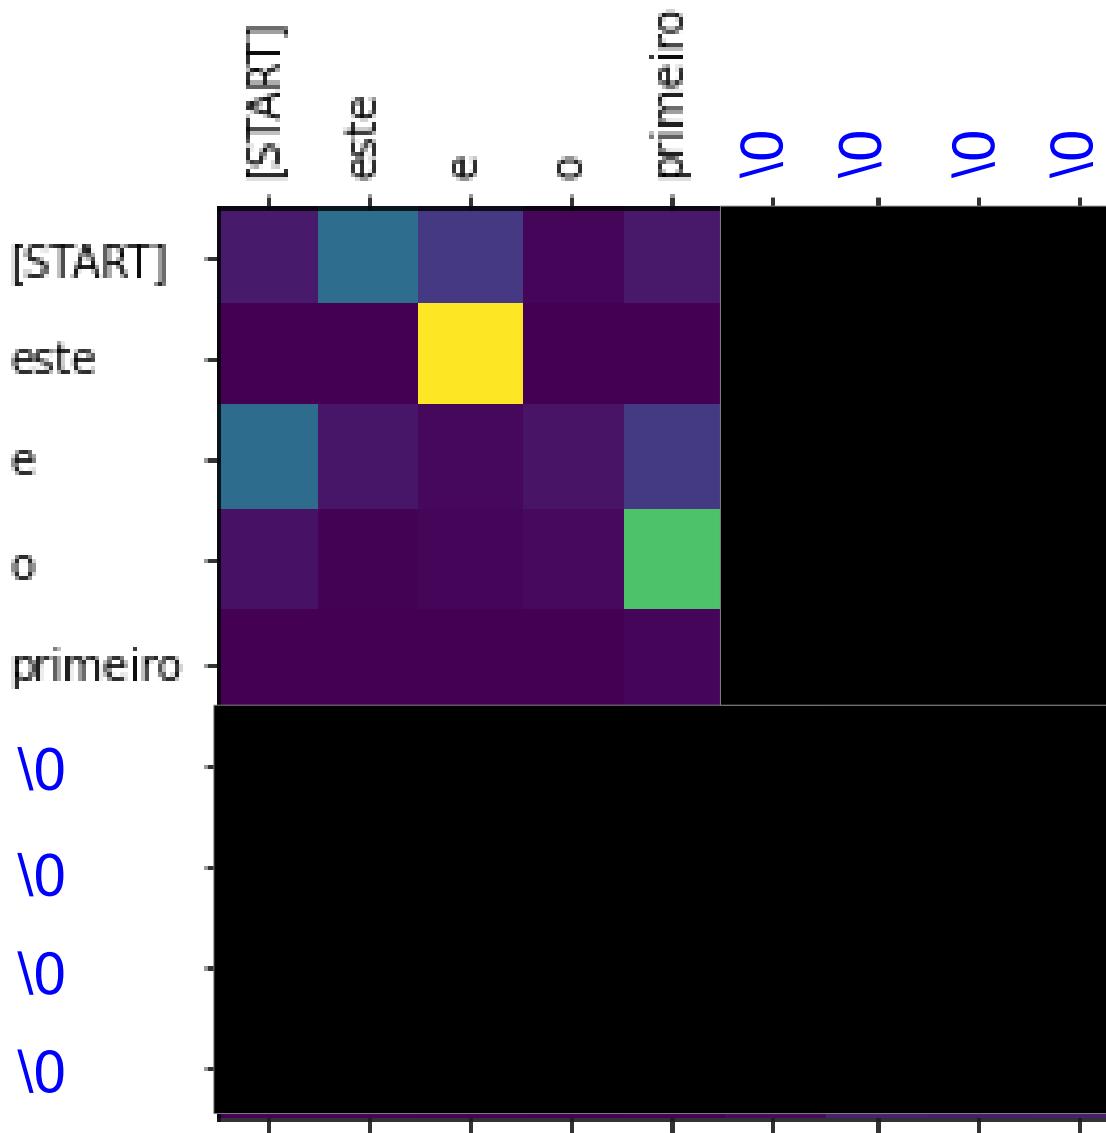


Figure 13: Masking padding words [11].

Image source:
https://www.tensorflow.org/static/text/tutorials/transformer_files/output_1kLCia68EIoE_1.png
(call date: 20.07.22)

3.3 Masking – Decoder: Left to right information flow

A word in the output sequence is not allowed to attend on words coming after it.




Figure 14: All values in the self-attention matrix above the diagonal are zero [11].

3.3 Masking – Decoder: Left to right information flow

A word in the output sequence is not allowed to attend on words coming after it.

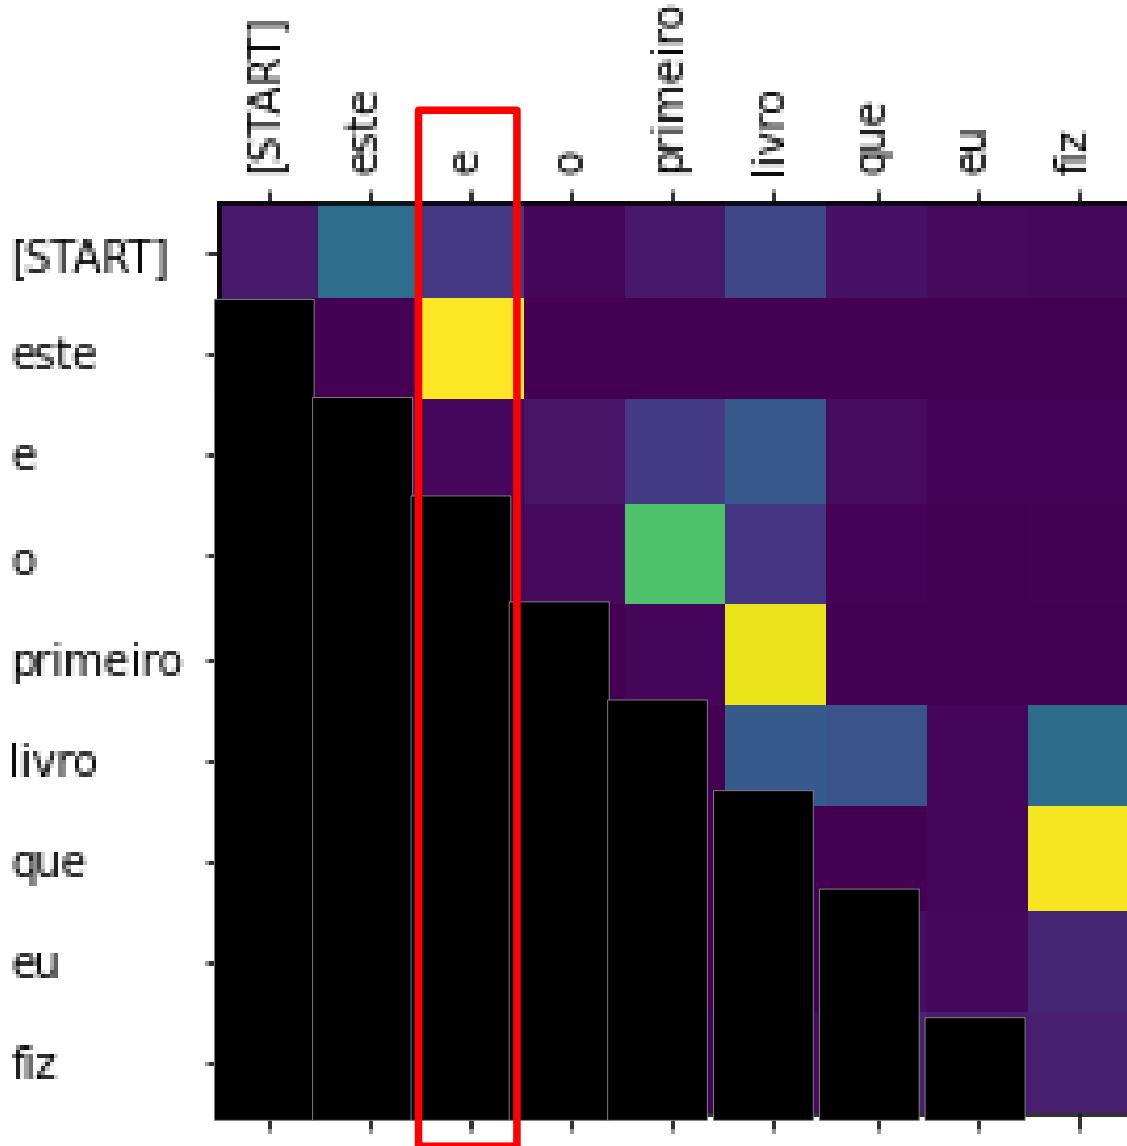


Figure 14: All values in the self-attention matrix above the diagonal are zero [11].

3.1 Let's build a transformer

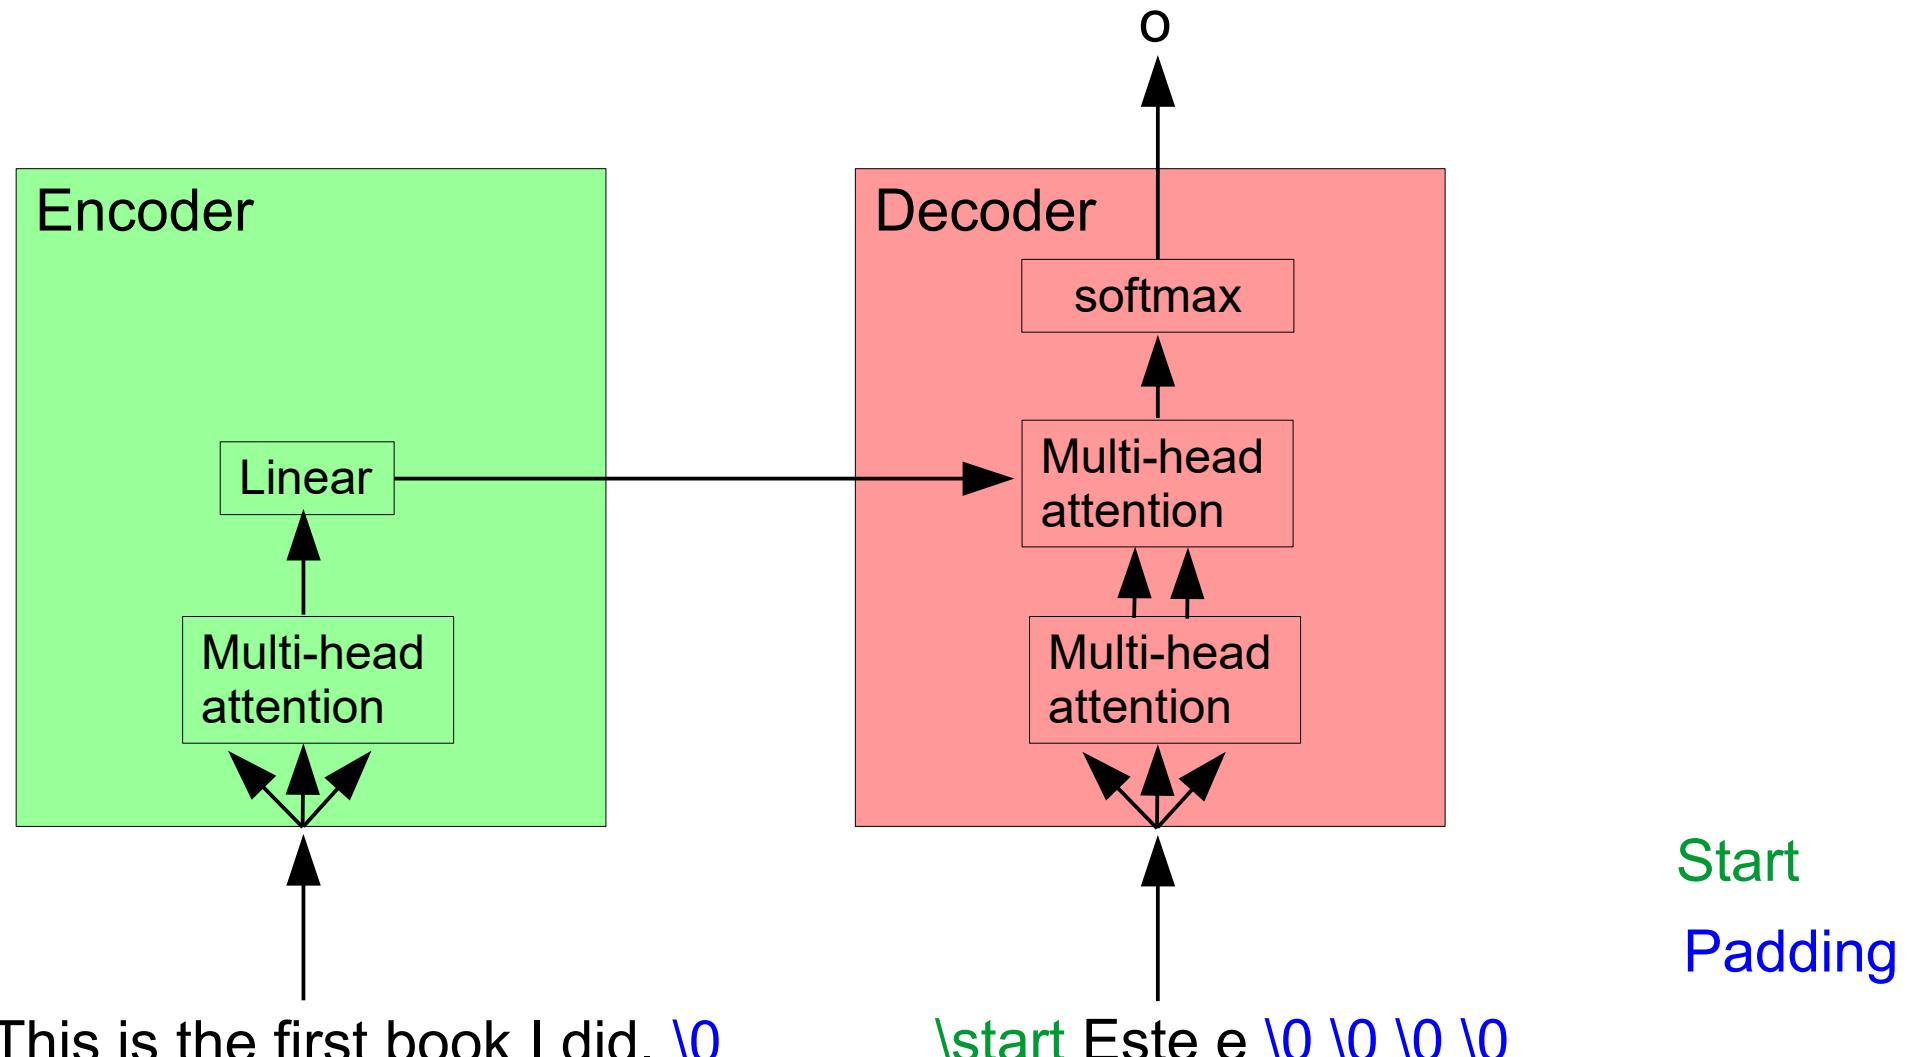


Figure 11: Neural machine translation based on a transformer [2].

3.4 Transformer

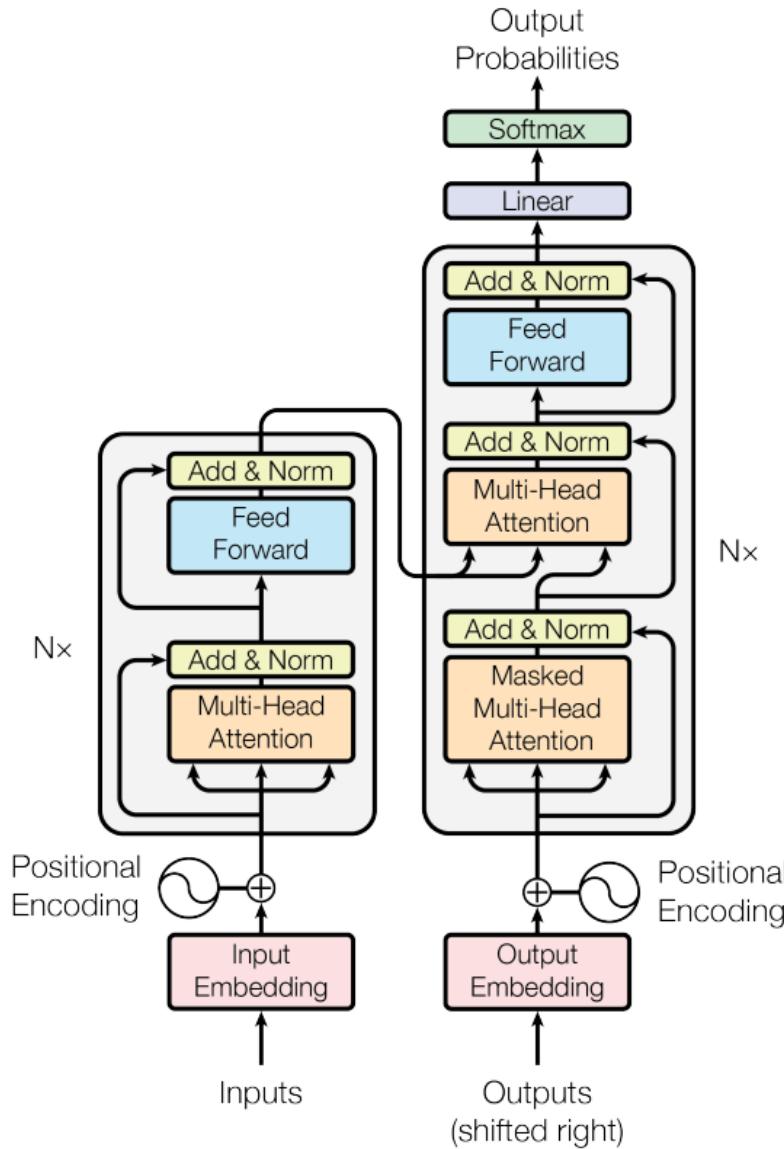


Figure 15: Transformer Architecture [2].

3.4 Transformer

Residual connection

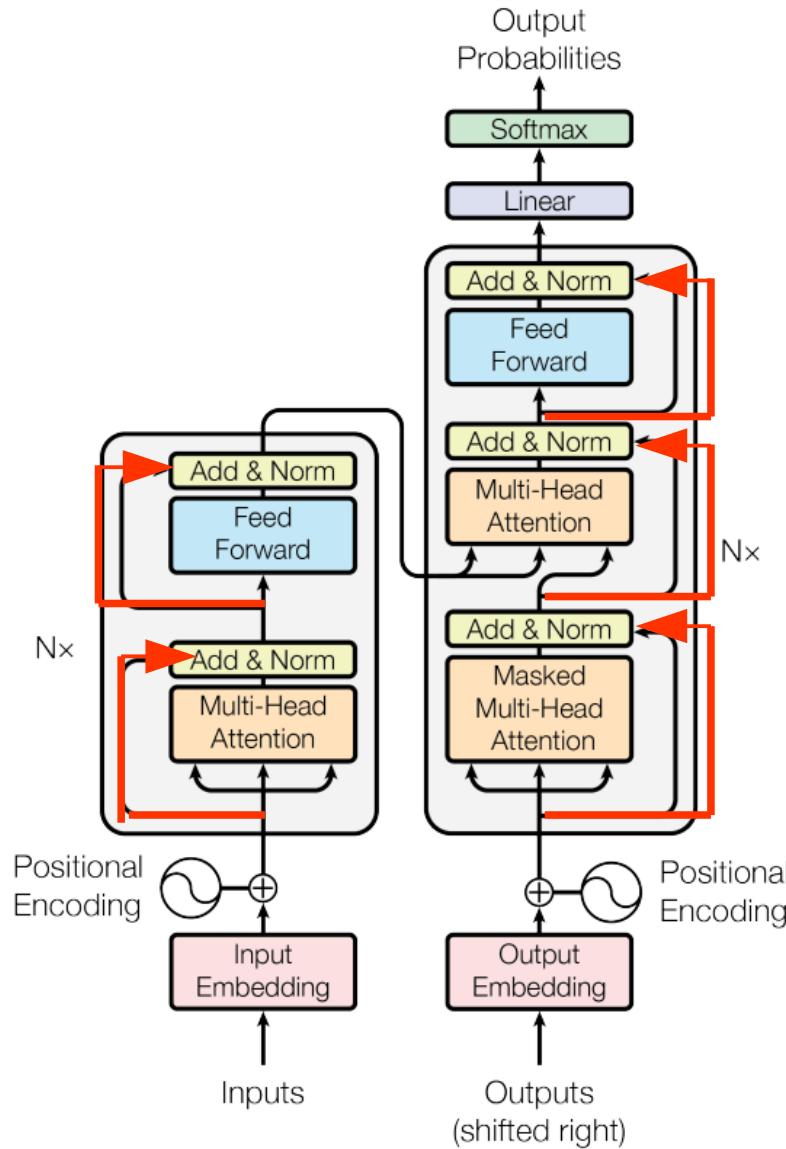


Figure 15: Transformer Architecture [2].

3.4 Transformer

Layer normalization

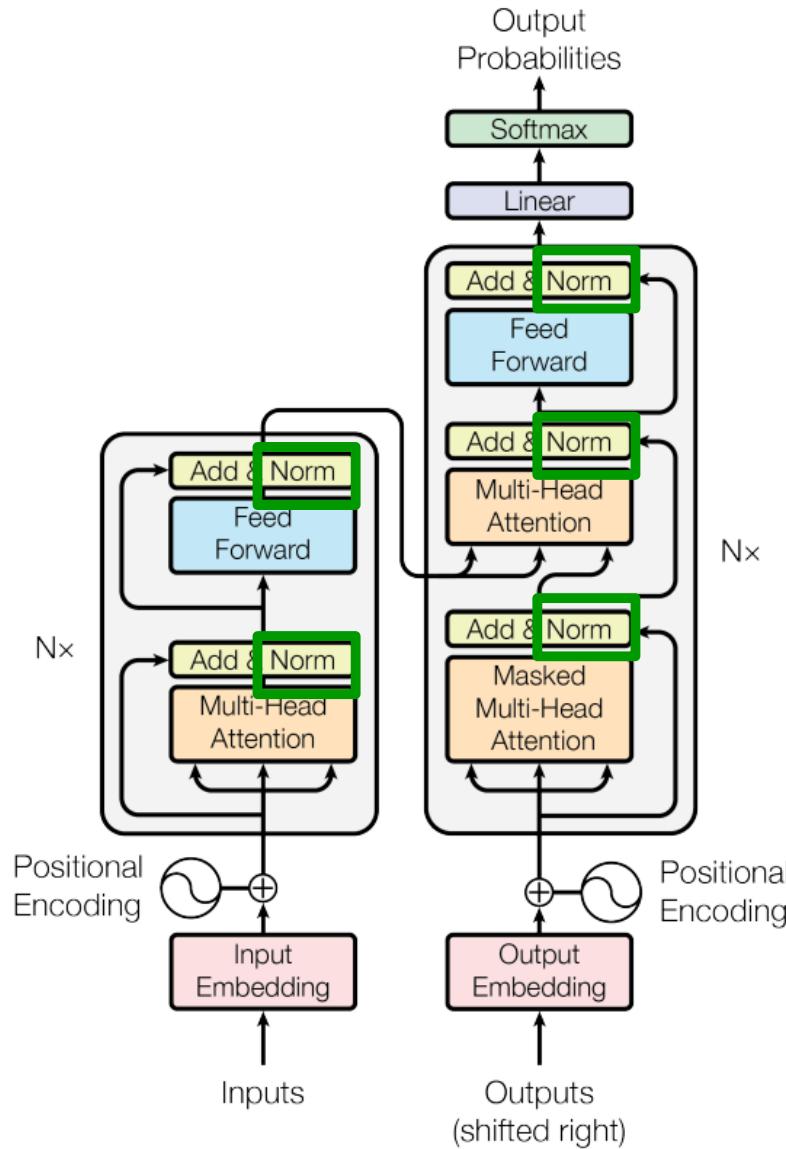


Figure 15: Transformer Architecture [2].

4. Transformers in Natural Language Processing

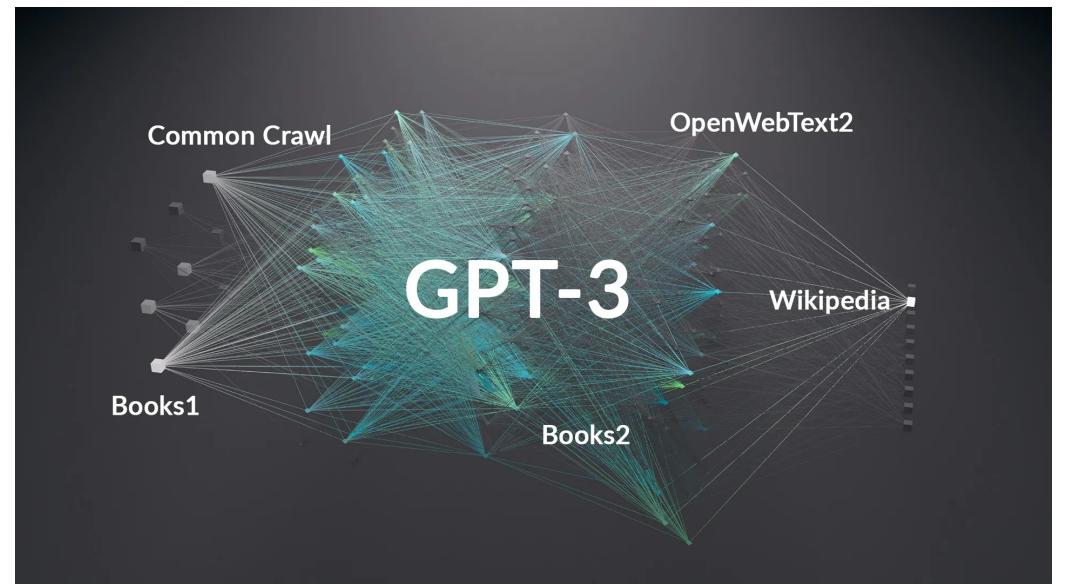


Image sources:
https://cdn-images-1.medium.com/max/428/1*IMlwRSQfzaYw1ckgS-jgZw.jpeg
<https://i0.wp.com/katzlberger.ai/wp-content/uploads/2021/04/GPT-3-Datenquellen.jpg?resize=1400%2C800&ssl=1>
(call dates: 20.07.22)

4.1 BERT: Bidirectional Encoder Representations from Transformers

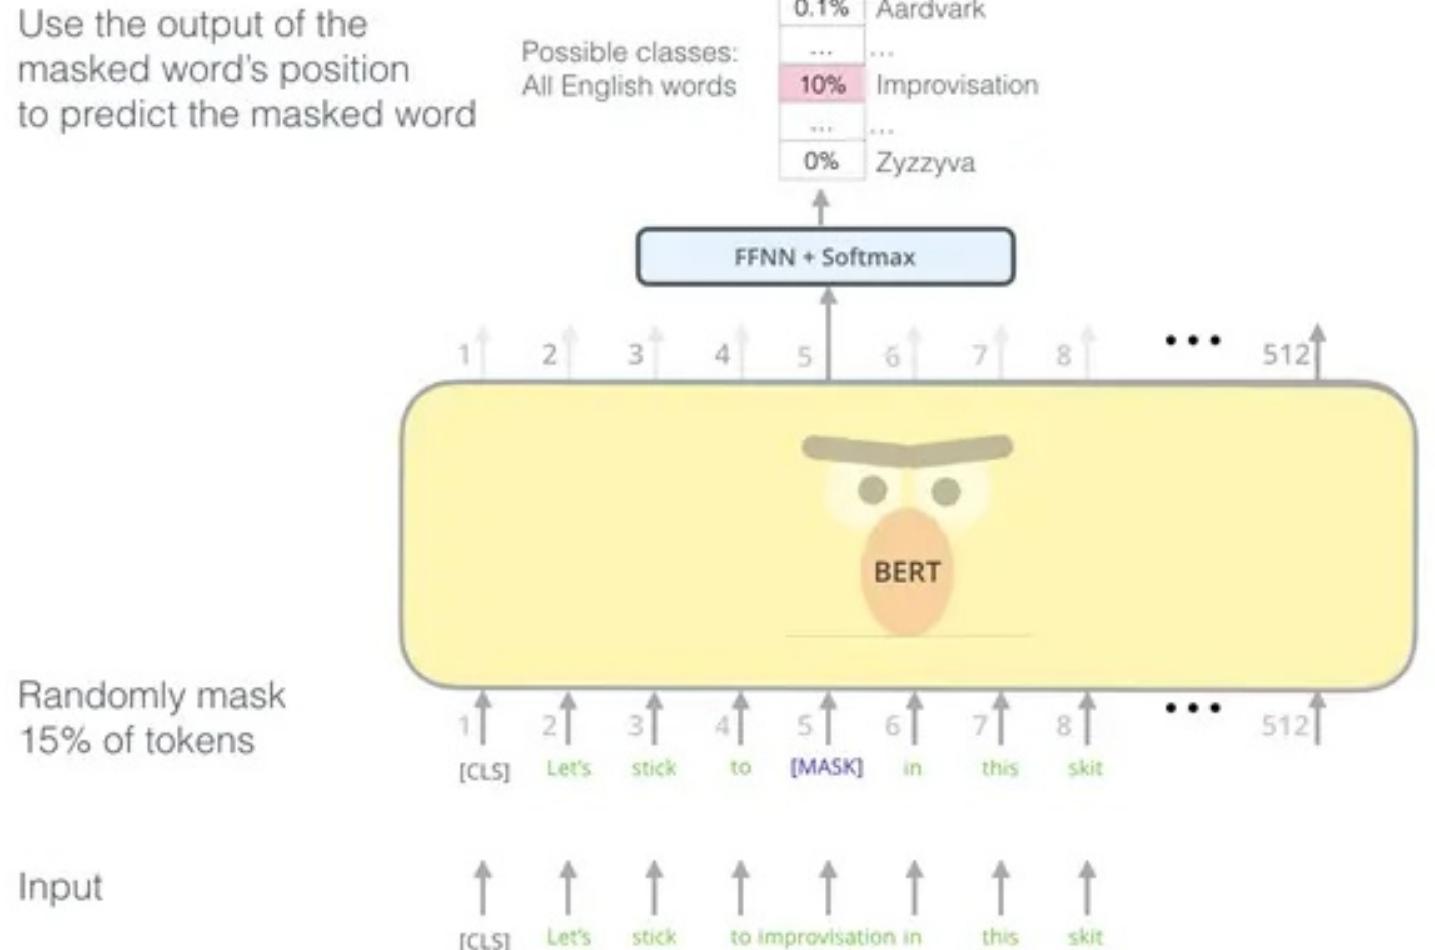


Figure 16: BERT Architecture [8].

Image source:
<https://www.alexanderthamm.com/wp-content/uploads/High-level-architecture-of-BERT-1.png.webp>
(call date: 20.07.22)

4.1 BERT: Bidirectional Encoder Representations from Transformers

Use the output of the masked word's position to predict the masked word

Possible classes:
All English words

0.1%	Aardvark
...	...
10%	Improvisation
...	...
0%	Zyzyva

Not autoregressive

Randomly mask 15% of tokens

Input




Figure 16: BERT Architecture [8].

4.1 BERT: Bidirectional Encoder Representations from Transformers

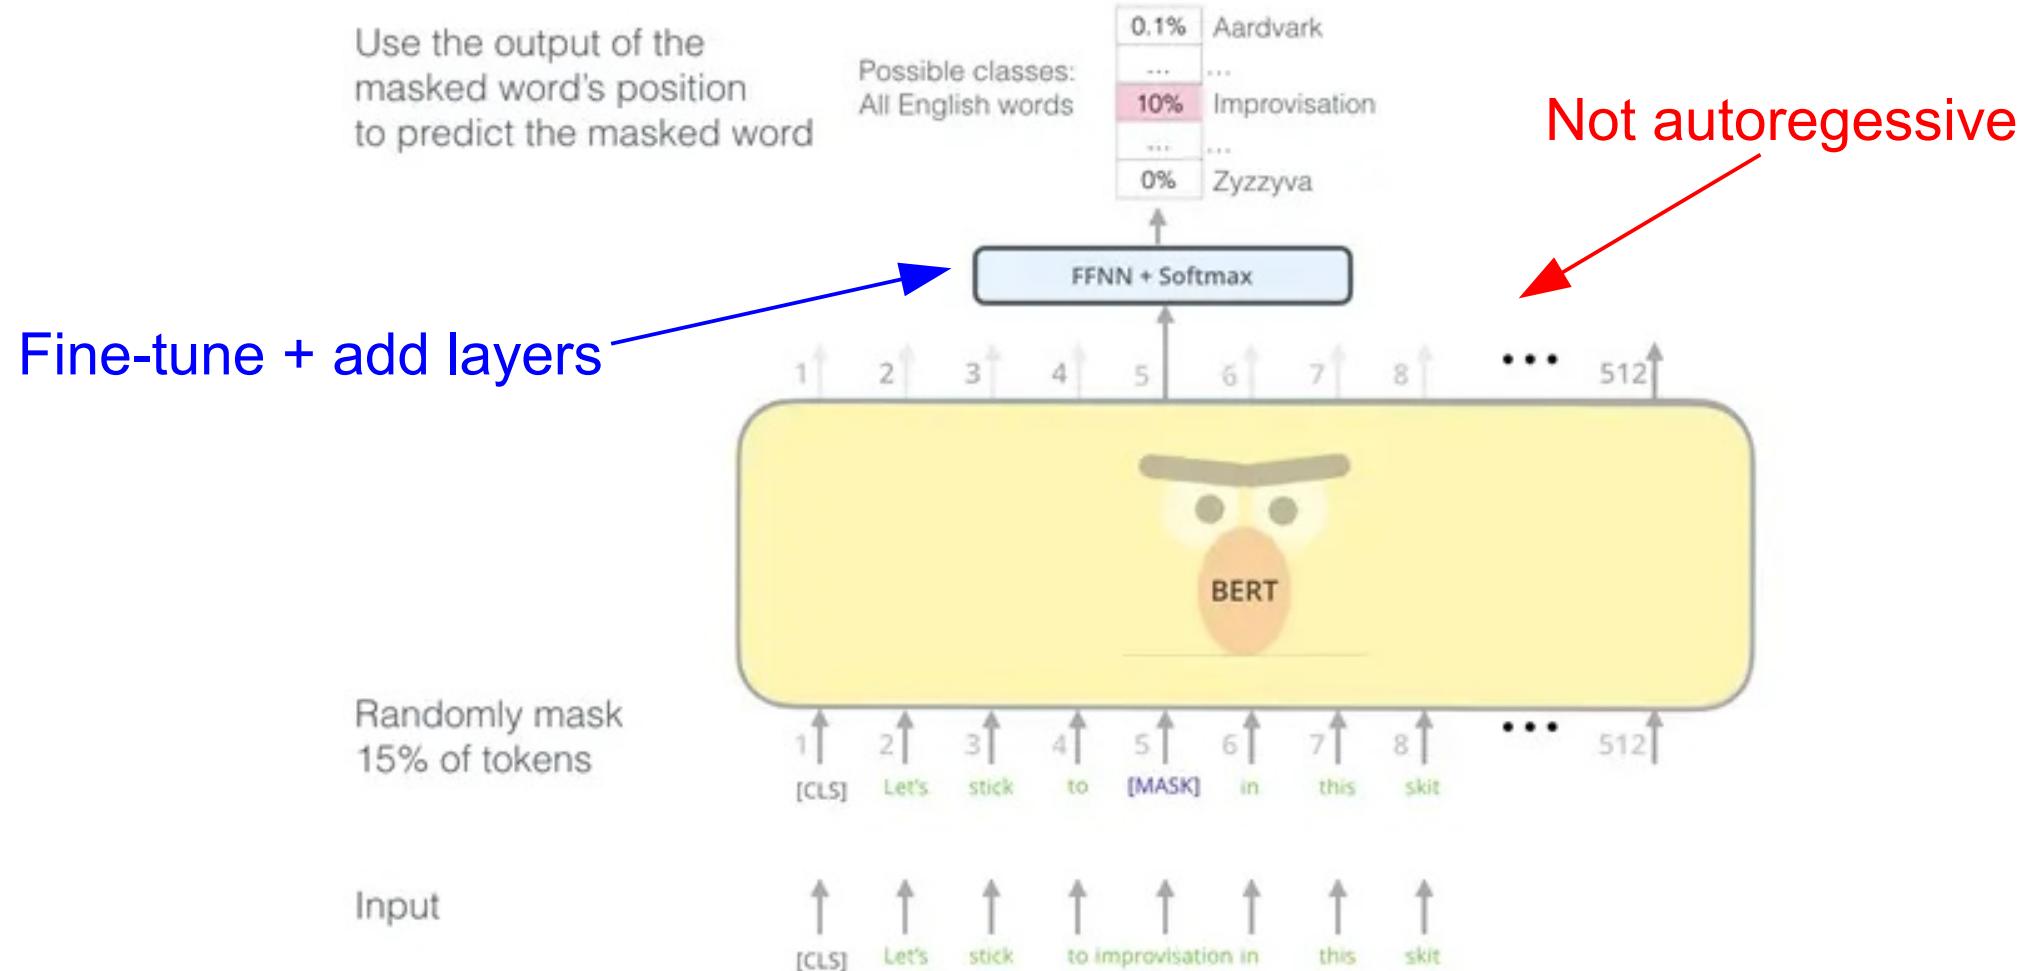


Figure 16: BERT Architecture [8].

Image source:
<https://www.alexanderthamm.com/wp-content/uploads/High-level-architecture-of-BERT-1.png.webp>
(call date: 20.07.22)

4.2 GPT-3: Generative Pre-trained Transformer 3

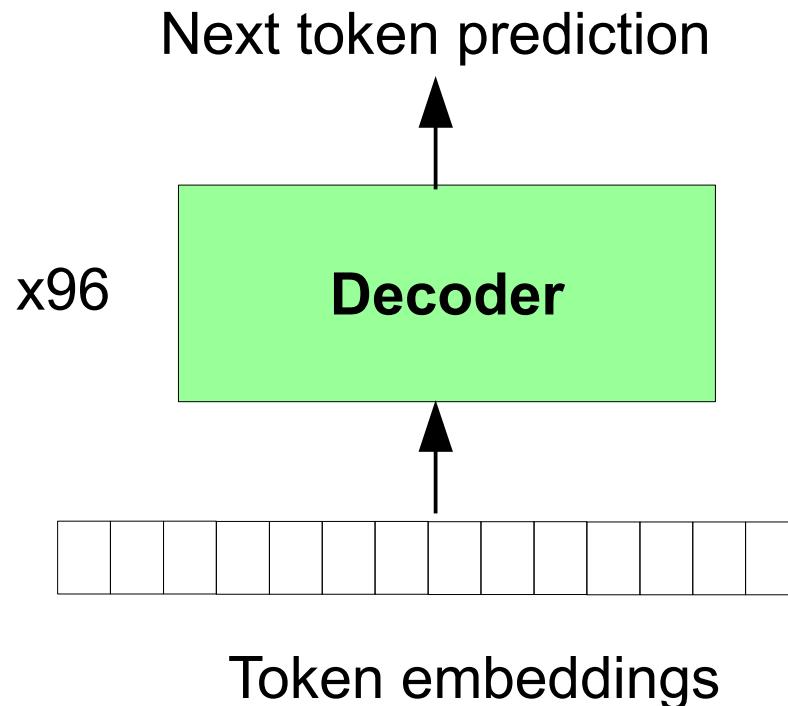


Figure 17: GPT-3 Overview [9].

4.2 GPT-3: Generative Pre-trained Transformer 3





Figure 17: GPT-3 Overview [9].

Image sources:
<https://upload.wikimedia.org/wikipedia/commons/thumb/8/80/Wikipedia-logo-v2.svg.png>
<https://cdn-icons-png.flaticon.com/512/29/29302.png>
<https://cdn-icons-png.flaticon.com/512/493/493805.png>
(call dates: 20.07.22)

4.2 GPT-3: Generative Pre-trained Transformer 3

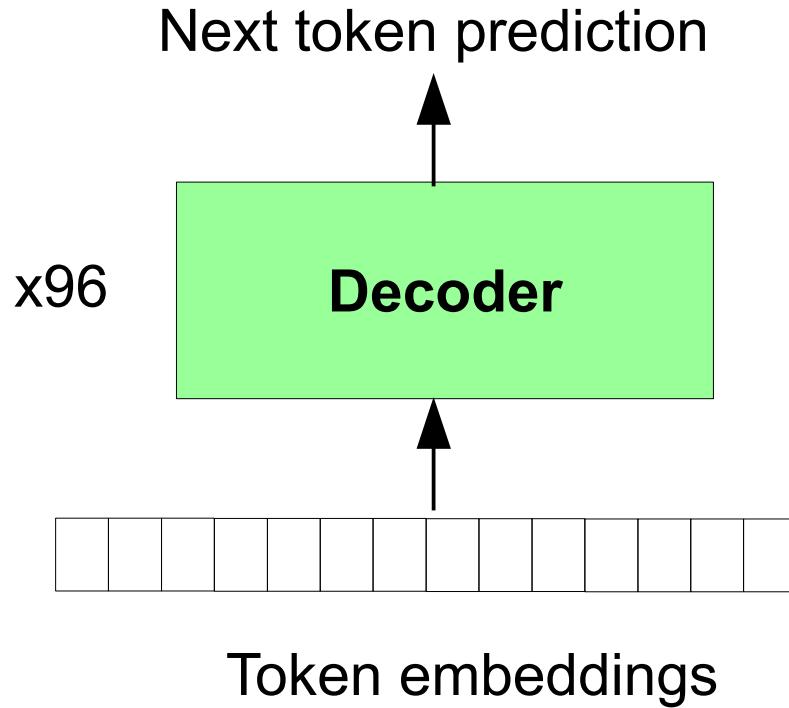


Figure 17: GPT-3 Overview [9].

Fine-tuning: **No gradient updates**

Input:

*Translate English to French:
cheese =>*

Image sources:
<https://upload.wikimedia.org/wikipedia/commons/thumb/8/80/Wikipedia-logo-v2.svg.png>
<https://cdn-icons-png.flaticon.com/512/29/29302.png>
<https://cdn-icons-png.flaticon.com/512/493/493805.png>
(call dates: 20.07.22)

4.3 BERT vs. GPT-3

	BERT	GPT-3
release year	2018	2020
#params	340M	175B
#attention heads	16	96
hidden size	1024	12288
stacked transformer blocks	encoder (non-autoregressive) bidirectional	decoder (autoregressive) left-to-right
attention direction		
#stacked transformer blocks	12	96
input	512 token embeddings	2048 token embeddings
output	512 token embeddings + classifiers	1 token embedding + classifier

Table 1: Differences between BERT and GPT-3 [8,9].


5. Transformers in Computer Vision

Classification



CAT

Object Detection



CAT, DOG, DUCK

5.1.1 ViT: Vision Transformer – Image Classification

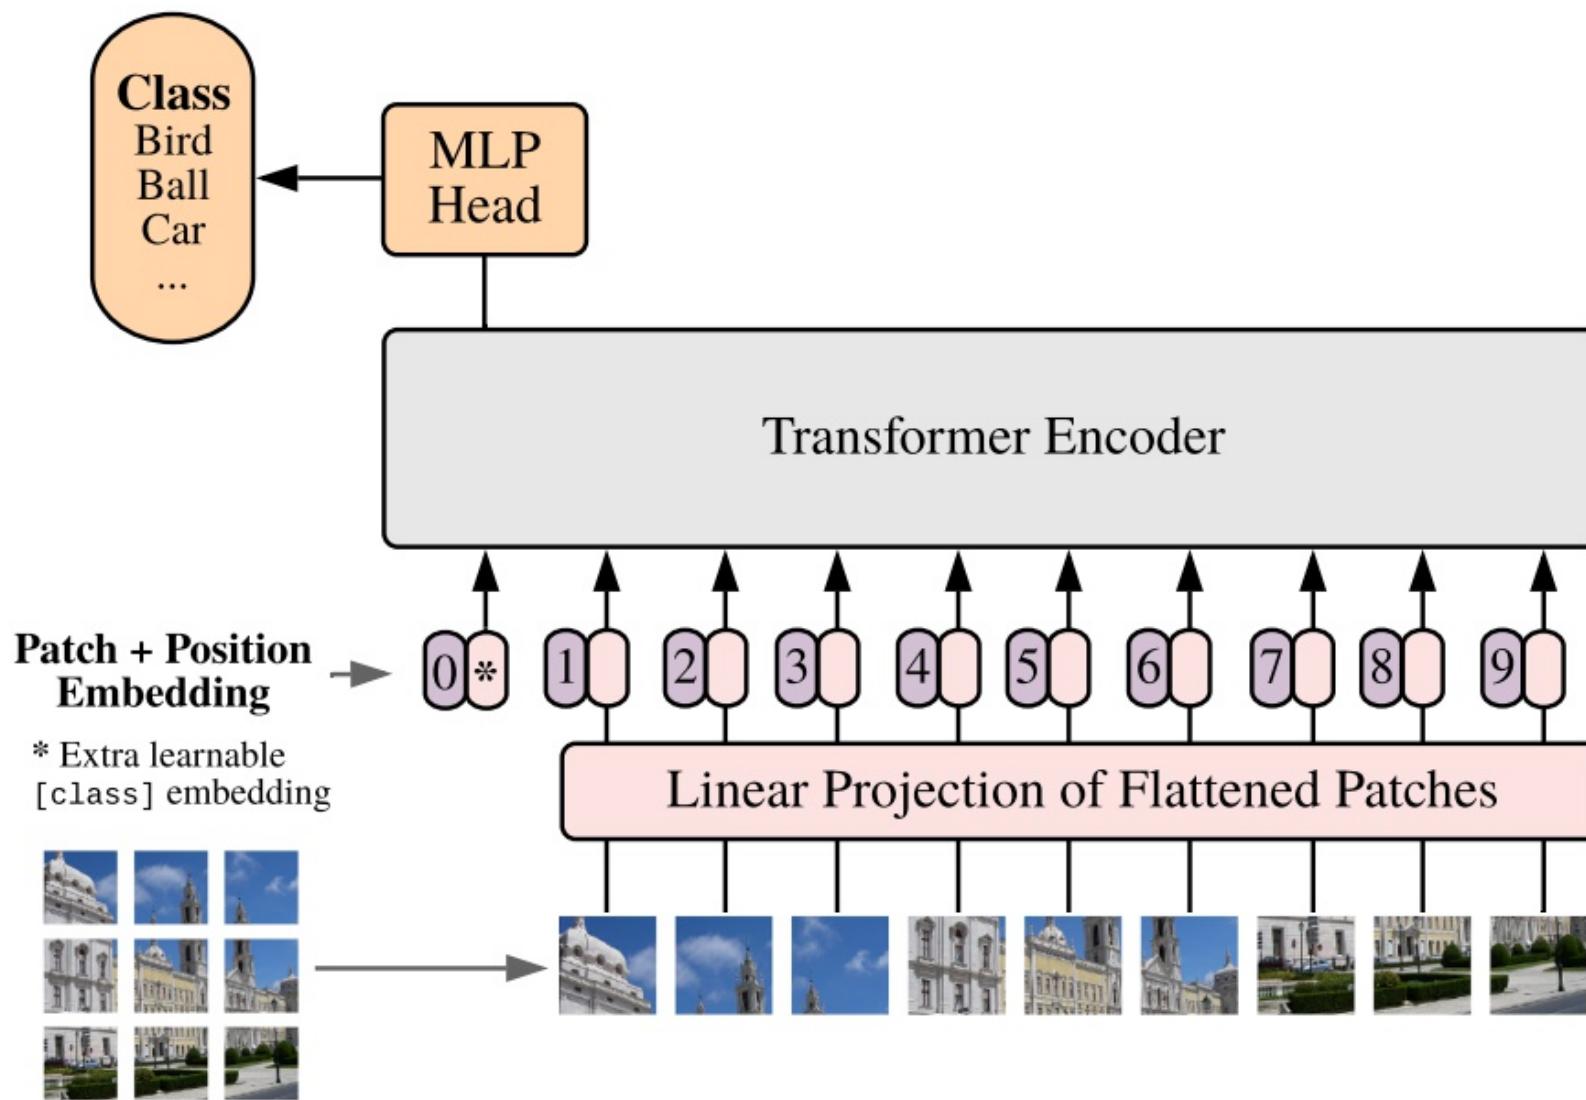


Figure 18: ViT Architecture [3].

5.1.2 ViT: Vision Transformer – Evaluation

	ViT-H/14	ResNet152x4
ImageNet	88.55	87.54
ImageNet ReaL	90.72	90.54
CIFAR-10	99.50	99.37
CIFAR-100	94.55	93.51
Oxford-IIIT Pets	97.56	96.62
Oxford Flowers-102	99.68	99.63
VTAB (19 tasks)	77.63	76.29
#param	632M	14M
TPUv3-core-days	2.5k	9.9k

Table 2: Performance of the ViT and ResNet [3].

5.1.3 ViT: Vision Transformer – “Attention in an image”

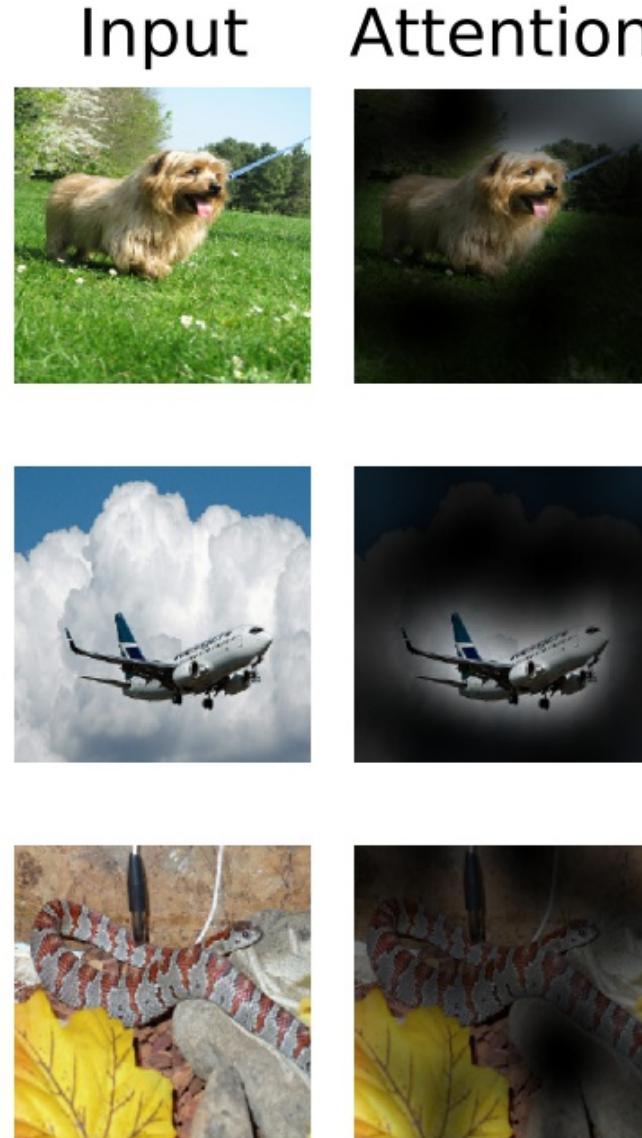
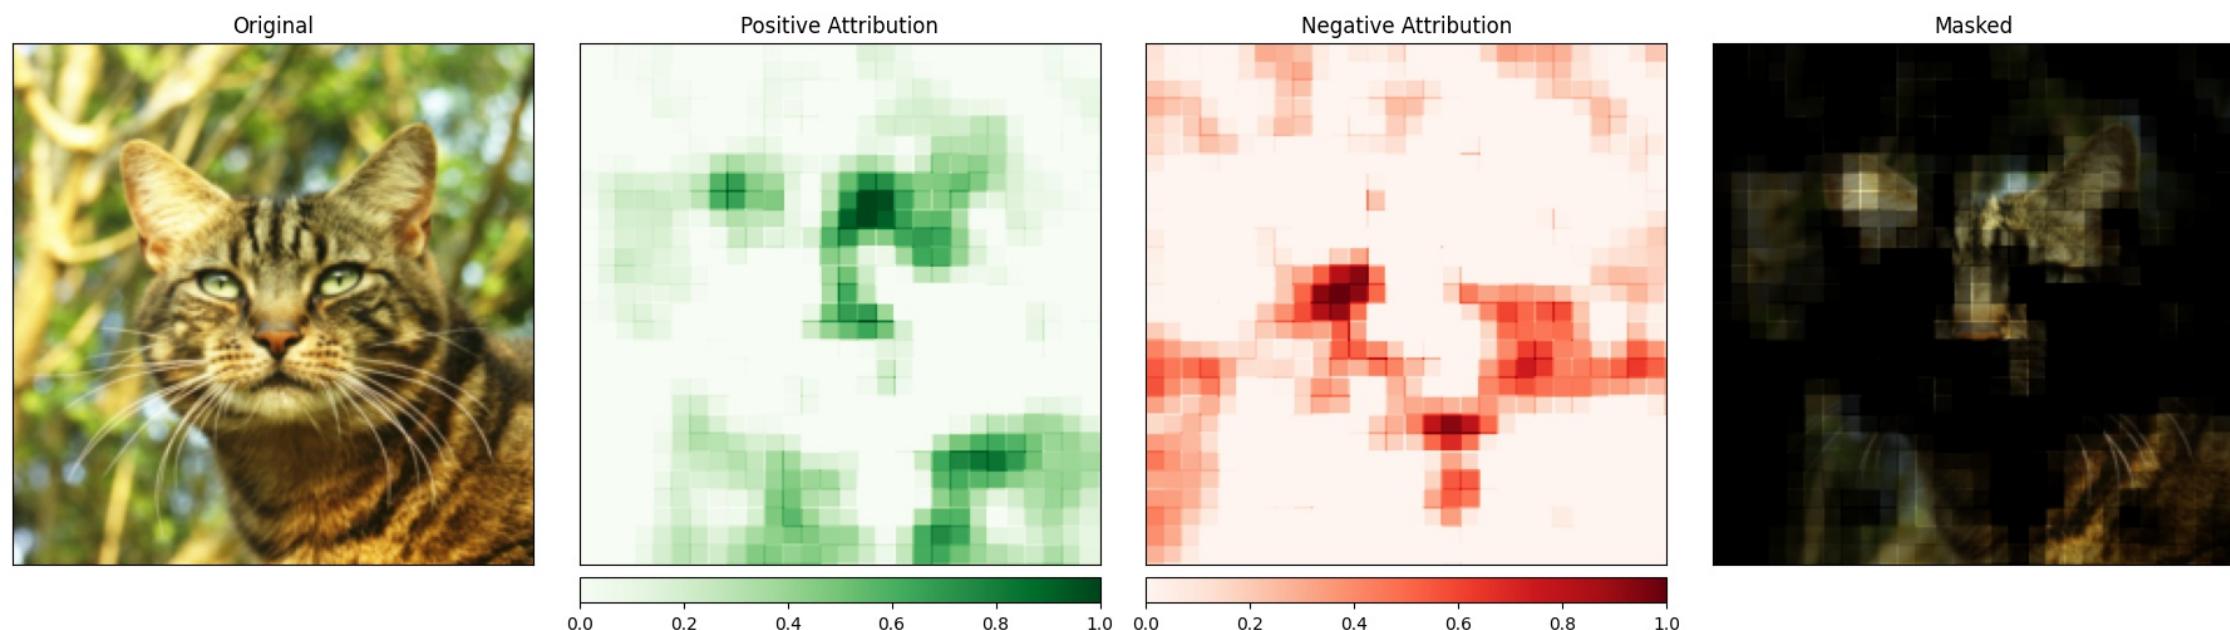


Figure 19: Visualization of self-attention
in an image [3].

Side note: Feature Attribution with Occlusion



5.1.4 ViT: Vision Transformer – “Learn to be a CNN”

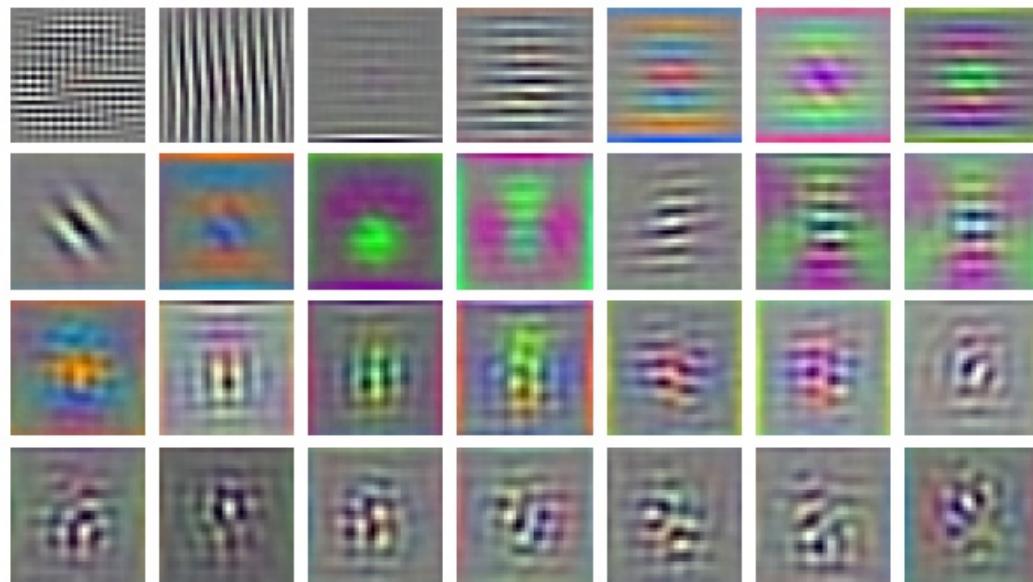


Figure 20: Learned patterns to which the self-attention mechanism reacts to [3].




Figure 21: Mean attention distance of each layer [3].

Side note: Receptive Field in a CNN




Image source:
<https://theaisummer.com/static/490be17ee7f19b78003c3fdf5a6bbafc/83b75/receptive-field-in-convolutional-networks.png> (call date: 20.07.22)

5.2.1 DETR: DEtection Transformer – Object detection

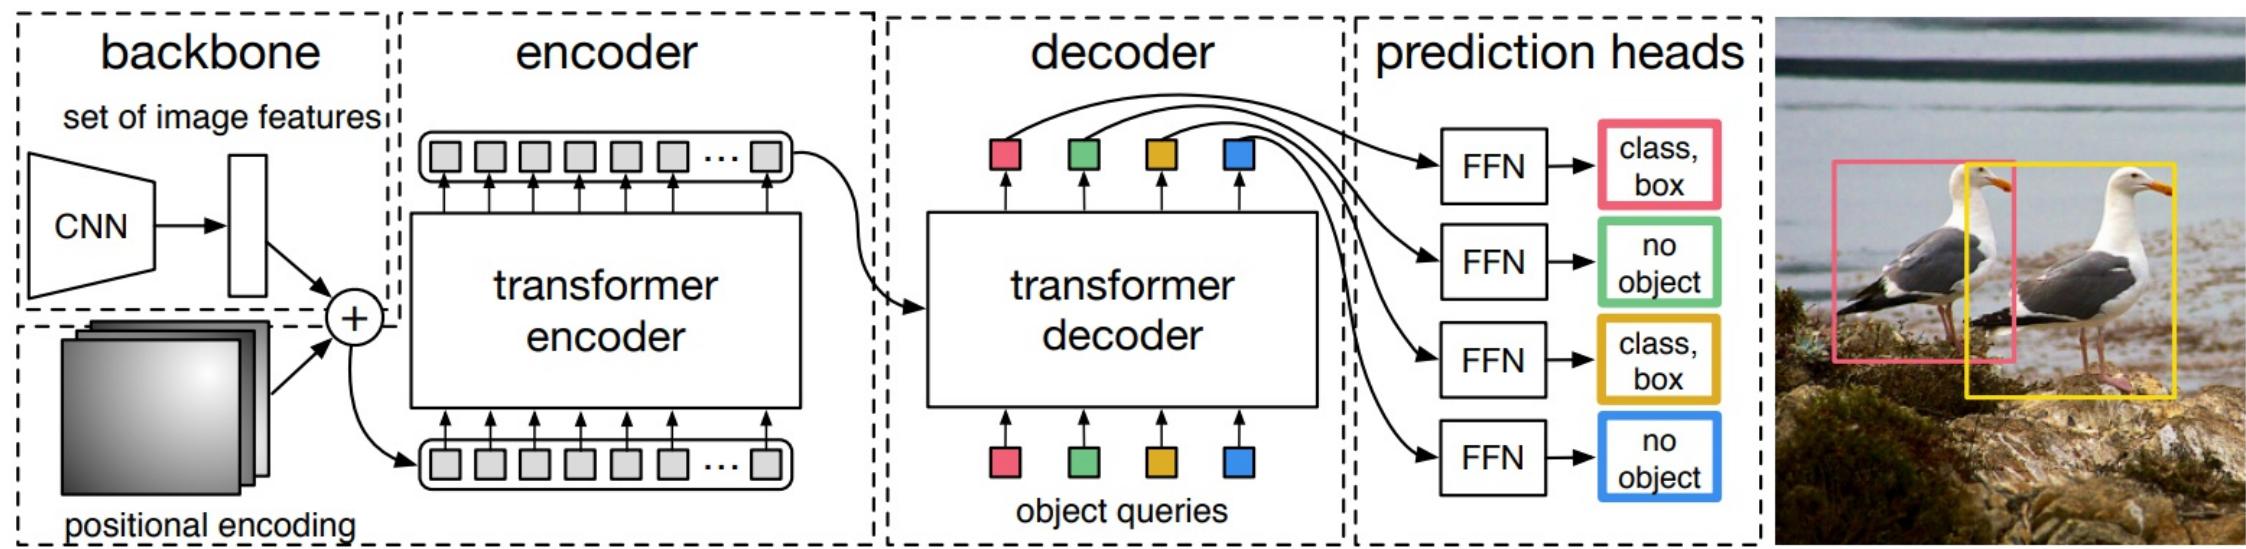


Figure 22: DETR architecture [4].

5.2.1 DETR: DEtection Transformer – Object detection

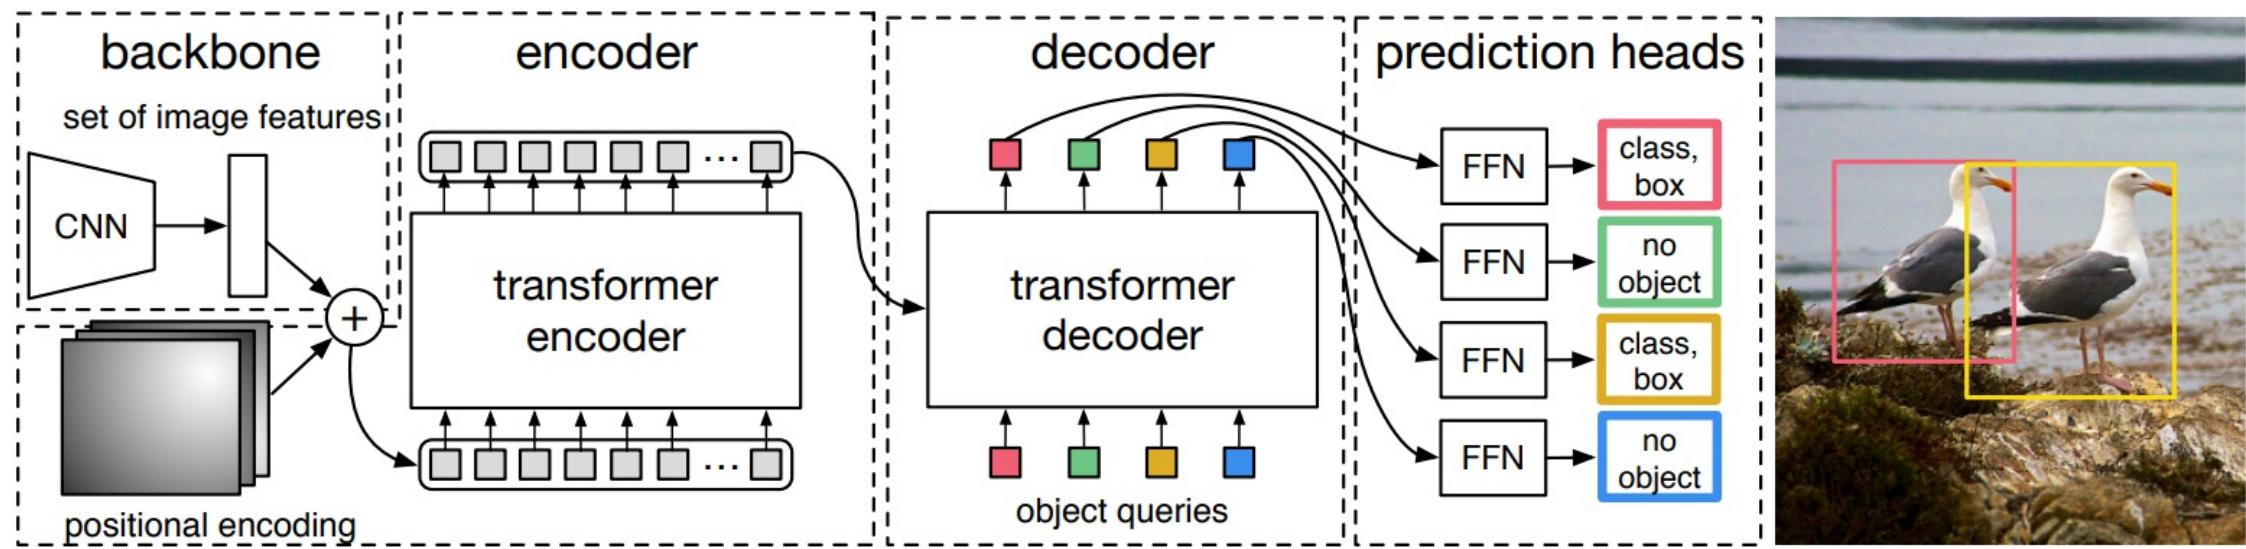


Figure 22: DETR architecture [4].




Figure 23: Object queries [4].

5.2.1 DETR: DEtection Transformer – Object detection

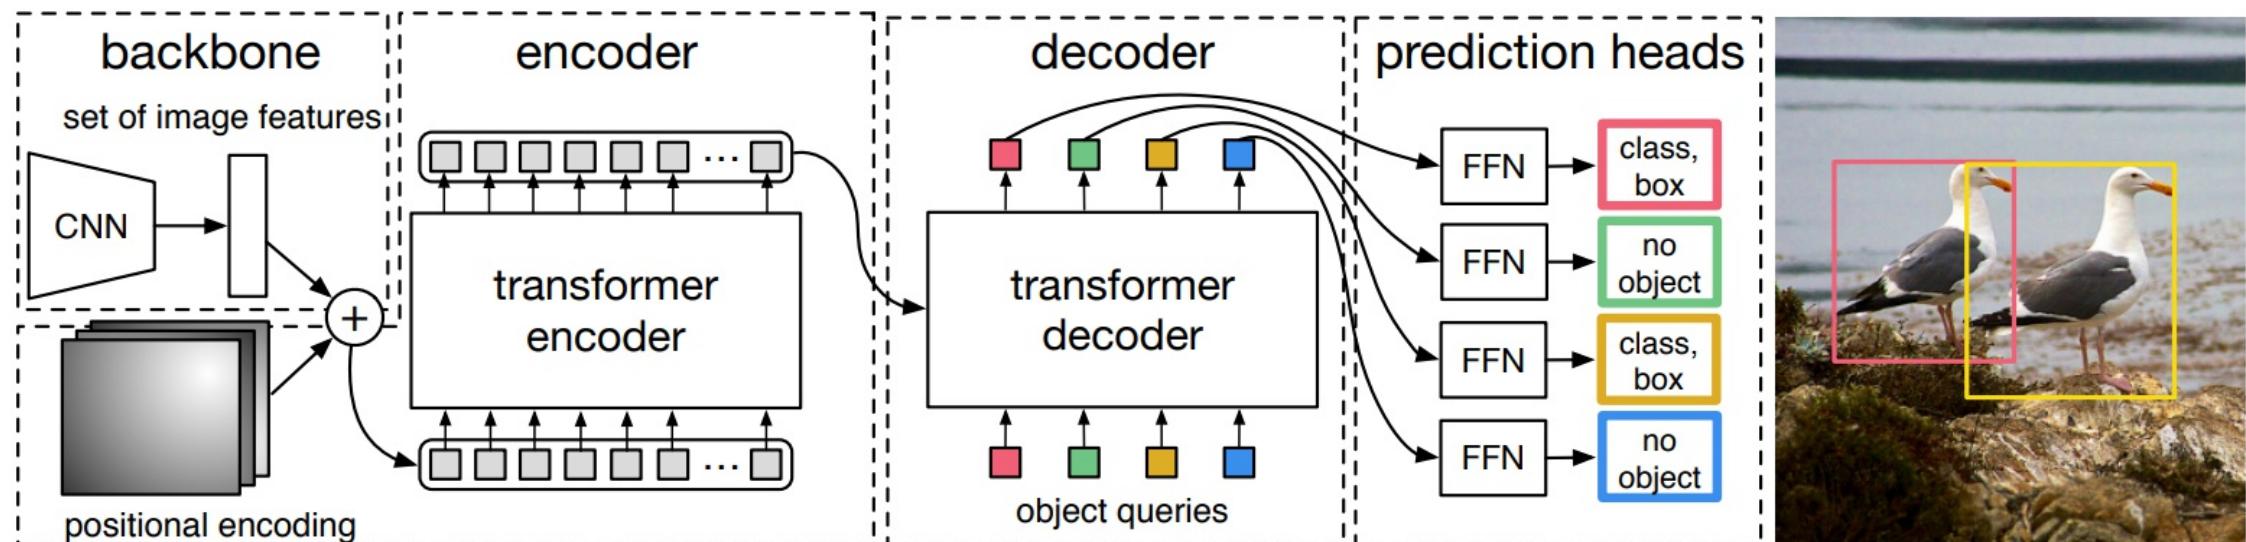


Figure 22: DETR architecture [4].

Training:

Loss = best matching loss between target and prediction → Hungarian Algorithm

$$\hat{\sigma} = \arg \min_{\sigma \in \mathfrak{S}_N} \sum_i^N \mathcal{L}_{\text{match}}(y_i, \hat{y}_{\sigma(i)})$$

N = amount of boxes

5.2.1 DETR: DEtection Transformer – Object detection

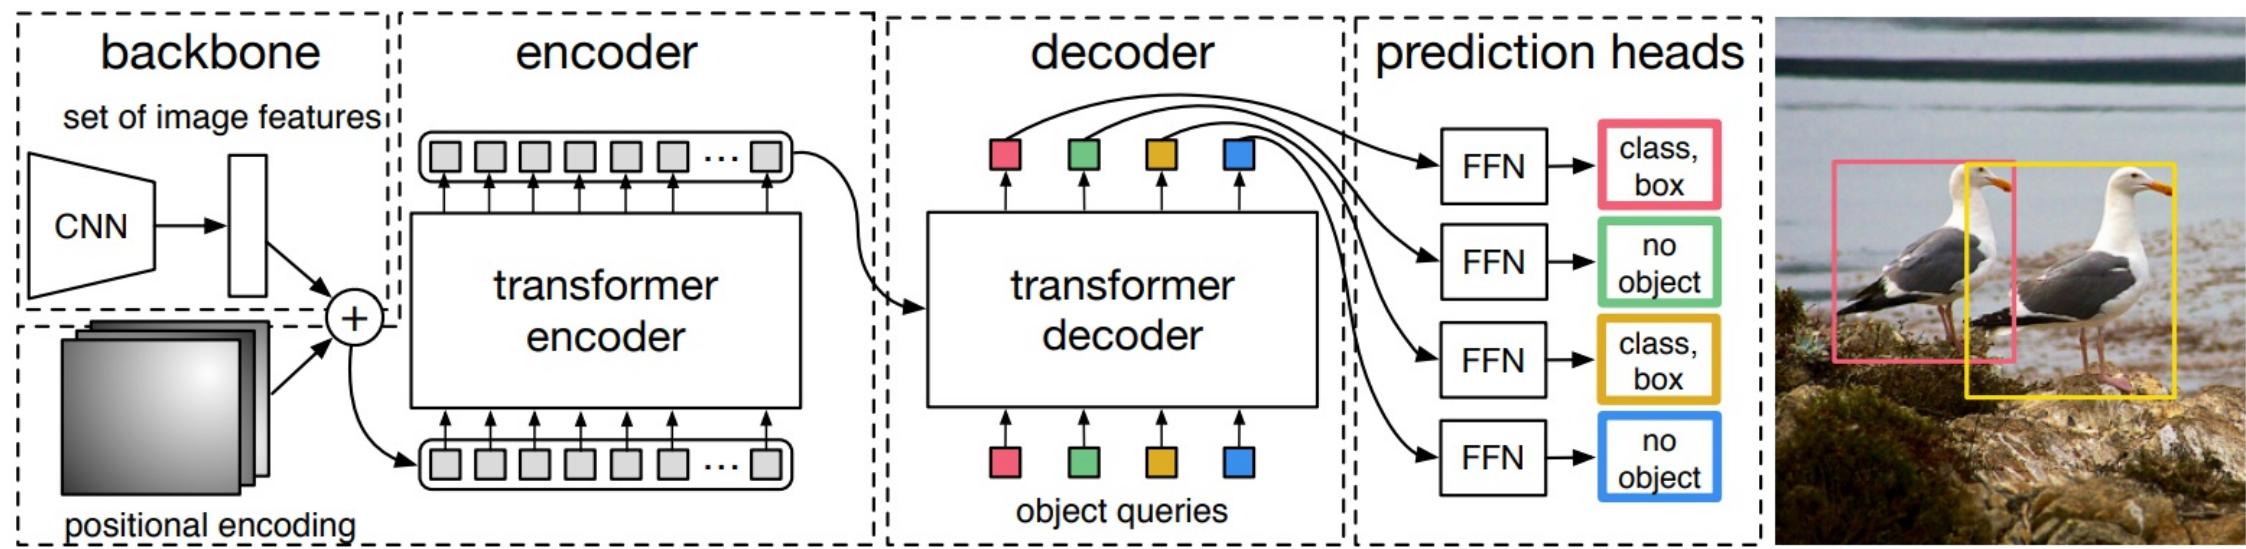


Figure 22: DETR architecture [4].

Training:

Loss = best matching loss between target and prediction → Hungarian Algorithm

$$\hat{\sigma} = \arg \min_{\sigma \in \mathfrak{S}_N} \sum_i^N \mathcal{L}_{\text{match}}(y_i, \hat{y}_{\sigma(i)})$$

N = amount of boxes

$$\mathcal{L}_{\text{Hungarian}}(y, \hat{y}) = \sum_{i=1}^N \left[-\log \hat{p}_{\hat{\sigma}(i)}(c_i) + \mathbb{1}_{\{c_i \neq \emptyset\}} \mathcal{L}_{\text{box}}(b_i, \hat{b}_{\hat{\sigma}(i)}) \right]$$

Classification loss
(cross entropy loss)

Box loss
(intersection over union loss)

5.2.2 DETR: DEtection Transformer “Attention-head communication”

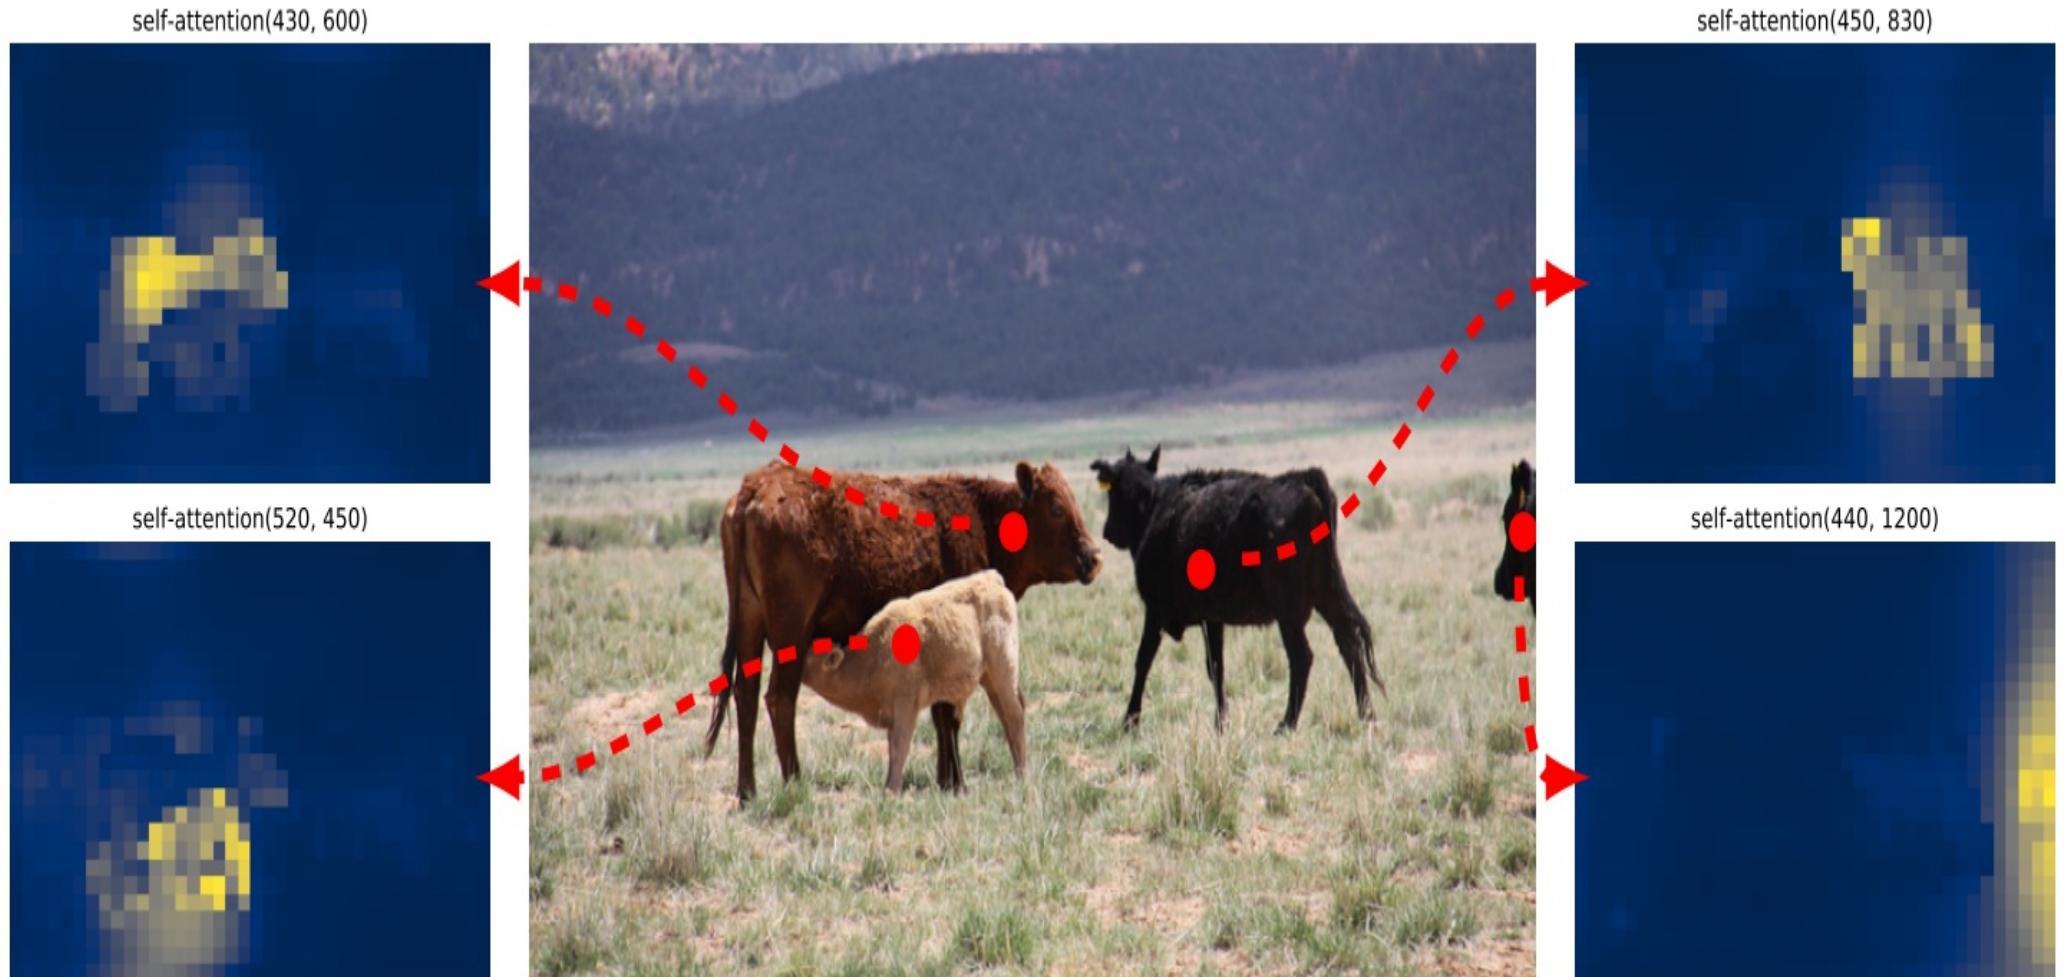


Figure 24: Different attention-heads attend to different objects in the image [4].



This Person Does Not Exist [25].

6. GAN based on a Transformer

Without transposed
convolutions!!!

6.1 TransGAN: Architecture

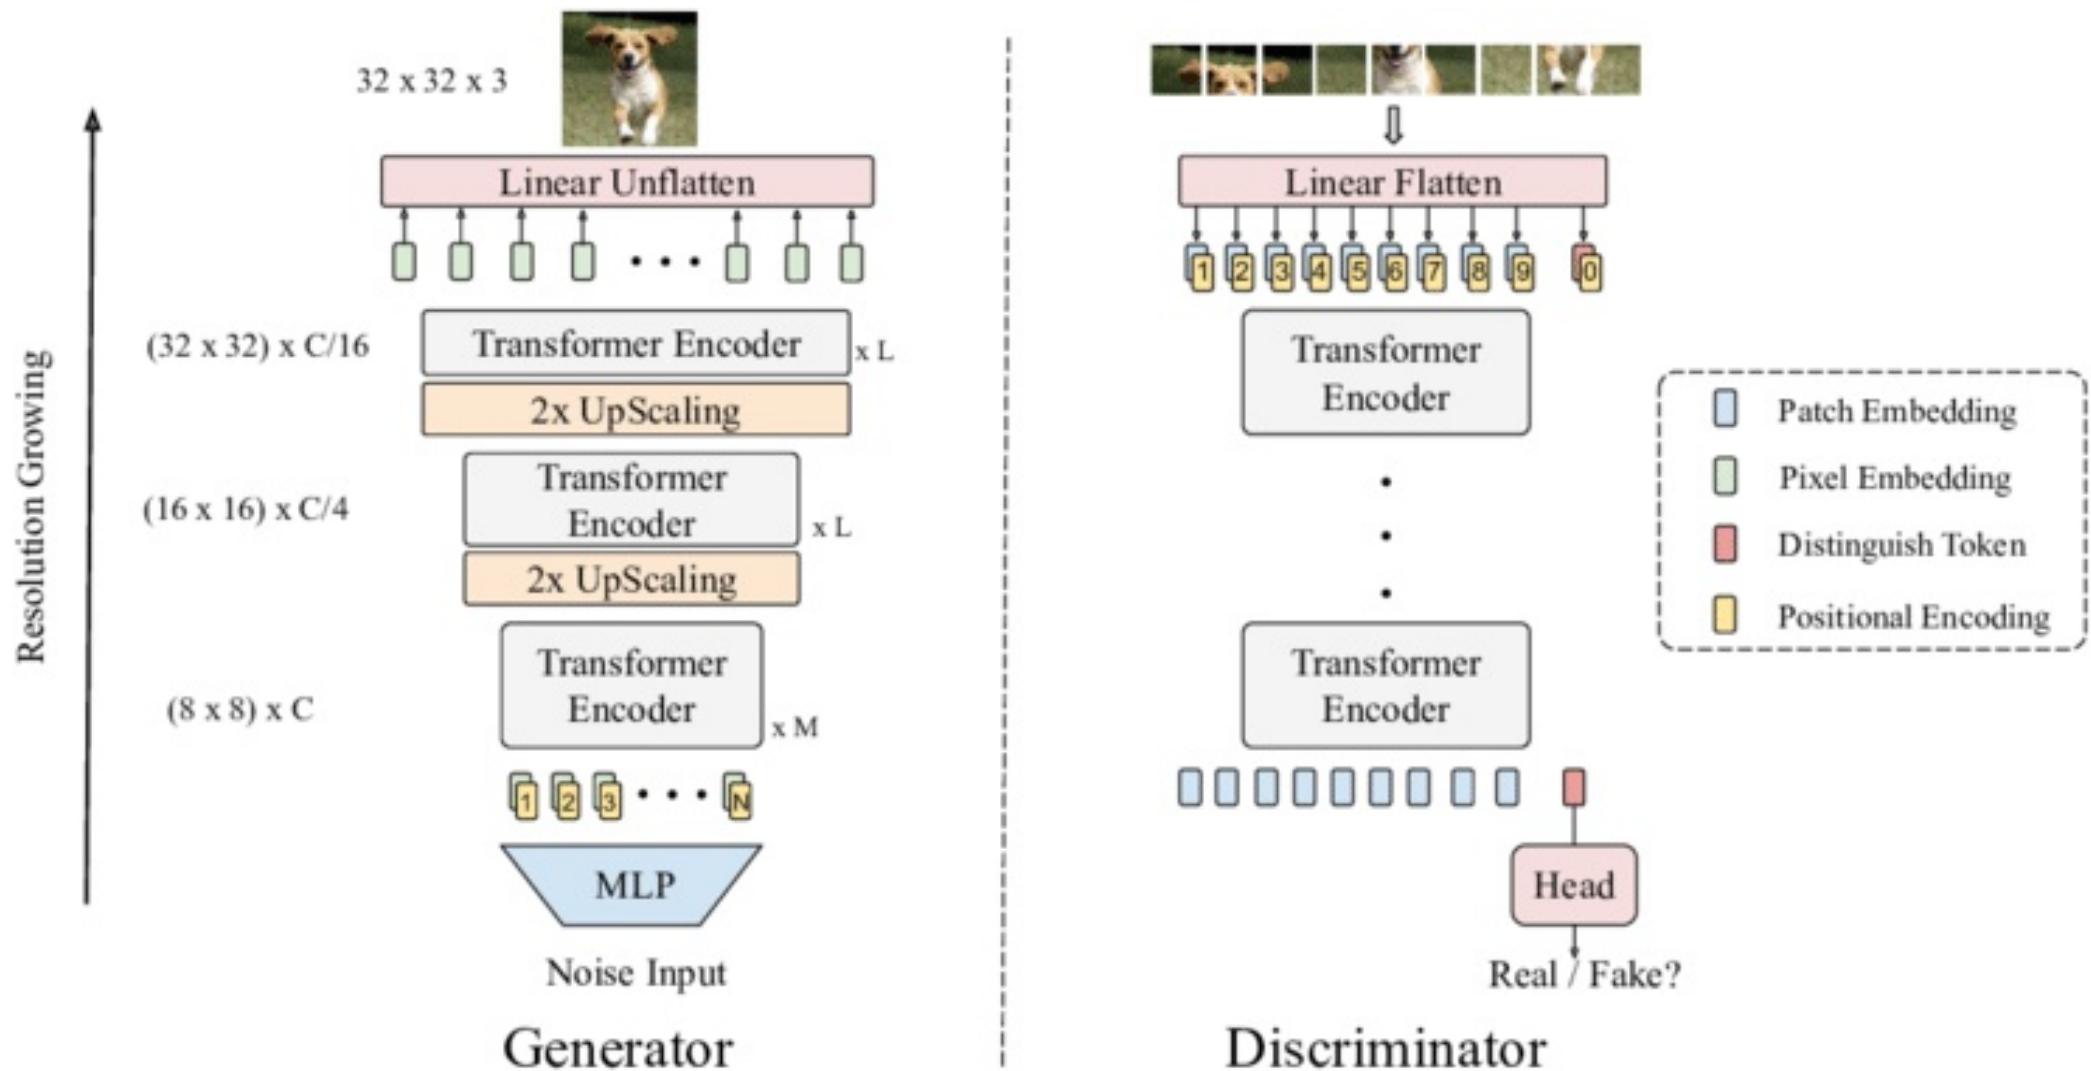


Figure 25: TransGAN architecture [5].

Side note: PixelShuffle




Image source:

Shi, Wenzhe et al. "Real-Time Single Image and Video Super-Resolution Using an Efficient Sub-Pixel Convolutional Neural Network." 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (2016): 1874-1883.

6.2 TransGAN: Evaluation



Figure 26: Generated images by the TransGAN [5].

6.2 TransGAN: Evaluation

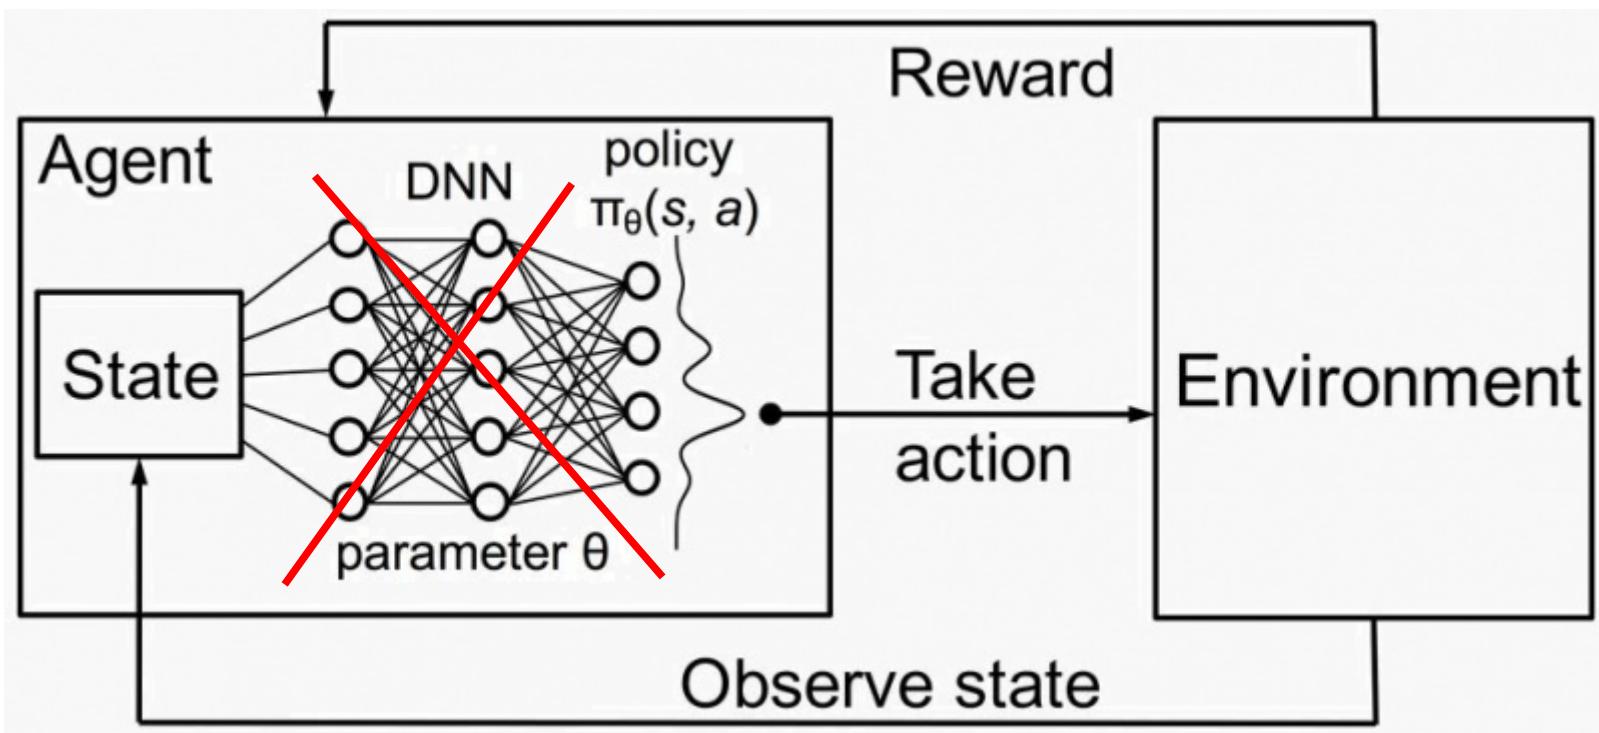
Methods	WGAN-GP		AutoGAN		StyleGAN-V2		TransGAN	
	IS ↑	FID ↓	IS ↑	FID ↓	IS ↑	FID ↓	IS ↑	FID ↓
Original + DiffAug [69]	6.49	39.68	8.55	12.42	9.18	11.07	8.36	22.53
	6.29	37.14	8.60	12.72	9.40	9.89	9.02	9.26

Table 3: Performance of the TransGAN compared with state-of-the-art GANs [5].

IS: Inception score “high = good”

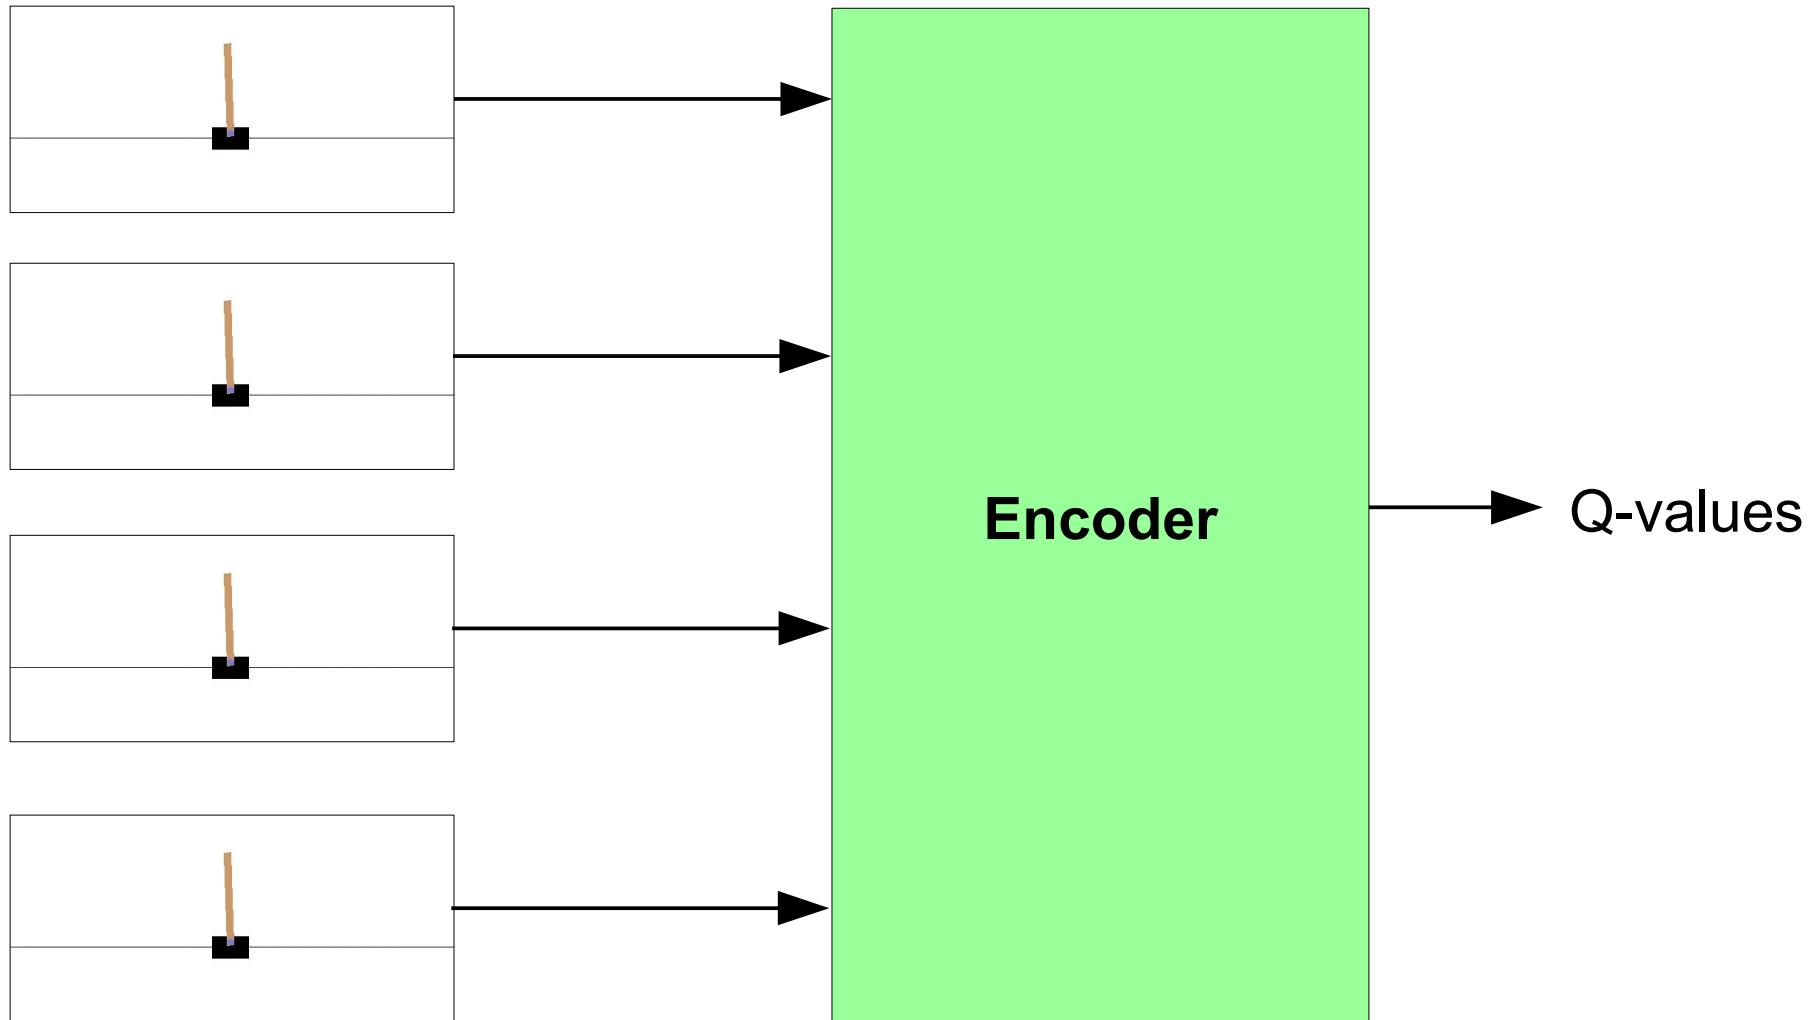
FID: Frechet Inception Distance “low = good”

7. Transformers in Deep Reinforcement Learning



Now: Transformer

7.1.1 DTQN: Deep Transformer Q-network



Last four states

Figure 27: DTQN architecture [6].

7.1.2 DTQN: Evaluation




Figure 28: DTQN performance [6].

7.2.1 Decision Transformer: Architecture

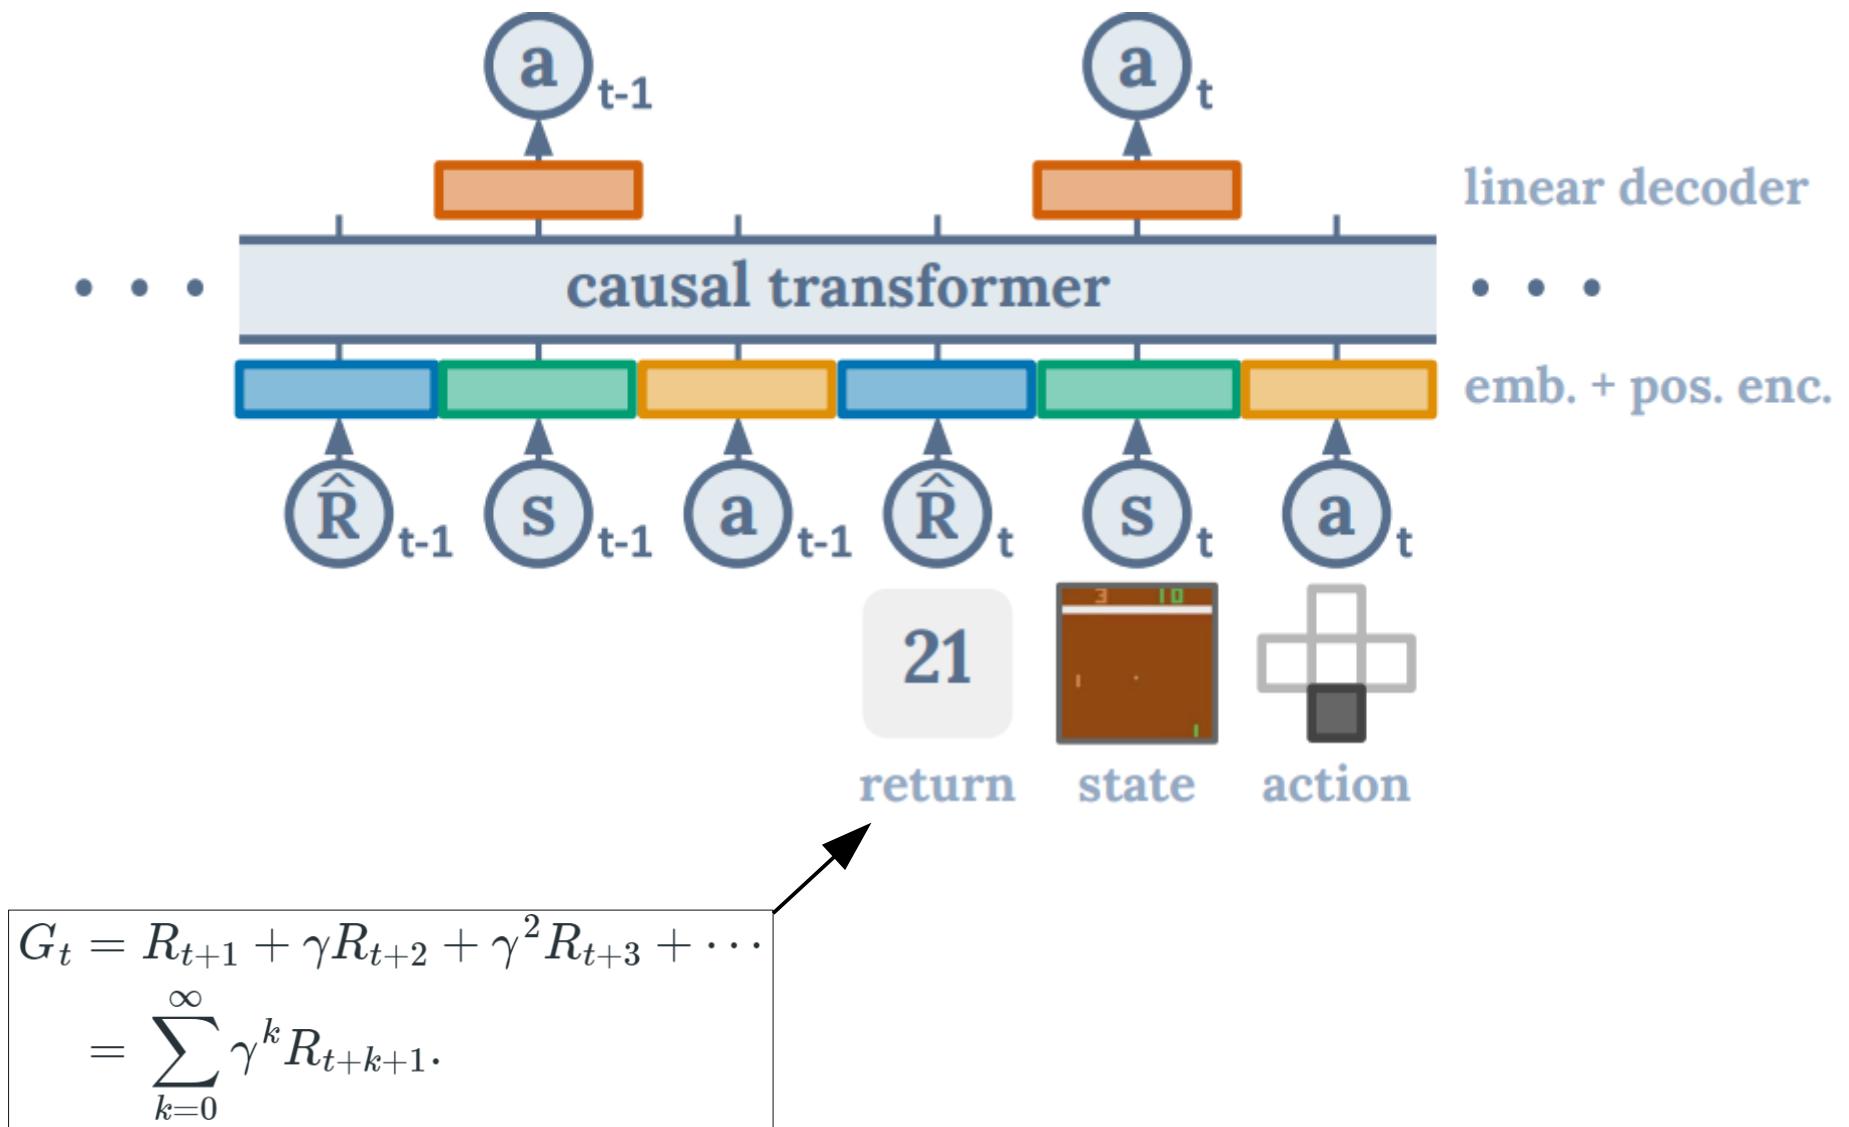
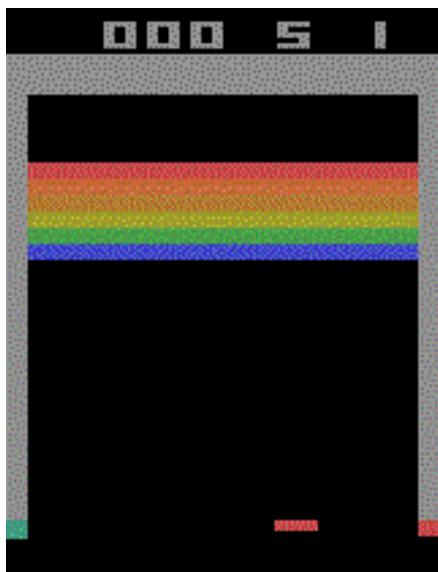
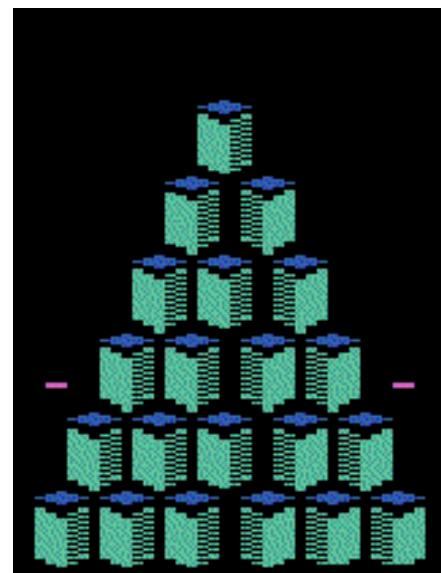


Figure 29: Decision Transformer architecture [7].


7.2.2 Decision Transformer: Evaluation




Breakout



Qbert



Pong



Seaquest

Figure 30: Evaluated environments.

Image sources:

https://www.gymlibrary.ml/_images/breakout.gif

https://www.gymlibrary.ml/_images/qbert.gif

https://www.gymlibrary.ml/_images/pong.gif

https://www.gymlibrary.ml/_images/seaquest.gif (call dates: 20.07.22)

7.2.3 Decision Transformer: Evaluation

CQL: conservative Q-Learning

QR-DQN: distributional RL

REM: random ensemble mixture

Game	DT	CQL	QR-DQN	REM
Breakout	267.5 ± 97.5	211.1	21.1	32.1
Qbert	25.1 ± 18.1	104.2	1.7	1.4
Pong	106.1 ± 8.1	111.9	20.0	39.1
Seaquest	2.4 ± 0.7	1.7	1.4	1.0

Table 4: Performance of the Decision Transformer compared with state-of-the-art DRL methods [7].

7.2.4 Decision Transformer: Why attention?

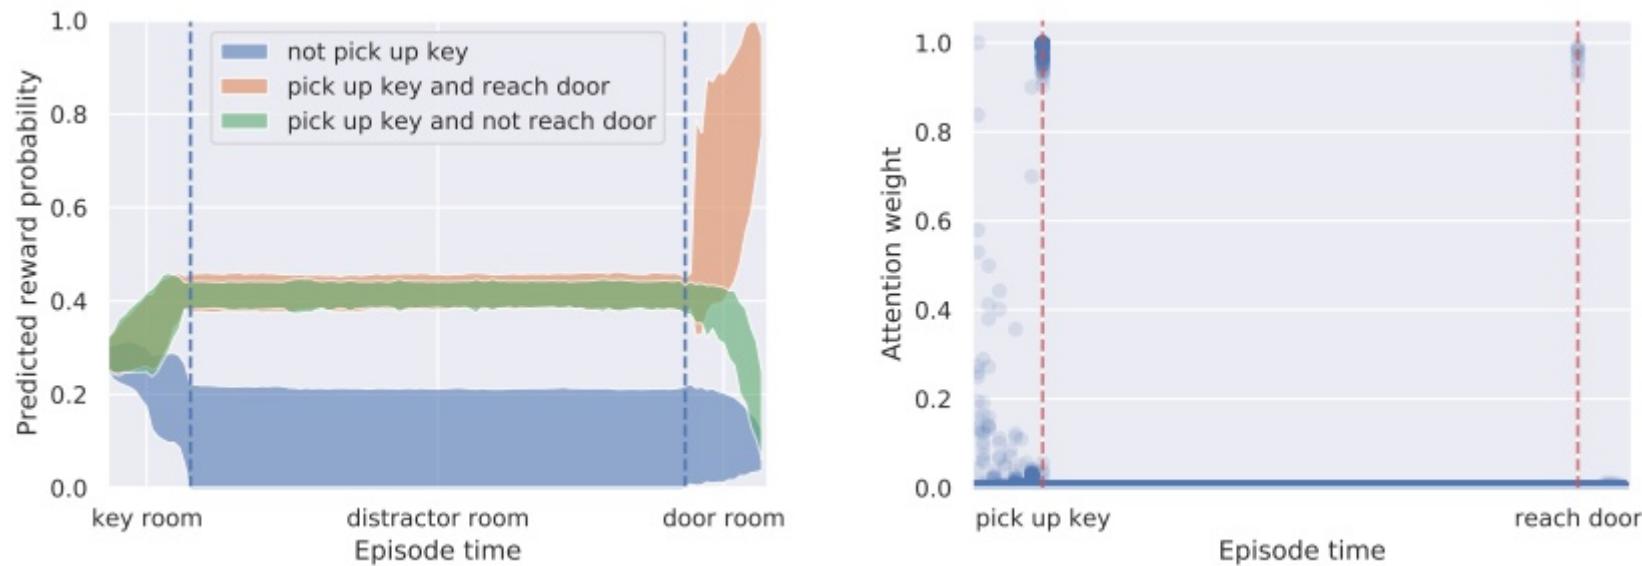




Figure 31: Decision Transformer attend to pivotal events [7].



8. Modifications of attention: Synthesizer, Linformer, Reformer


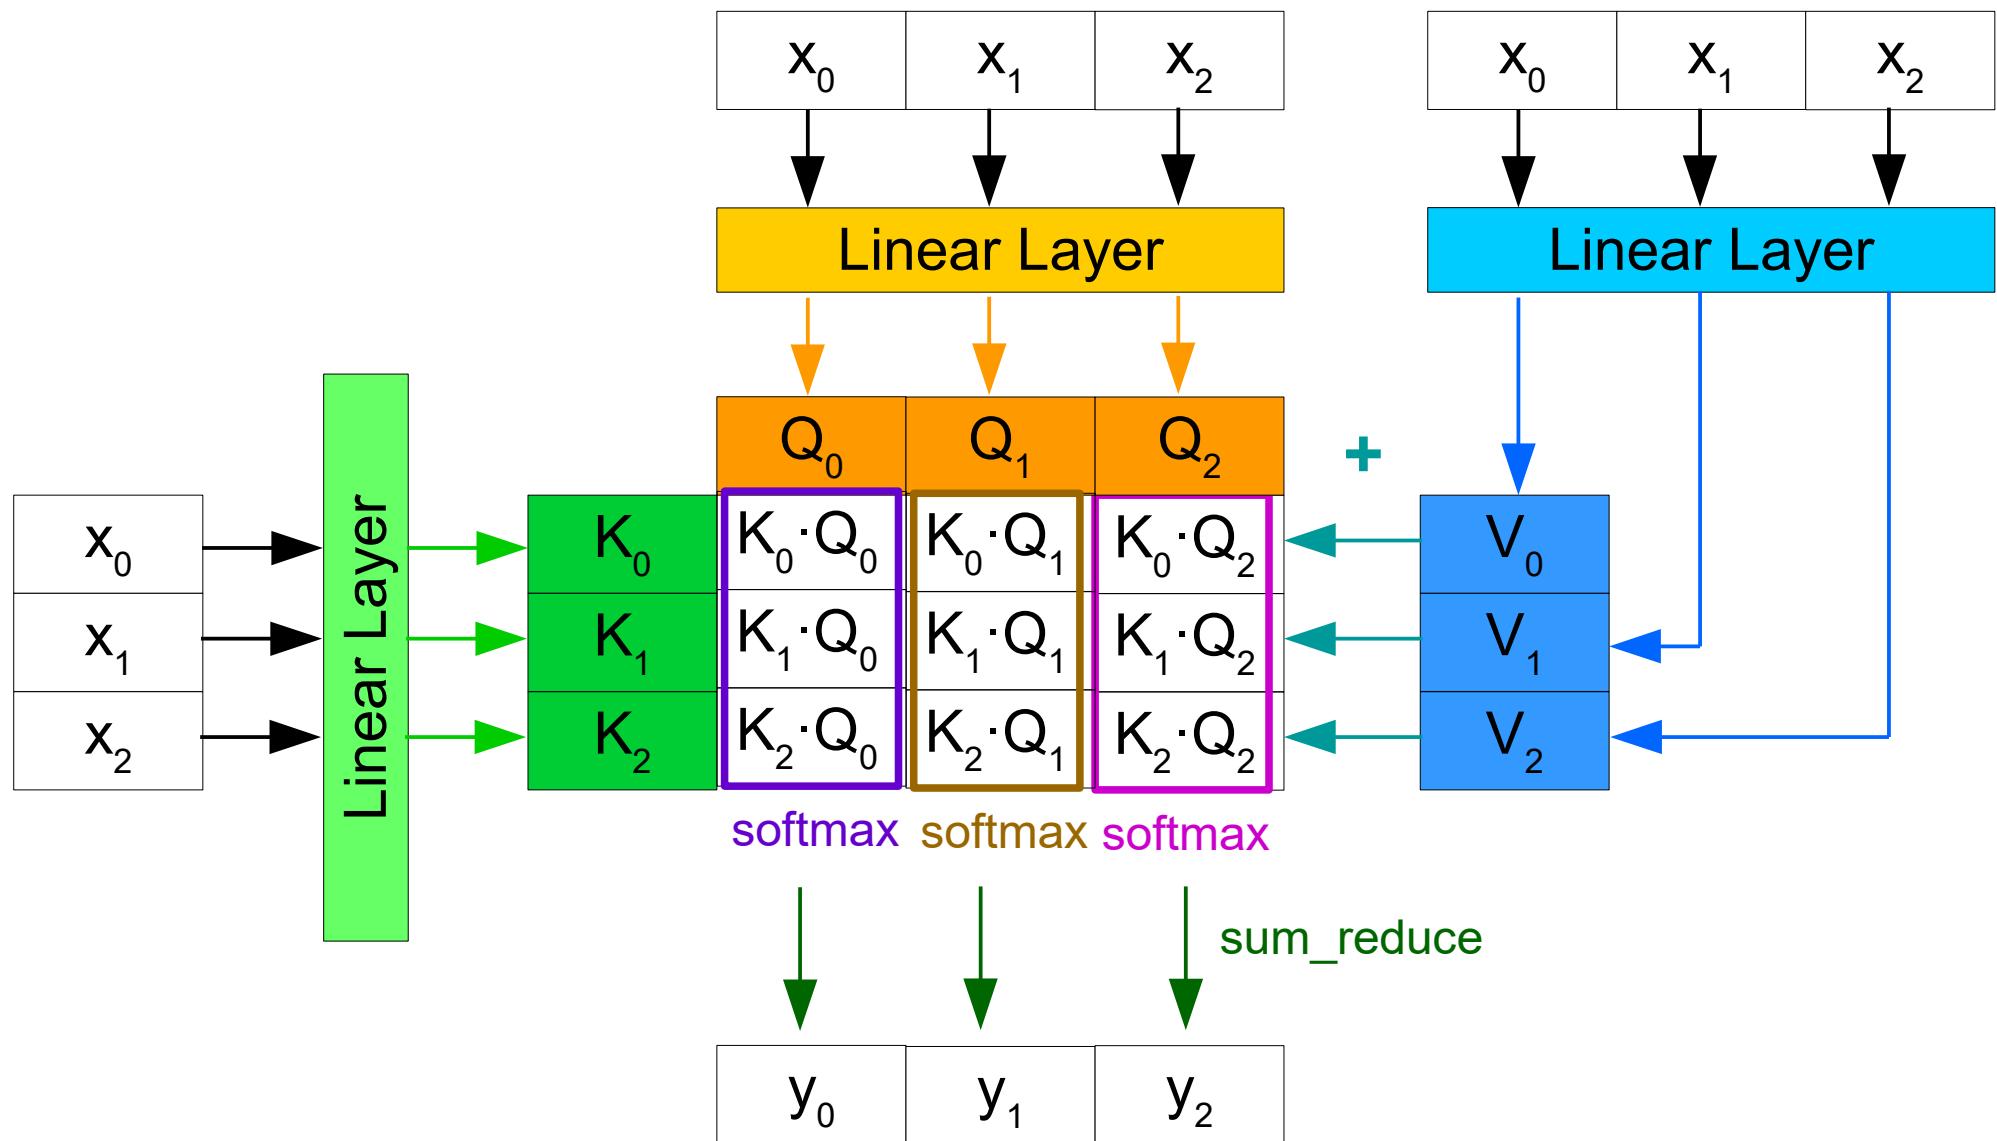


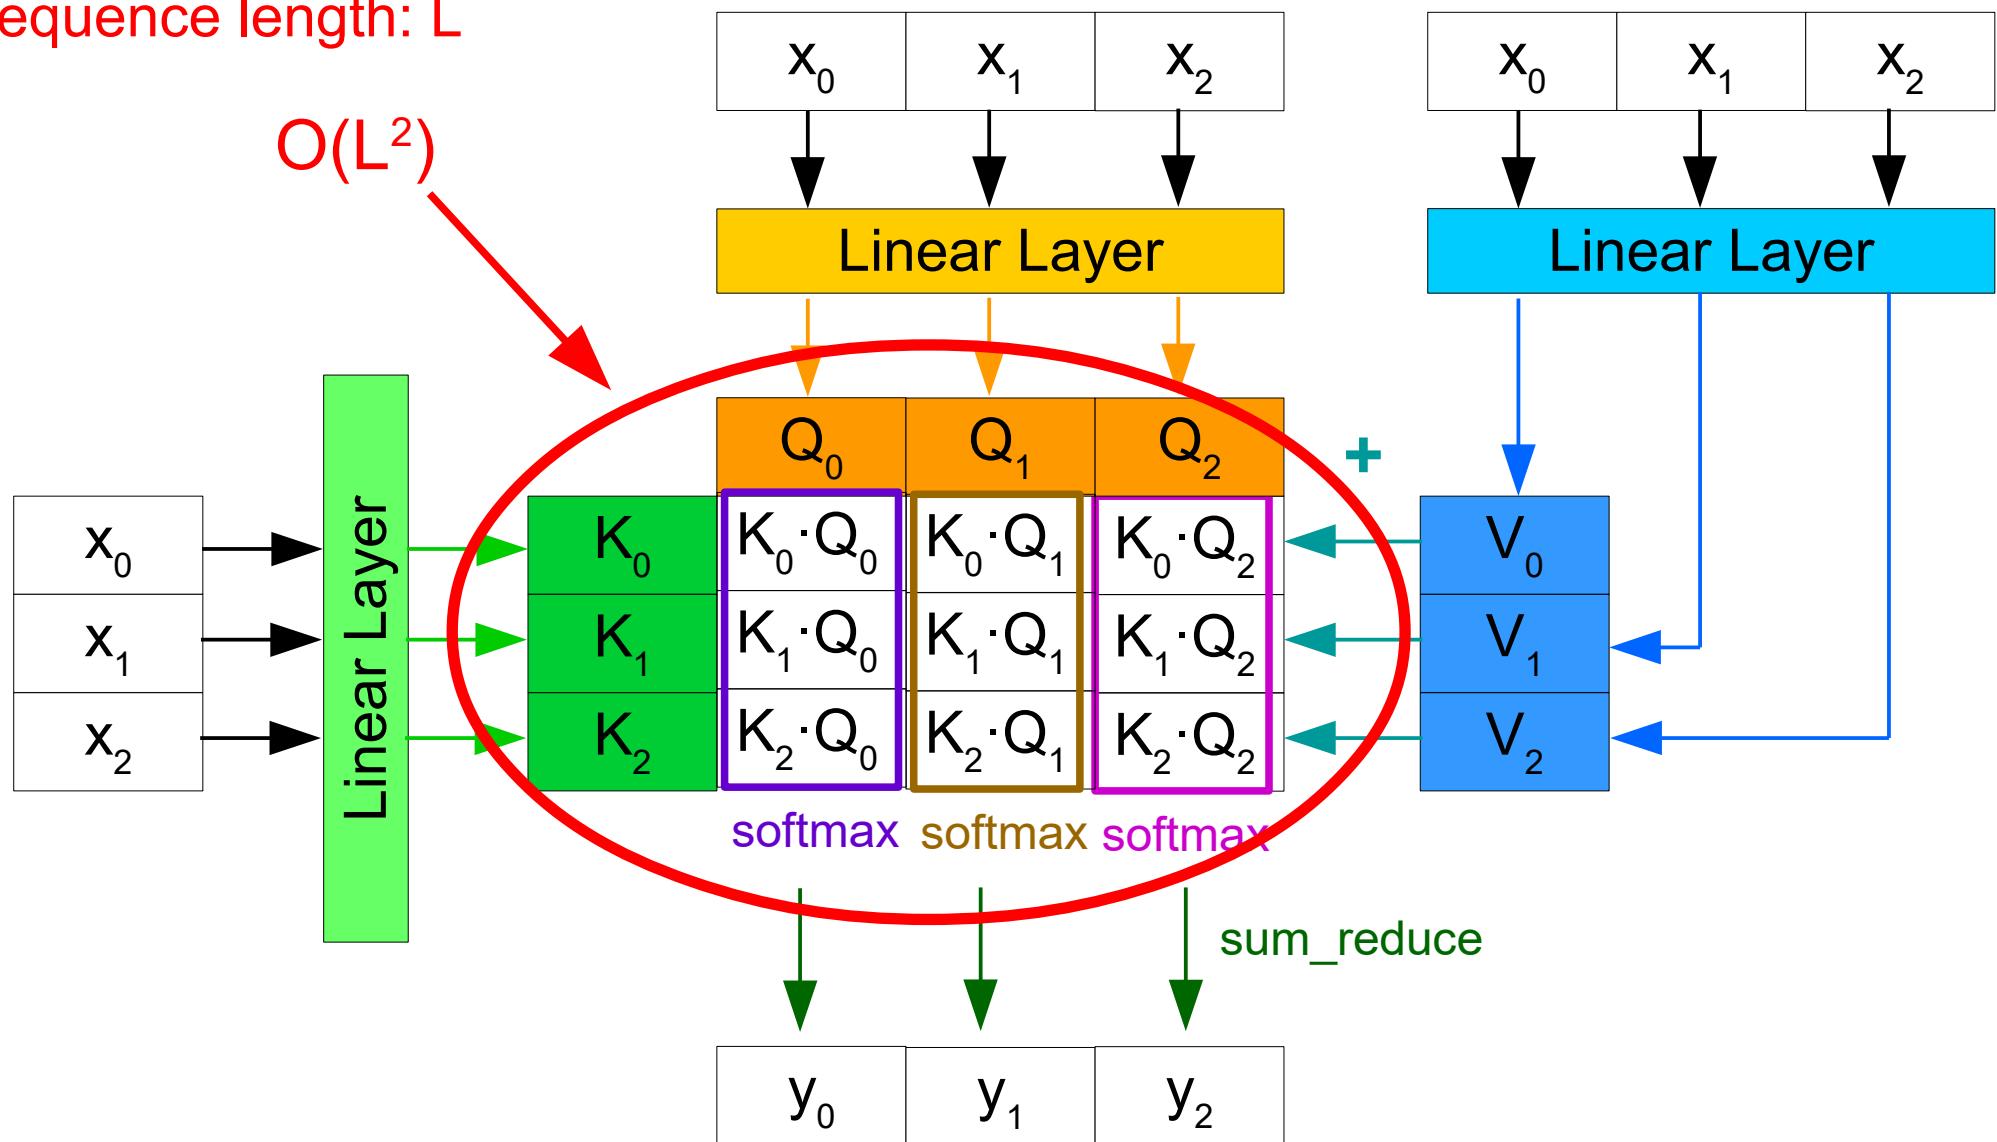
Image sources:
<https://cdn-icons-png.flaticon.com/512/5116/5116423.png>
 (call dates: 21.01.23)
 Tay, Yi et al. "Synthesizer: Rethinking Self-Attention in Transformer Models." International Conference on Machine Learning (2020).
 Wang, Sinong et al. "Linformer: Self-Attention with Linear Complexity." ArXiv abs/2006.04768 (2020): n. pag.

8. Self-attention: Time and Space Complexity



8. Self-attention: Time and Space Complexity

Sequence length: L



8.1 Synthesizer: *Dense* variant


x_0	x_1	x_2
-------	-------	-------

8.1 Synthesizer: *Dense* variant


x_0	x_1	x_2
-------	-------	-------

Linear Layer

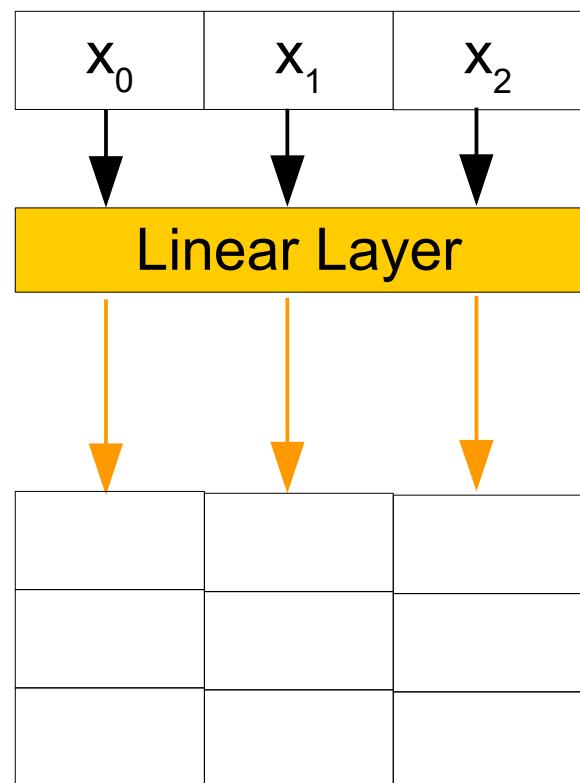
8.1 Synthesizer: *Dense* variant




8.1 Synthesizer: *Dense* variant



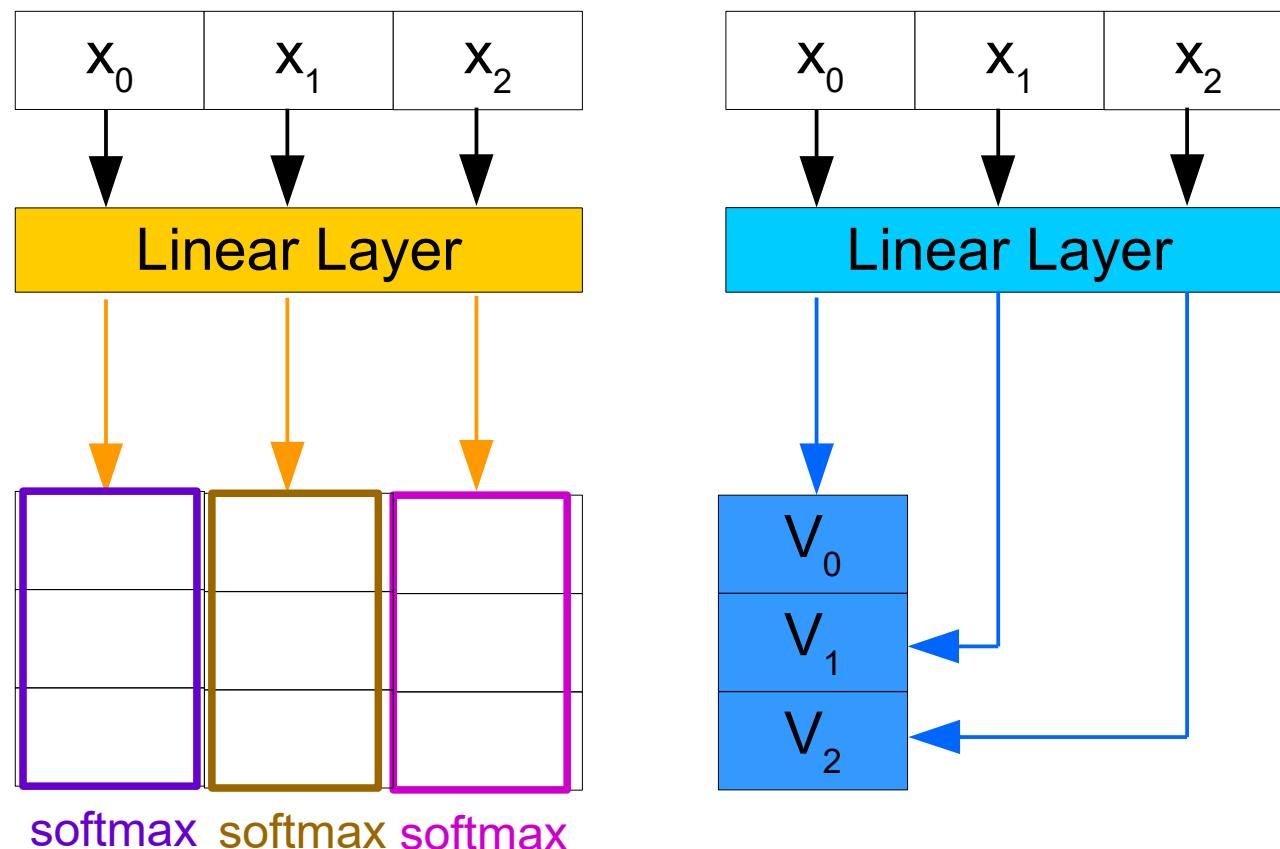
8.1 Synthesizer: *Dense* variant



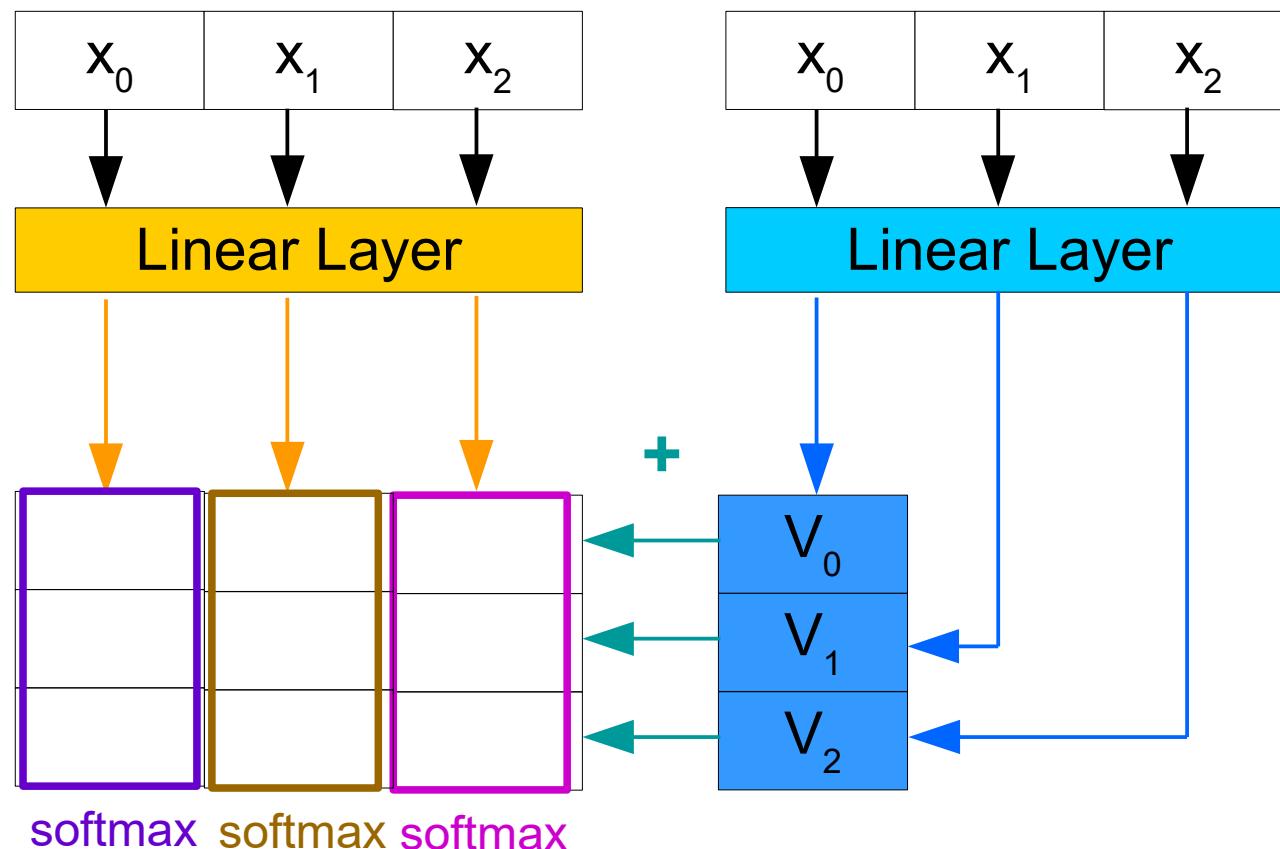
8.1 Synthesizer: *Dense* variant




8.1 Synthesizer: *Dense* variant



8.1 Synthesizer: *Dense* variant



8.1 Synthesizer: *Dense* variant



8.1 Synthesizer: *Random* variant


x_0	x_1	x_2
-------	-------	-------

Fixed


8.1 Synthesizer: *Random* variant

x_0	x_1	x_2
-------	-------	-------


Fixed




8.1 Synthesizer: *Random* variant




8.1 Synthesizer: *Random* variant



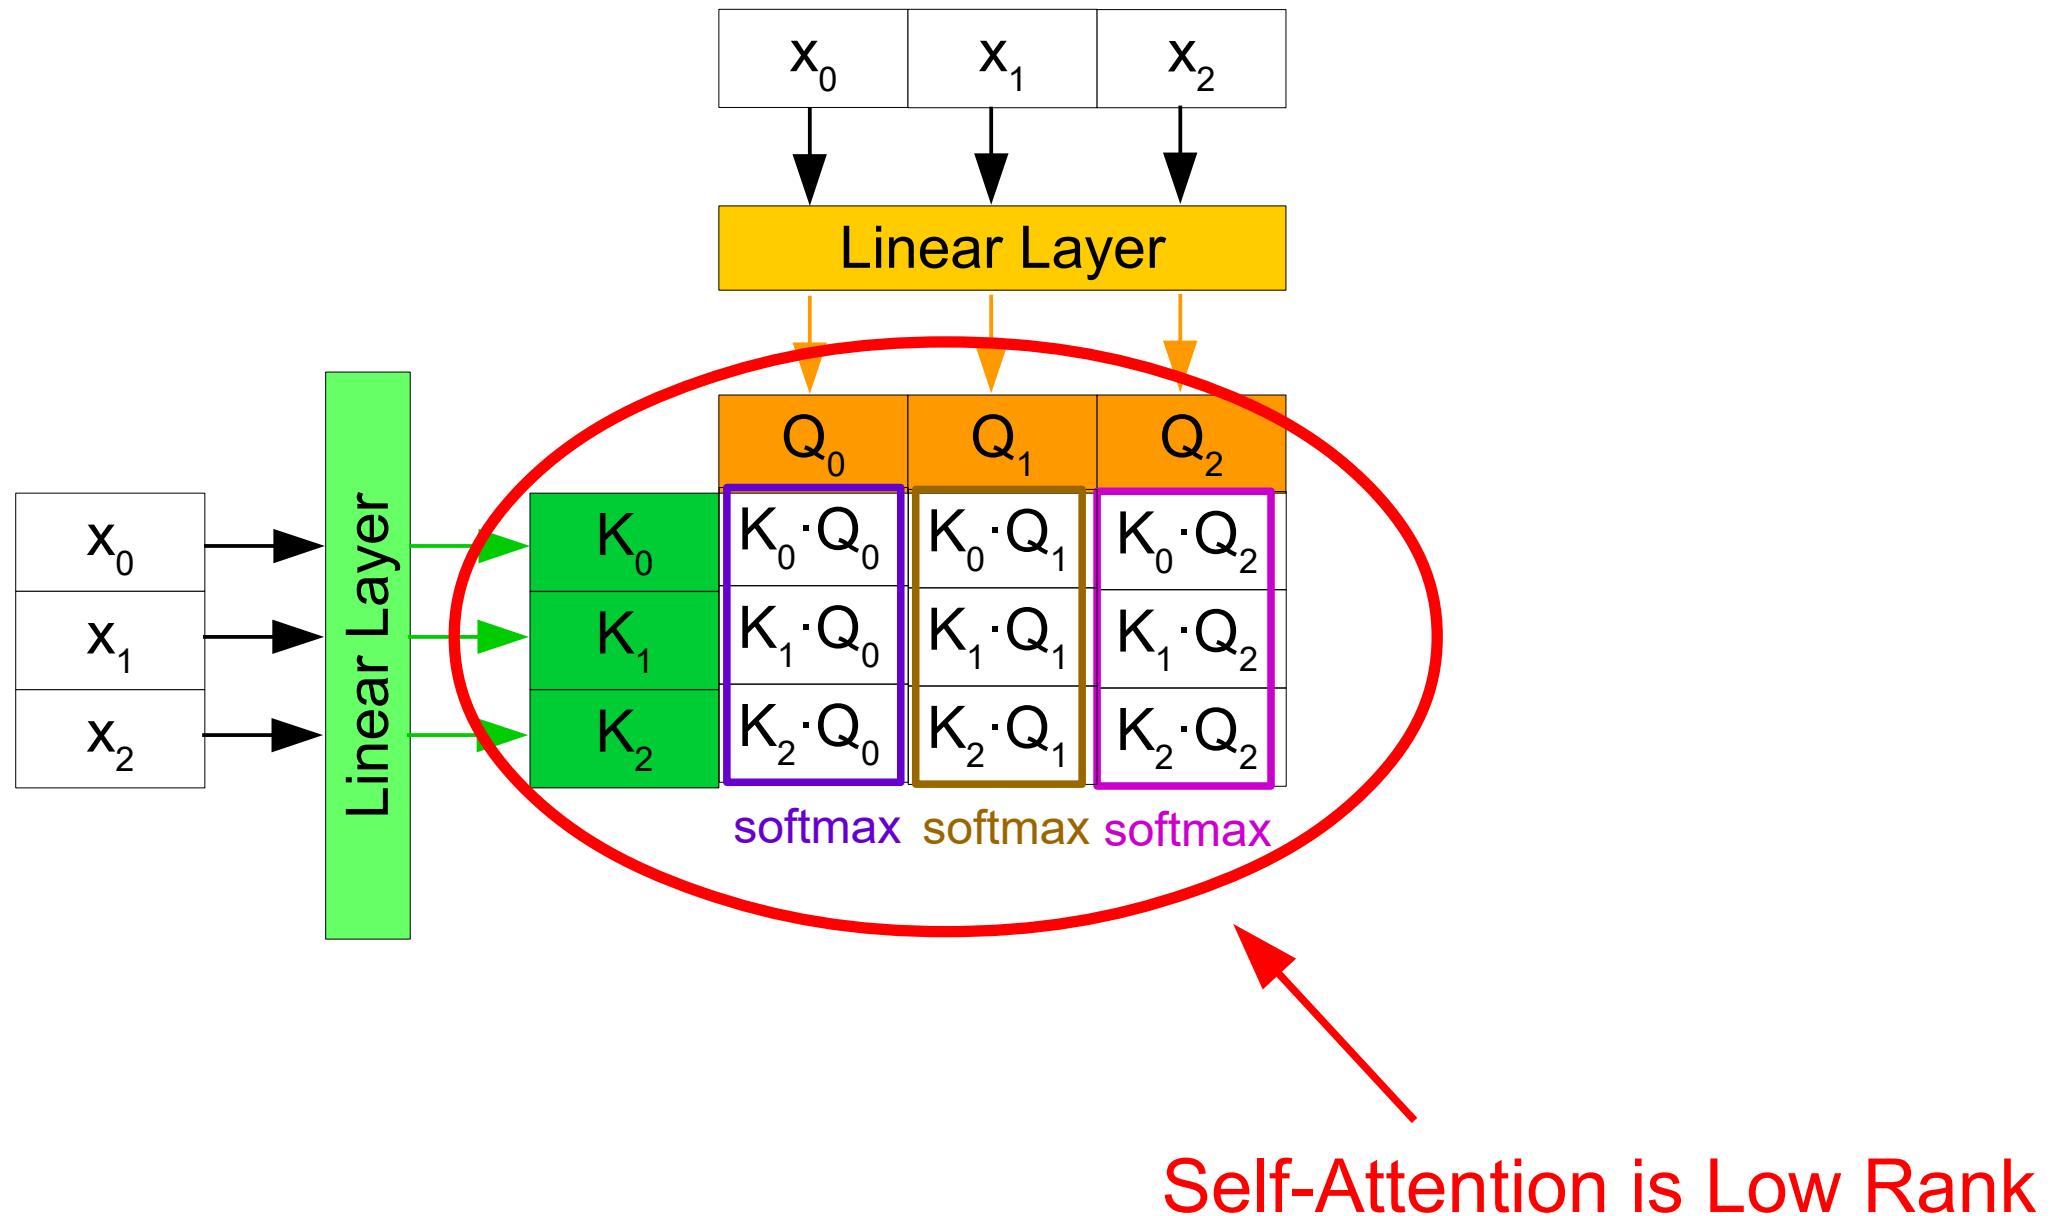
8.1 Synthesizer: *Random* variant




8.2 Linformer: Self-Attention with Linear Complexity



8.2 Linformer: Self-Attention with Linear Complexity



8.2 Linformer: Self-Attention with Linear Complexity




Eigenvalue spectrum of a 256×256 matrix (full rank).

8.2 Linformer: Self-Attention with Linear Complexity

“Order a set of uniform distributed numbers”

→ All dimensions have the same amount of information





Eigenvalue spectrum of a 256x256 matrix (full rank).

8.2 Linformer: Self-Attention with Linear Complexity

“Order a set of uniform distributed numbers”

→ All dimensions have the same amount of information




Eigenvalue spectrum of a 256x256 matrix (full rank).

8.2 Linformer: Self-Attention with Linear Complexity

$A = 256 \times 128$

$B = 128 \times 256$

$M = A \cdot B$



Eigenvalue spectrum of a 256×256 matrix (low rank).


8.2 Linformer: Self-Attention with Linear Complexity

$$A = 256 \times 128$$

$$B = 128 \times 256$$

$$M = A \cdot B$$

Information is concentrated along few dimensions



Eigenvalue spectrum of a 256×256 matrix (low rank).



8.2 Linformer: Self-Attention with Linear Complexity

$$A = 256 \times 128$$

$$B = 128 \times 256$$

$$M = A \cdot B$$

Information is concentrated along few dimensions



Eigenvalue spectrum of a 256×256 matrix (low rank).

8.2 Linformer: Self-Attention is Low Rank




Image sources:

Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention is Low Rank






Image sources:

Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention is Low Rank

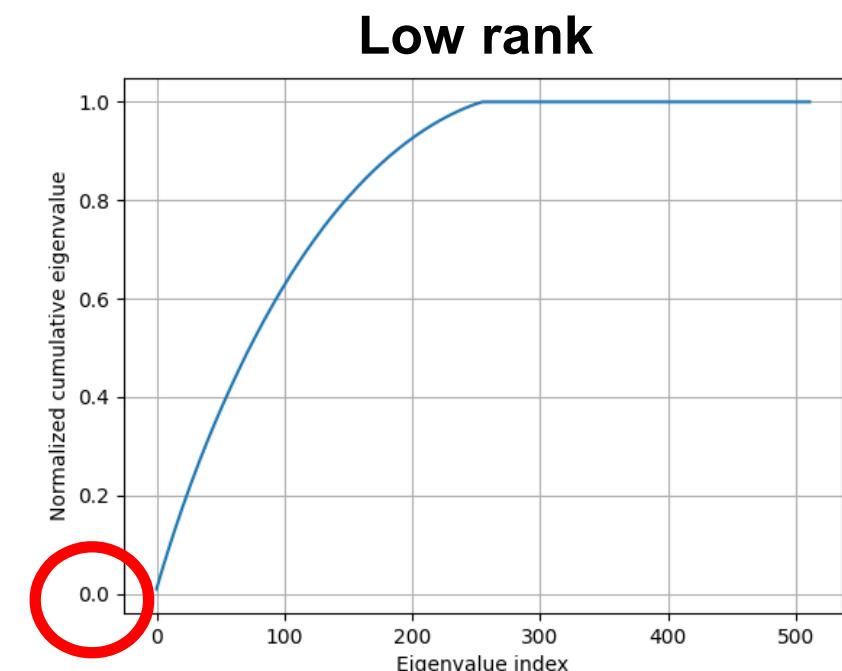
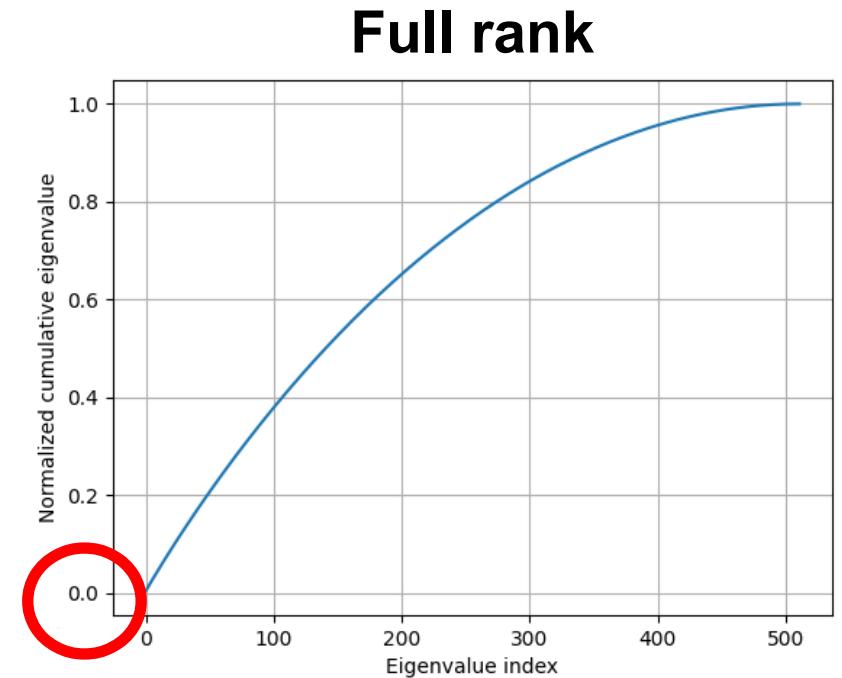



Image sources:

Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention with Linear Complexity

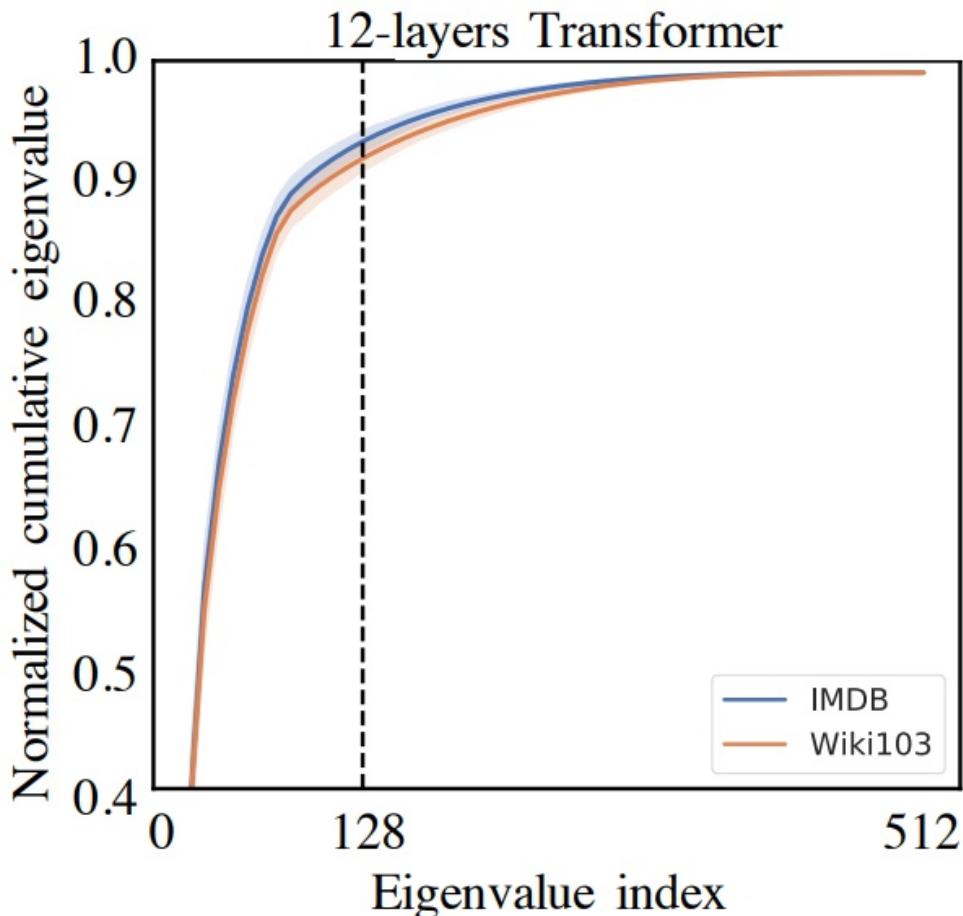


Image sources:

Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention with Linear Complexity







Image sources:

Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention with Linear Complexity




With increasing depth:
Put more and more
information into fewer
and fewer dimensions

Image sources:


Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention with Linear Complexity

More and more skewed



Eigenvalue index: 128




With increasing deep:
Put more and more
information into fewer
and fewer dimensions

Image sources:

Wang, S., Li, B.Z., Khabsa, M., Fang, H., & Ma, H. (2020). Linformer: Self-Attention with Linear Complexity. ArXiv, abs/2006.04768.

8.2 Linformer: Self-Attention with Linear Complexity



8.2 Linformer: Self-Attention with Linear Complexity

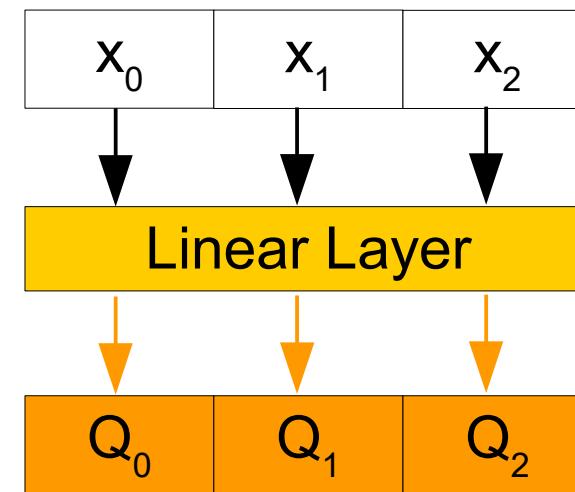
x_0	x_1	x_2
-------	-------	-------

8.2 Linformer: Self-Attention with Linear Complexity


x_0	x_1	x_2
-------	-------	-------

Linear Layer


8.2 Linformer: Self-Attention with Linear Complexity




8.2 Linformer: Self-Attention with Linear Complexity




8.2 Linformer: Self-Attention with Linear Complexity



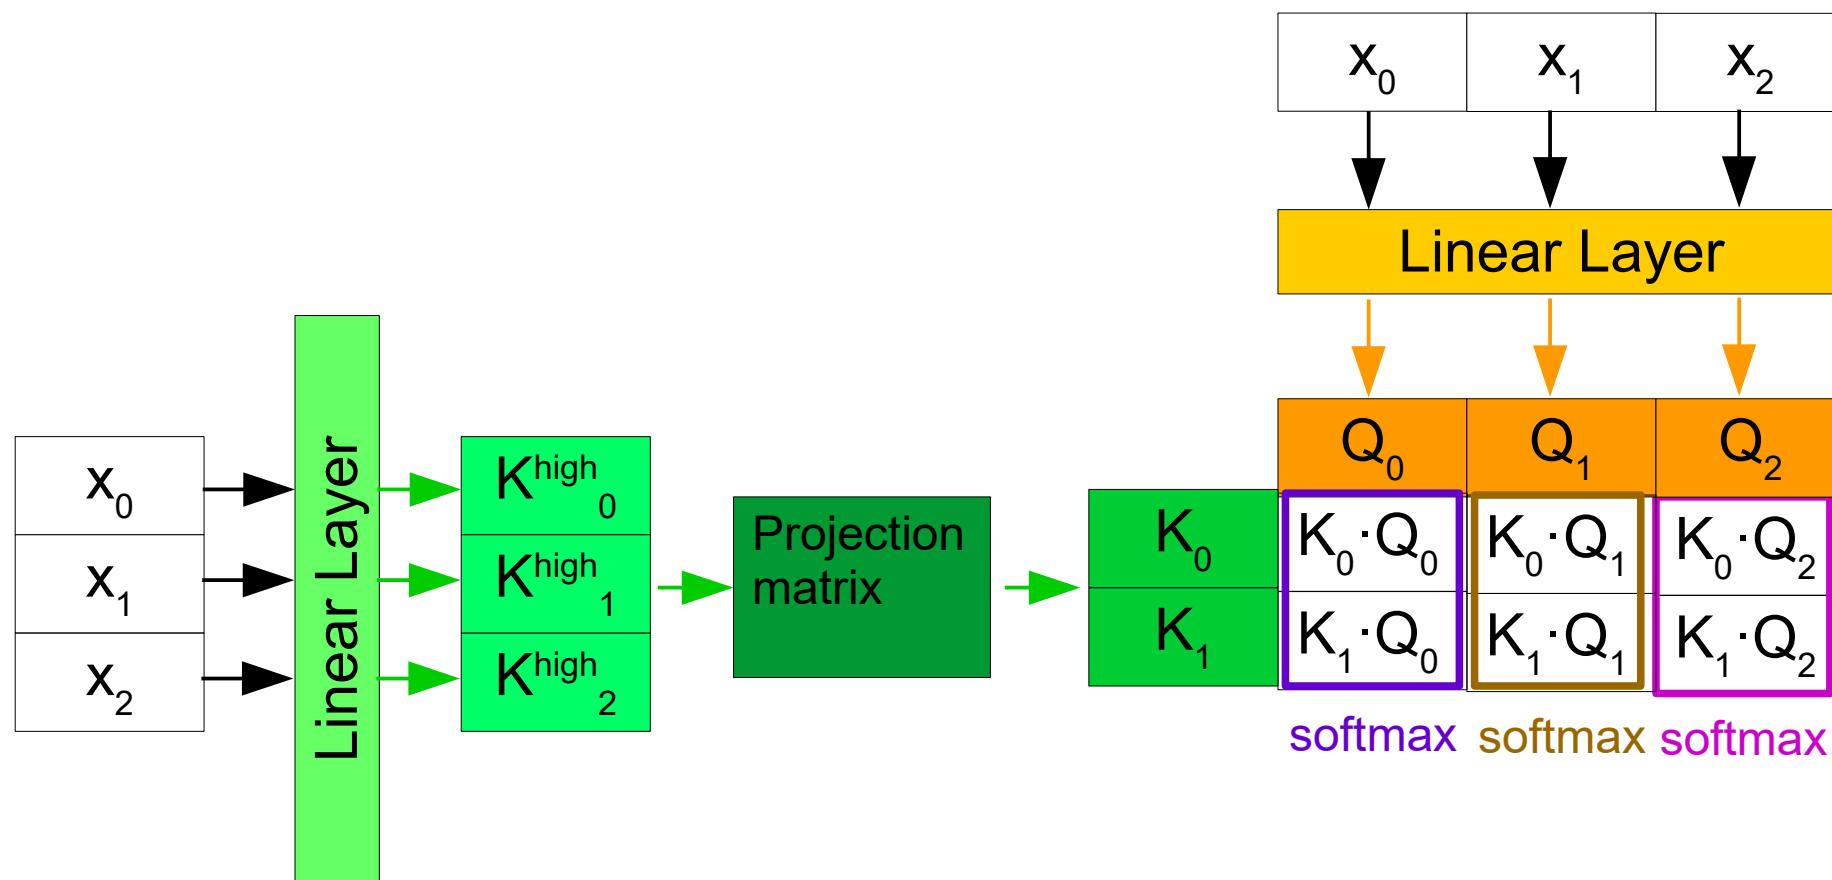
8.2 Linformer: Self-Attention with Linear Complexity



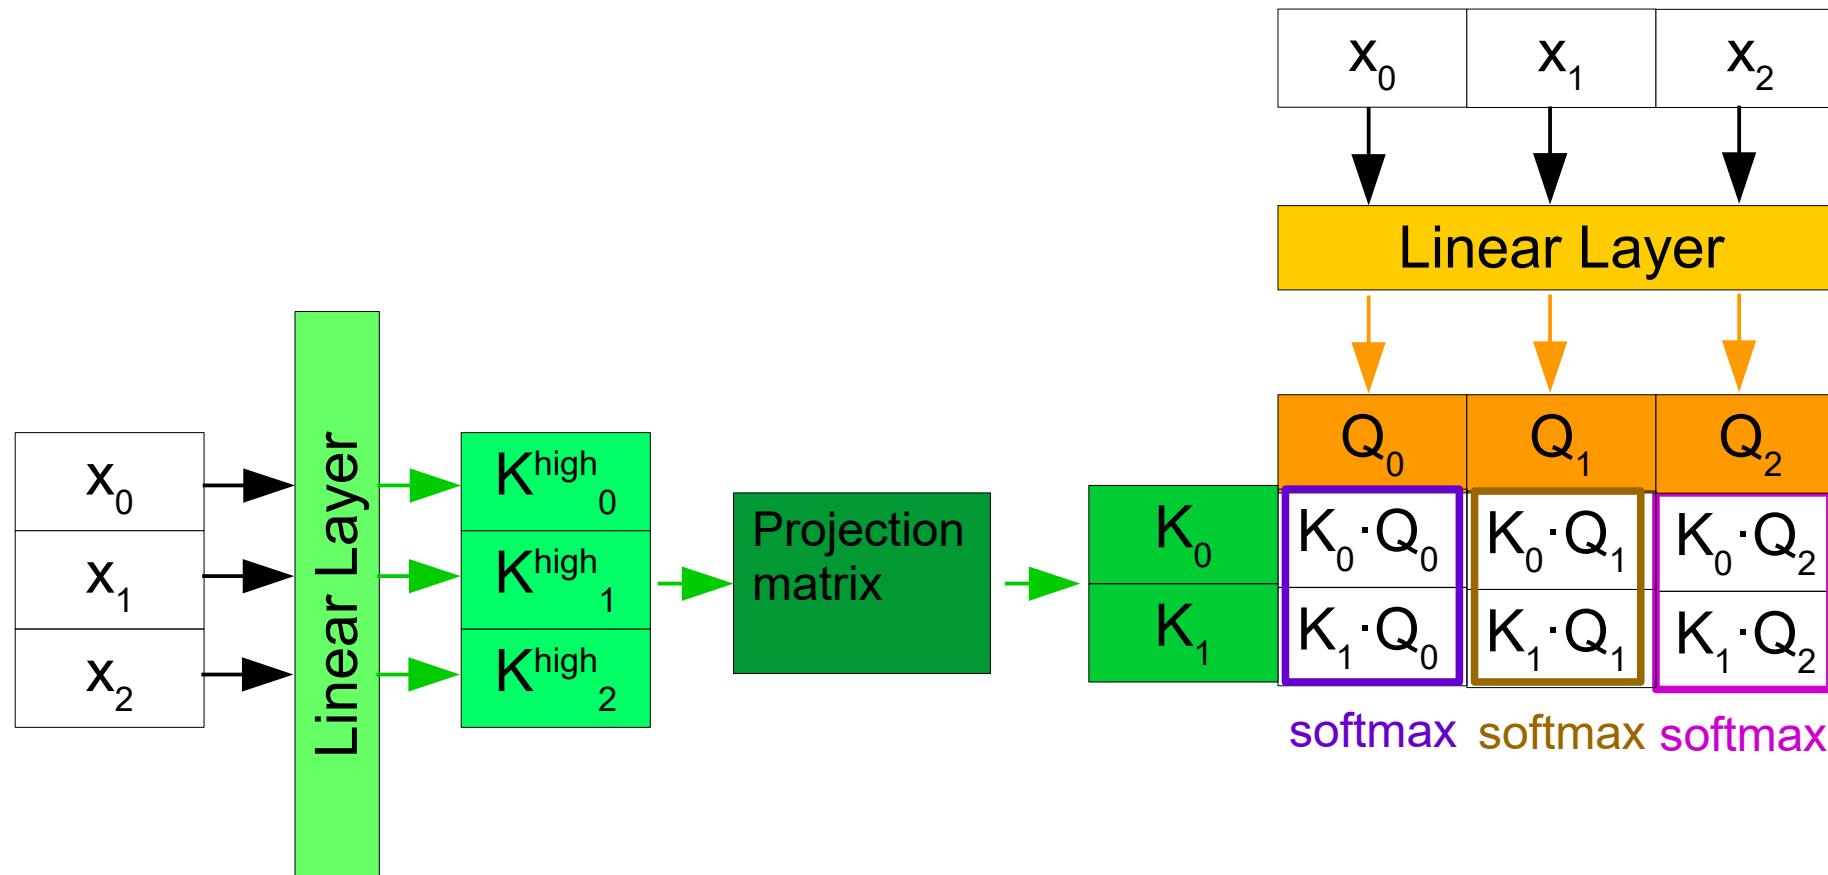

8.2 Linformer: Self-Attention with Linear Complexity



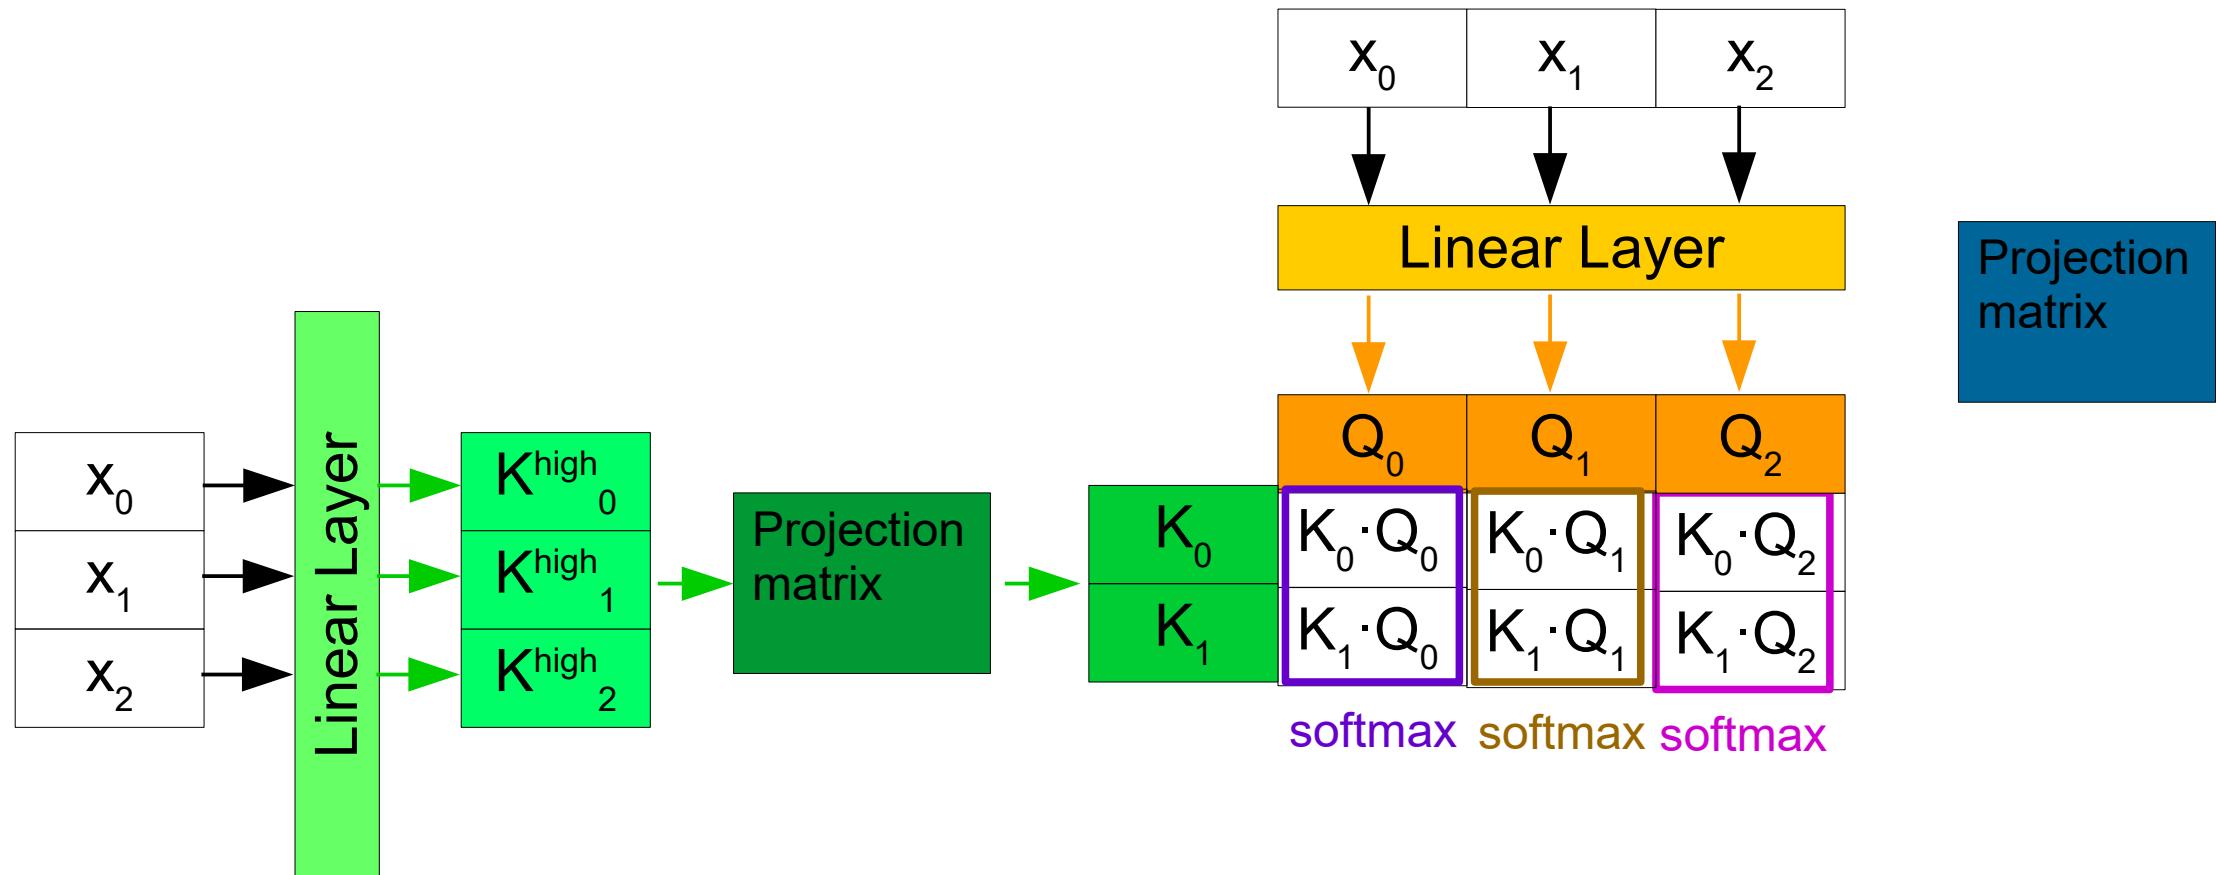

8.2 Linformer: Self-Attention with Linear Complexity




8.2 Linformer: Self-Attention with Linear Complexity



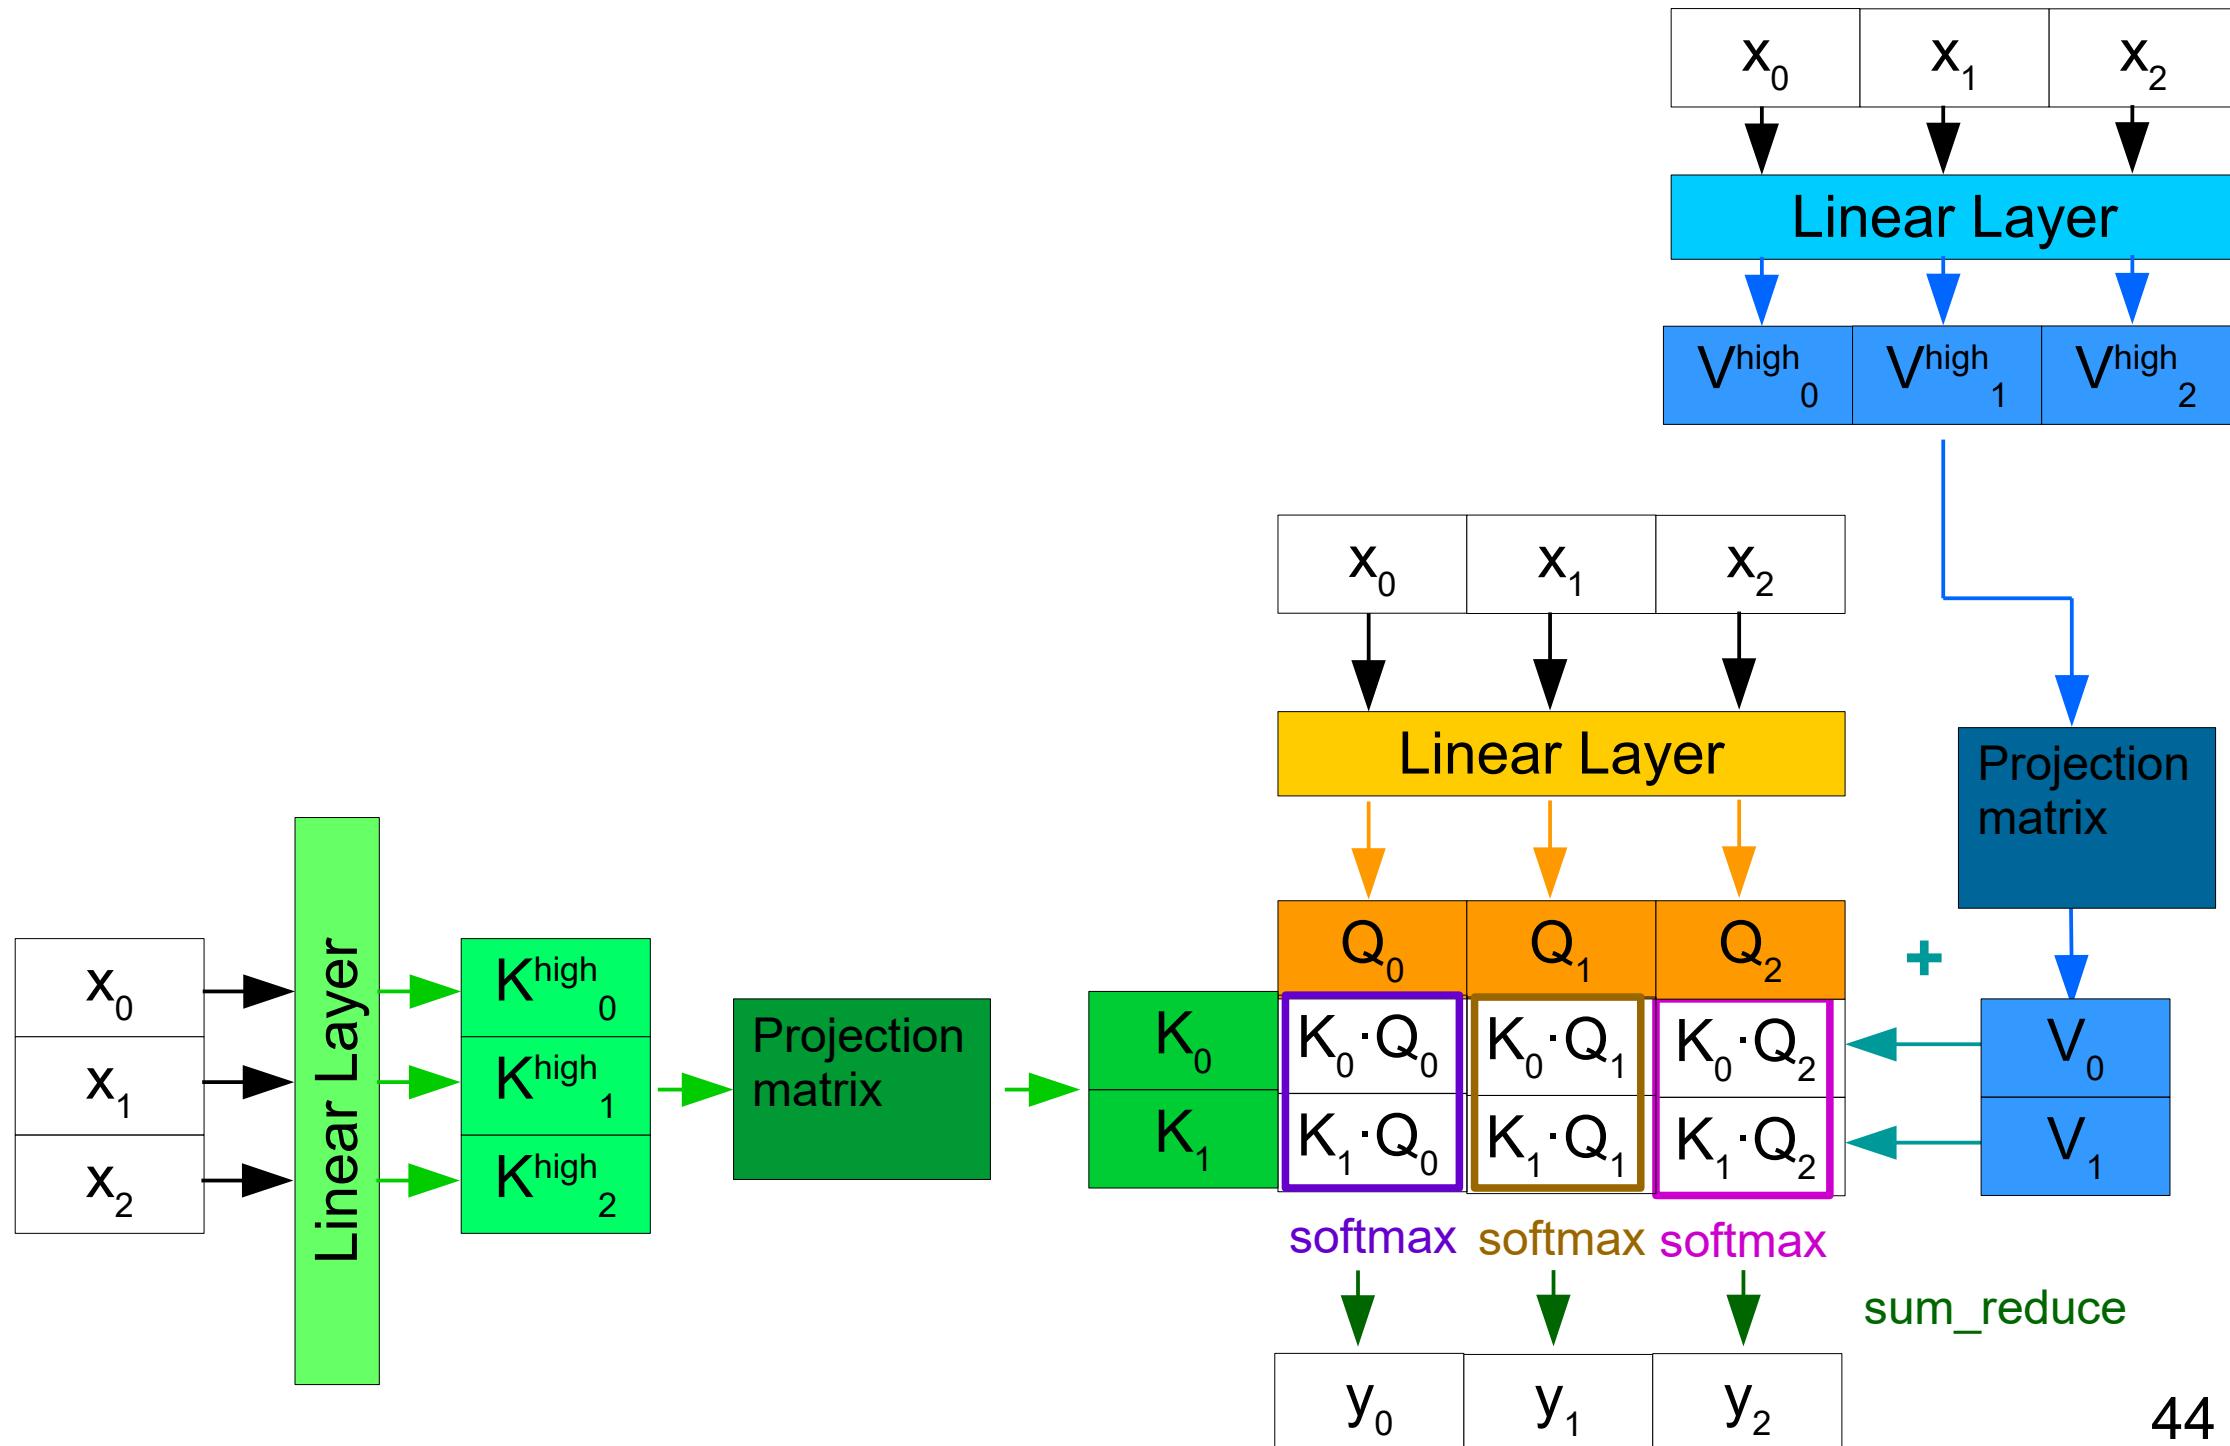
8.2 Linformer: Self-Attention with Linear Complexity



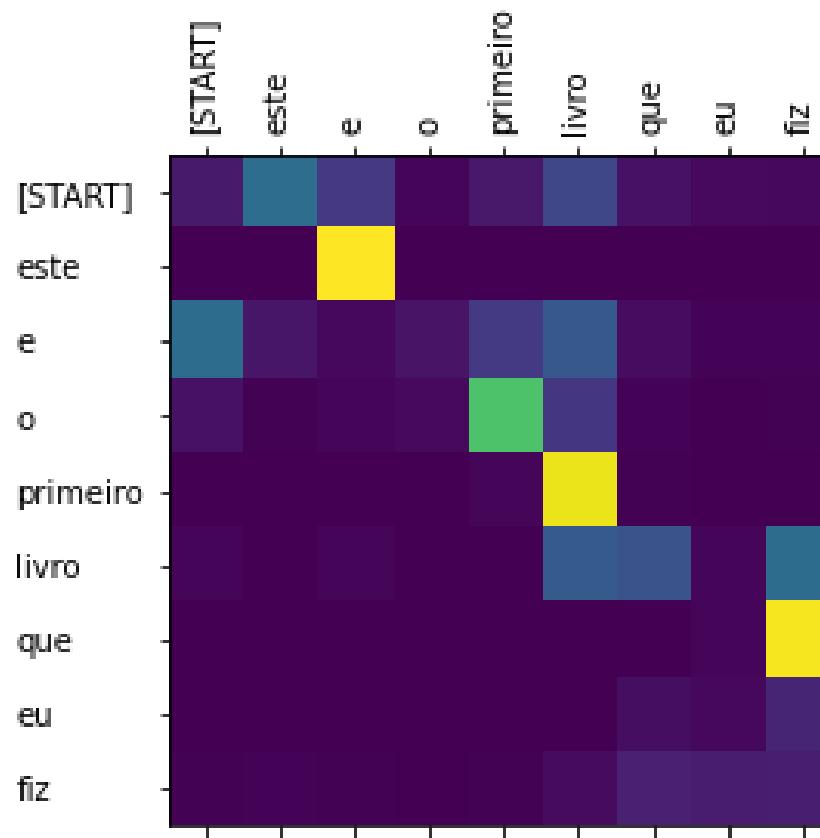
8.2 Linformer: Self-Attention with Linear Complexity



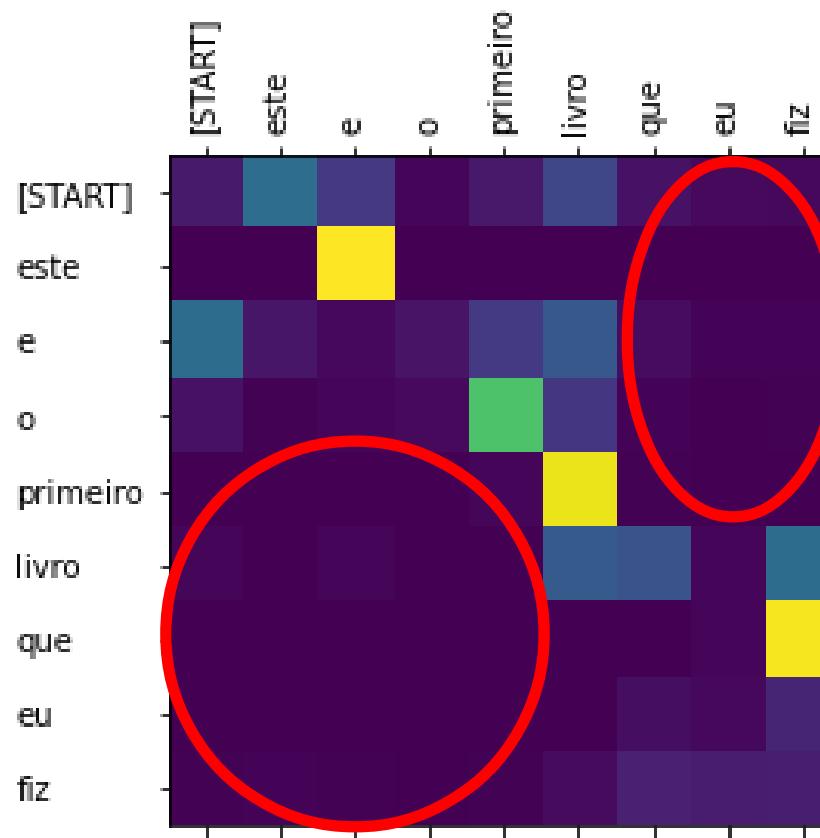
8.2 Linformer: Self-Attention with Linear Complexity



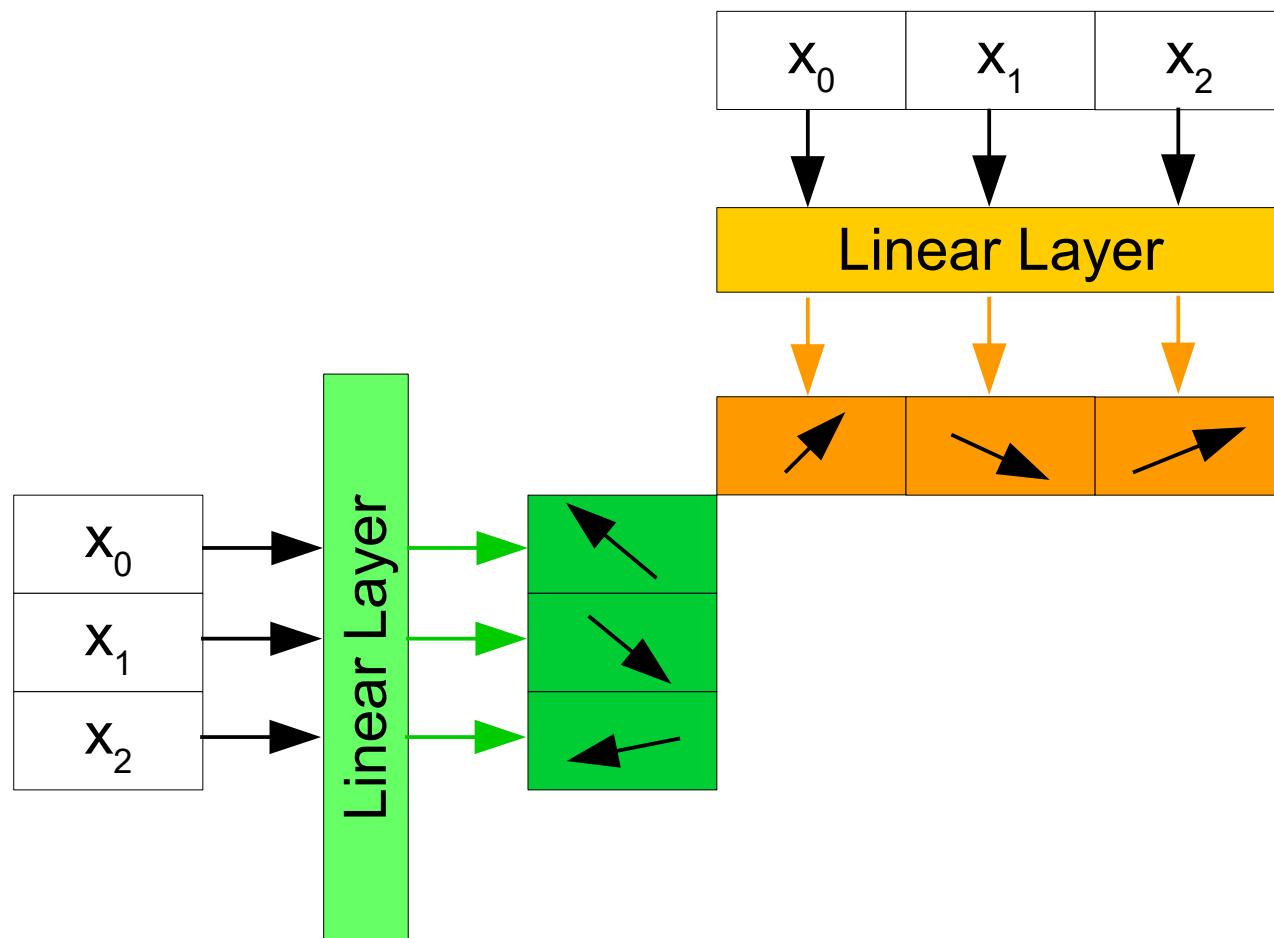
8.3 Reformer: The Efficient Transformer




8.3 Reformer: The Efficient Transformer




8.3 Reformer: The Efficient Transformer



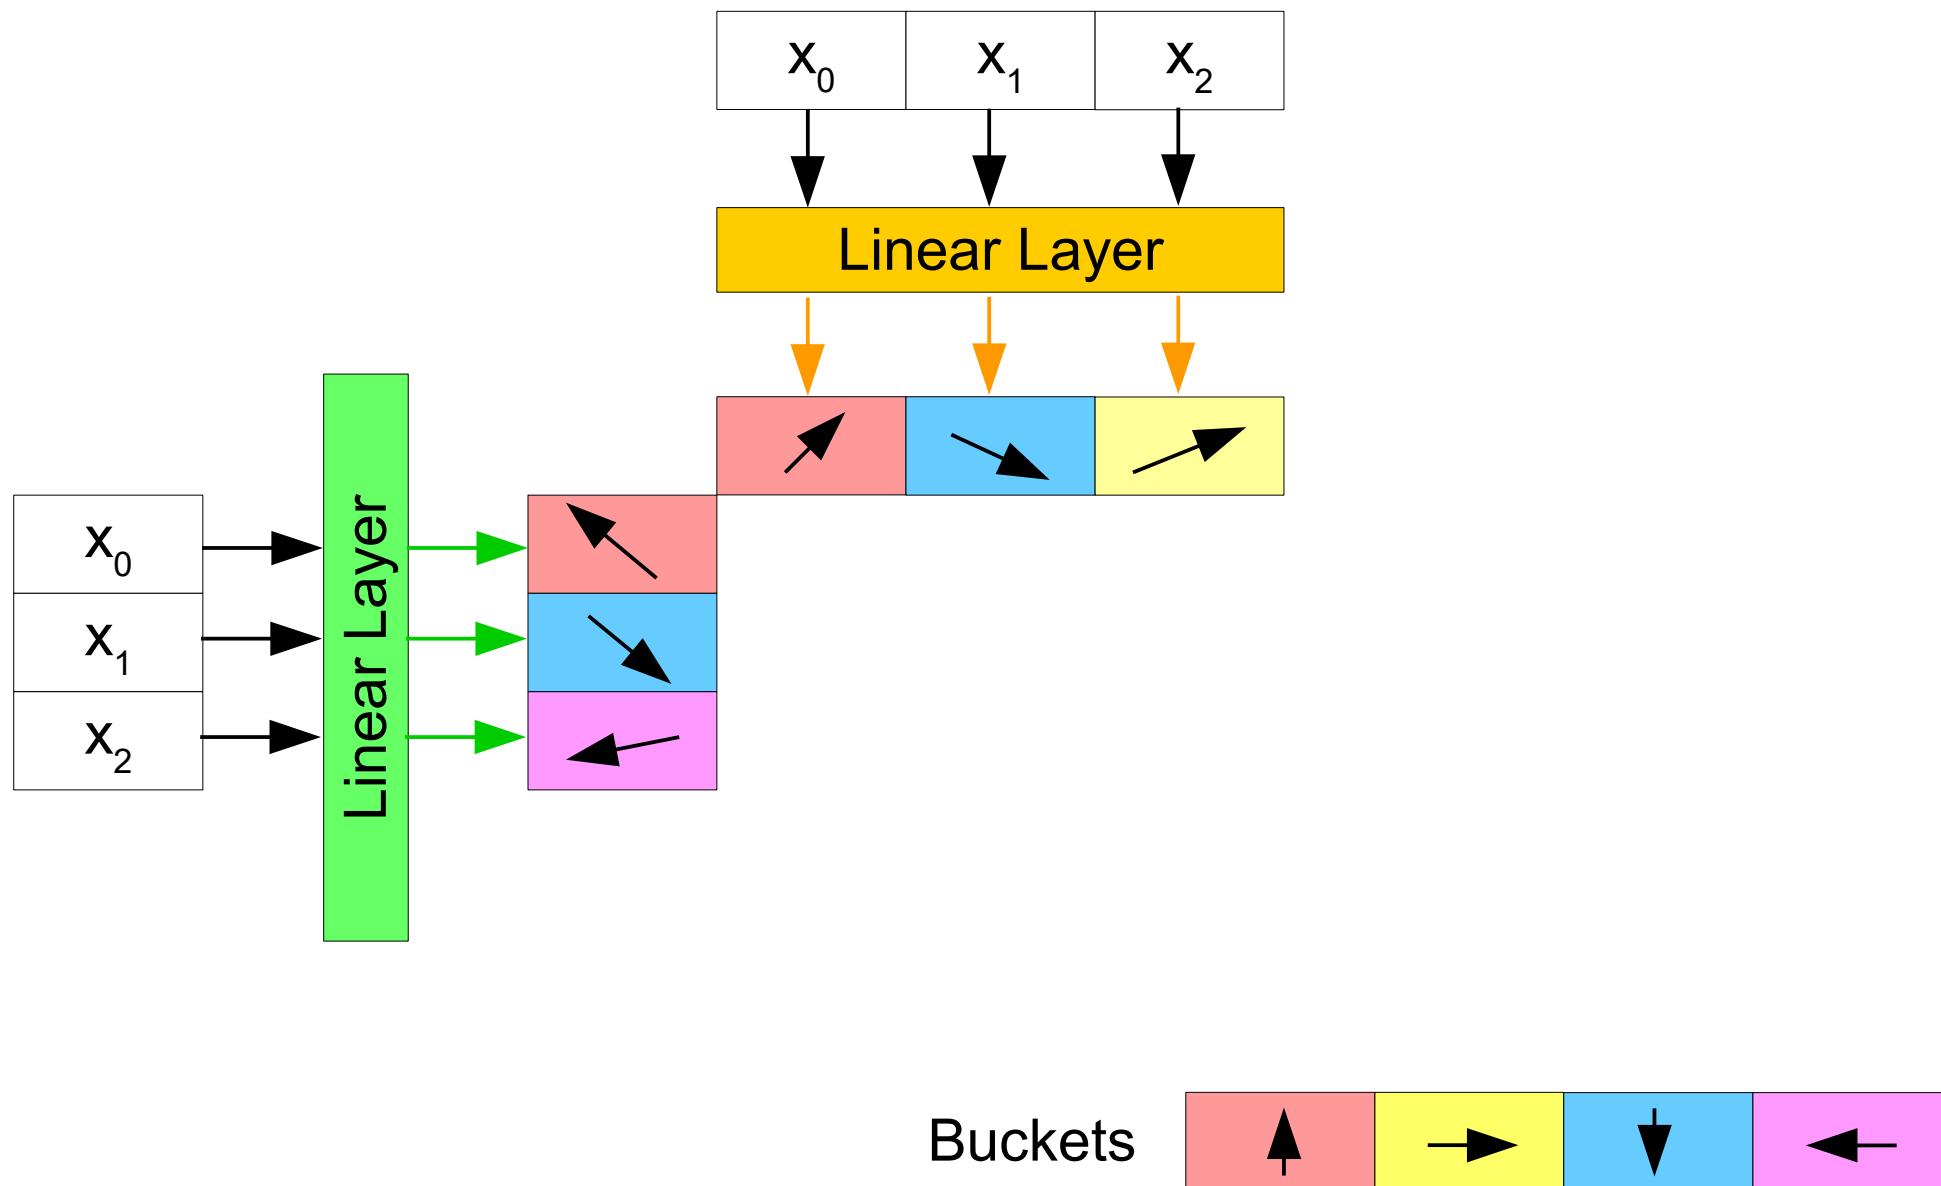
8.3 Reformer: The Efficient Transformer




8.3 Reformer: The Efficient Transformer



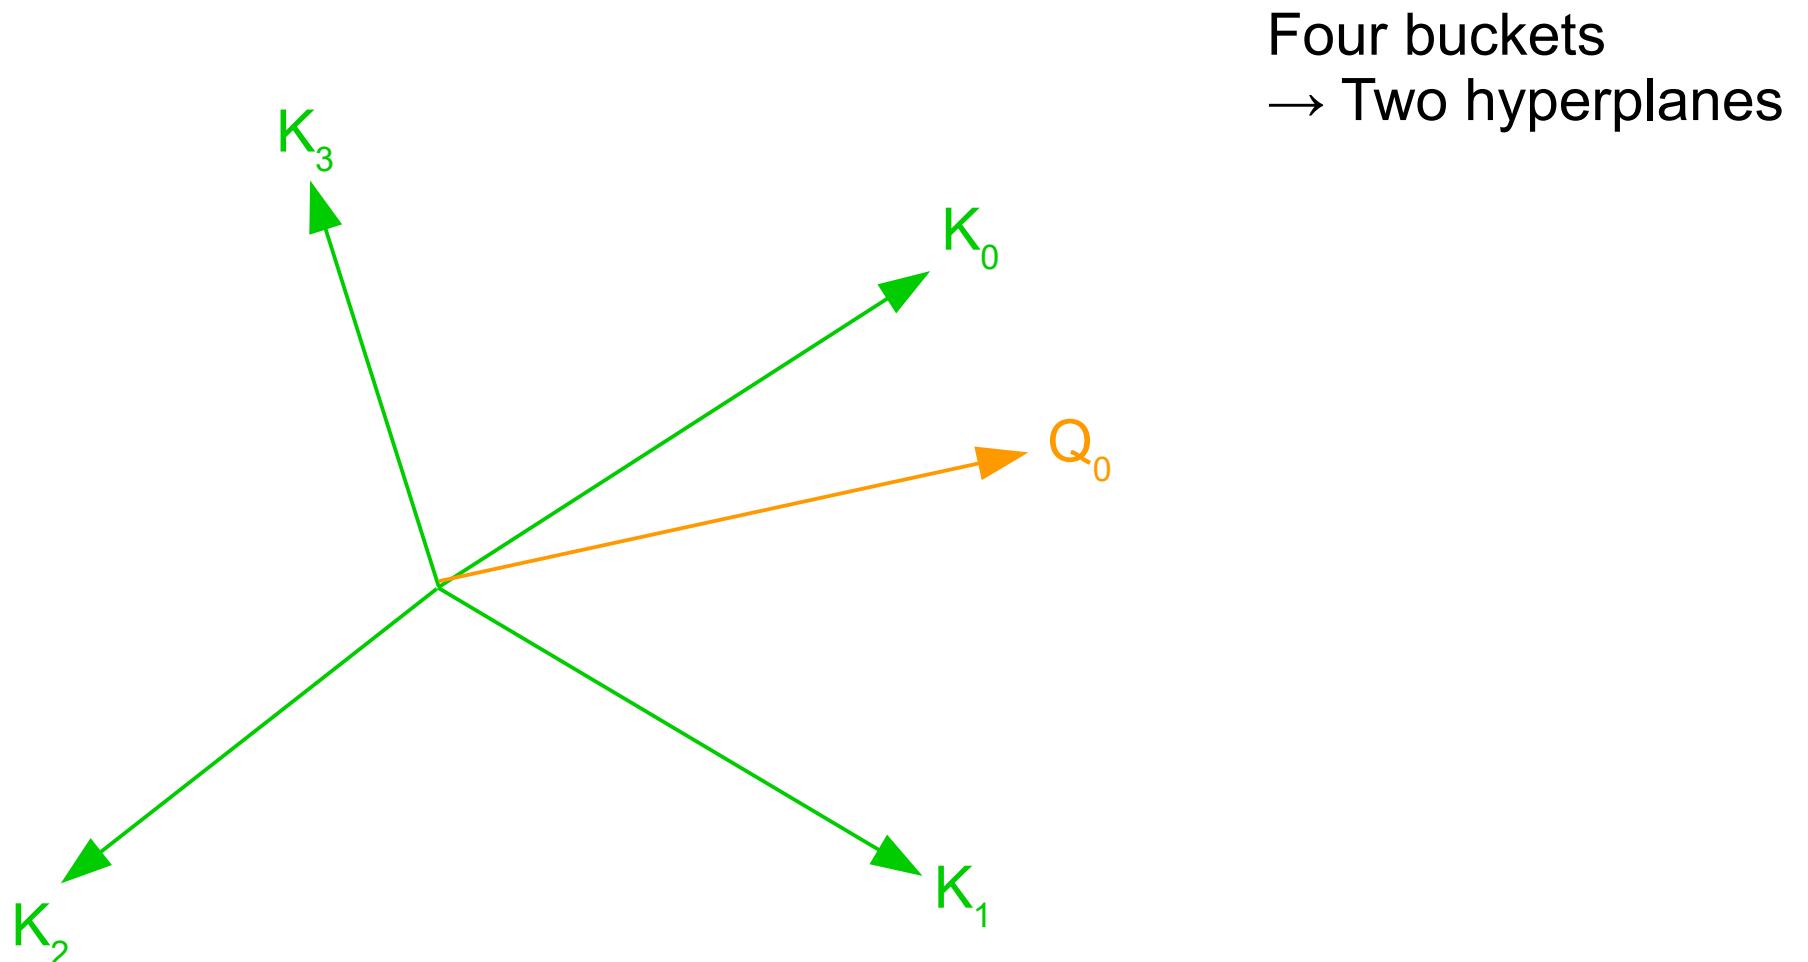
8.3 Reformer: The Efficient Transformer




8.3.1 Random Plane Projections




8.3.1 Random Plane Projections




8.3.1 Random Plane Projections




8.3.1 Random Plane Projections




8.3.1 Random Plane Projections




8.3.1 Random Plane Projections



8.3.1 Random Plane Projections



8.3.1 Random Plane Projections




8.3.2 Reformer: The Efficient Transformer

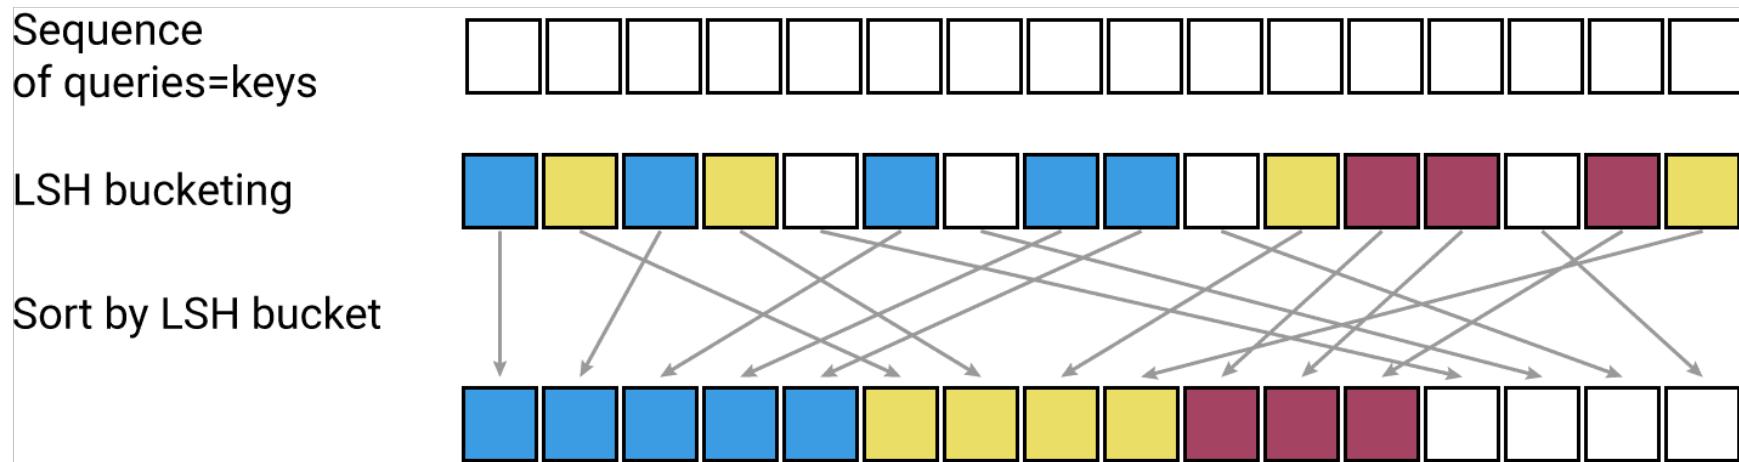
Sequence
of queries=keys




8.3.2 Reformer: The Efficient Transformer



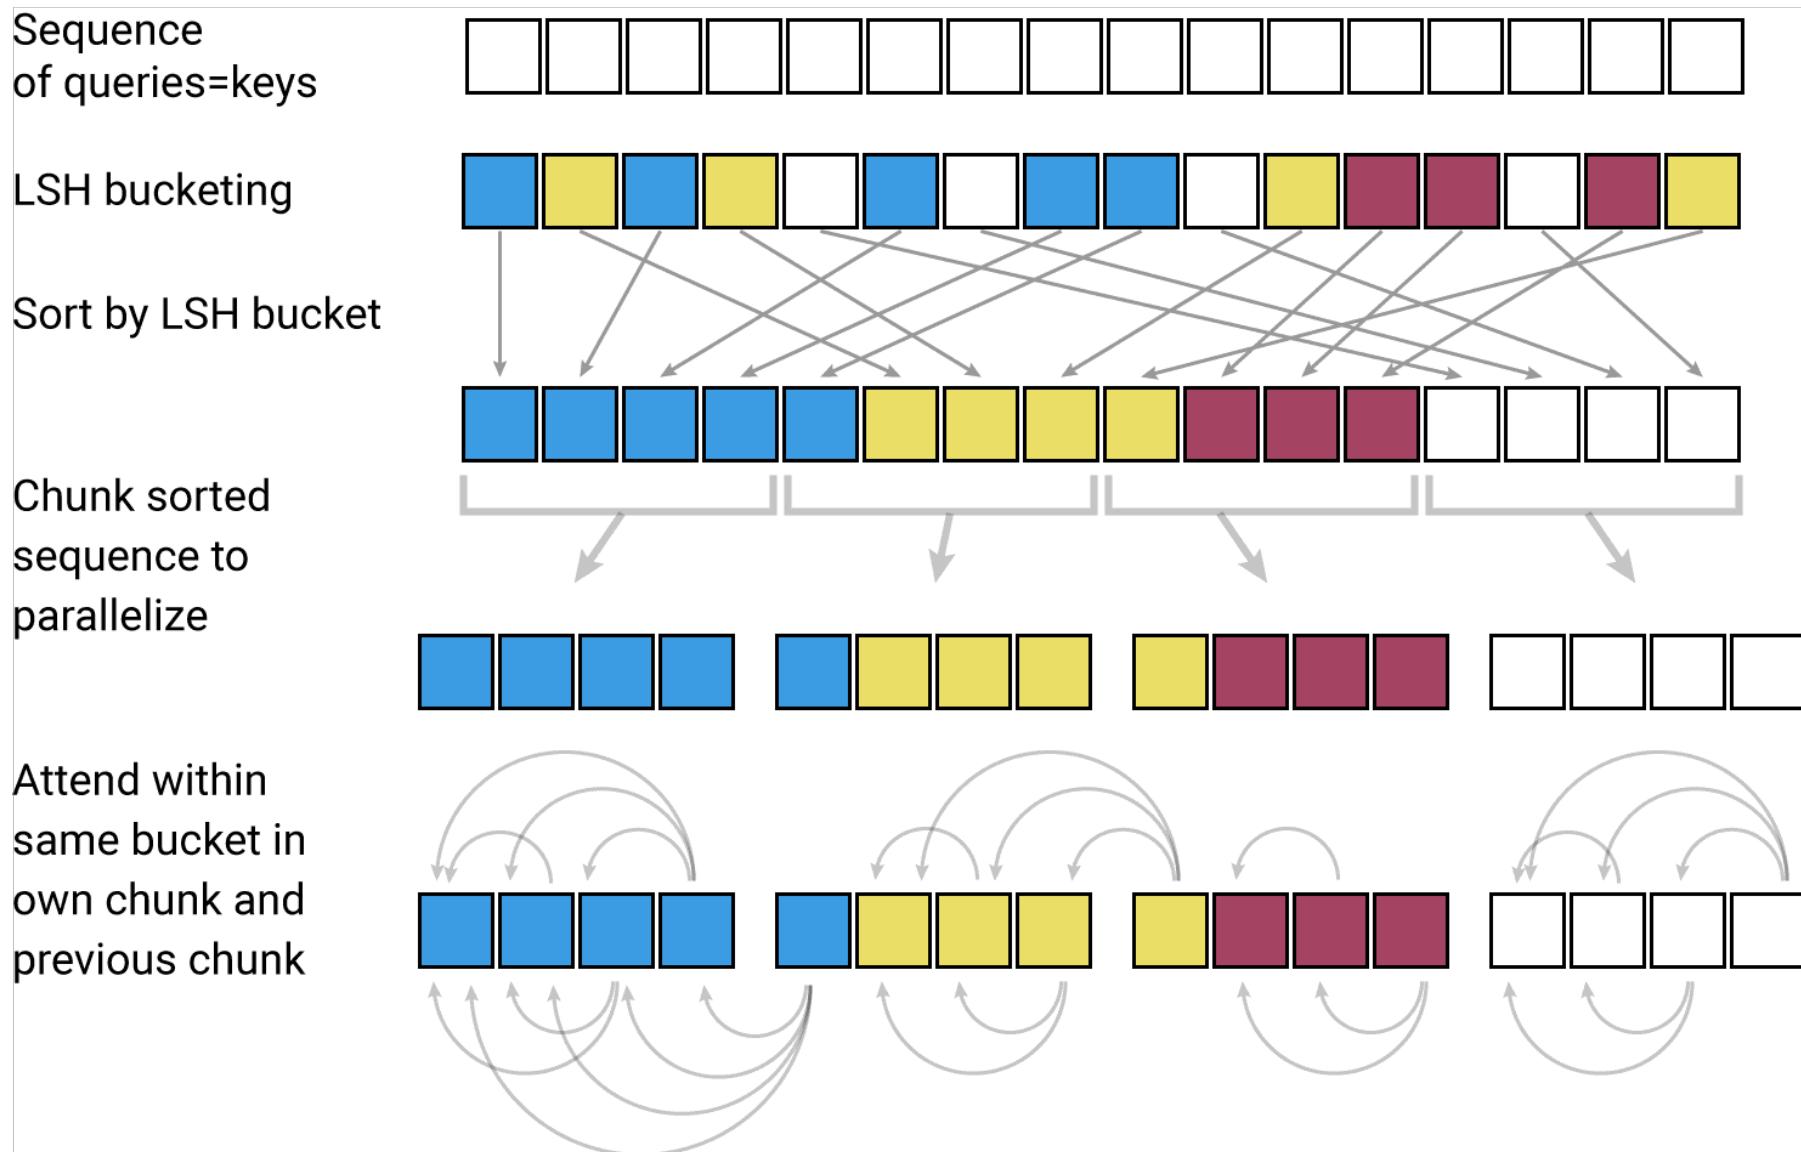
8.3.2 Reformer: The Efficient Transformer




8.3.2 Reformer: The Efficient Transformer




8.3.2 Reformer: The Efficient Transformer





8.3.3 Memory problem – Backpropagation



8.3.3 Memory problem – Backpropagation




8.3.3 Memory problem – Backpropagation




8.3.4 Reversible layer

Residual connection




8.3.4 Reversible layer

Residual connection




Reversible layer
(forward pass)

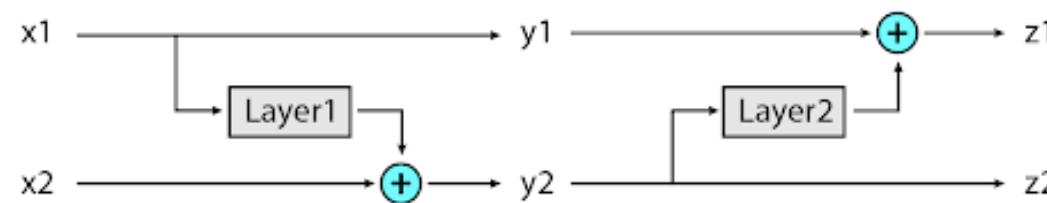


8.3.4 Reversible layer


Residual connection



Reversible layer
(forward pass)




Reversible layer
(backward pass)



9. Conclusions

- Transform data into a sequence
- More general



- Data hungry



- Quadratic scale of self-attention
- Outside of context window
- Performance depends on the situation → RNN better?



Image sources:

https://www.flaticon.com/free-icon/fangs_2068587?term=fangs&page=1&position=3&page=1&position=3&related_id=2068587&origin=search
https://www.flaticon.com/free-icon/file-storage_3616558?term=data&page=1&position=37&page=1&position=37&related_id=3616558&origin=search
https://emojipedia-us.s3.dualstack.us-west-1.amazonaws.com/thumbs/120/whatsapp/326/thinking-face_1f914.png
(call dates: 20.07.22)

AlphaCandy




Figure 35: AlphaCandy UI.

Image sources:
https://www.flaticon.com/free-icon/candy-jar_1075135?related_id=1075135&origin=tag
<https://emojipedia.org/de/toss-face/march-2022/gesicht-mit-umarmenden-h%C3%A4nden/>
(call dates: 20.07.22)

References - Papers

- [1] BAHDANAU, D., CHO, K., AND BENGIO, Y. Neural machine translation by jointly learning to align and translate. In 3rd ICLR, Conference Track Proceedings (2015), Y. Bengio and Y. LeCun, Eds.
- [2] VASWANI, A., SHAZER, N., PARMAR, N., USZKOREIT, J., JONES, L., GOMEZ, A. N., KAISER, L. U., AND POLOSUKHIN, I. Attention is all you need. In Advances in Neural Information Processing Systems (2017), I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, Eds., vol. 30, Curran Associates, Inc.
- [3] DOSOVITSKIY, A., BEYER, L., KOLESNIKOV, A., WEISSENBORN, D., ZHAI, X., UNTERTHINER, T., DEHGHANI, M., MINDERER, M., HEIGOLD, G., GELLY, S., USZKOREIT, J., AND HOULSBY, N. An image is worth 16x16 words: Transformers for image recognition at scale. In 9th ICLR (2021), OpenReview.net.
- [4] CARION, N., MASSA, F., SYNNAEVE, G., USUNIER, N., KIRILLOV, A., AND ZAGORUYKO, S. End-to-end object detection with transformers. In Computer Vision - ECCV - 16th European Conference, Proceedings, Part I (2020), A. Vedaldi, H. Bischof, T. Brox, and J. Frahm, Eds., vol. 12346 of Lecture Notes in Computer Science, Springer, pp. 213–229.
- [5] JIANG, Y., CHANG, S., AND WANG, Z. Transgan: Two pure transformers can make one strong gan, and that can scale up. In Advances in Neural Information Processing Systems 34: NeurIPS 2021 (2021), M. Ranzato, A. Beygelzimer, Y. N. Dauphin, P. Liang, and J. W. Vaughan, Eds., pp. 14745–14758.
- [6] UPADHYAY, U., SHAH, N., RAVIKANTI, S., AND MEDHE, M. Transformer based reinforcement learning for games. vol. Abs/1912.03918.

References - Papers

- [7] CHEN, L., LU, K., RAJESWARAN, A., LEE, K., GROVER, A., LASKIN, M., ABBEEL, P., SRINIVAS, A., AND MORDATCH, I. Decision transformer: Reinforcement learning via sequence modeling. In Advances in Neural Information Processing Systems 34: NeurIPS 2020 (2021), M. Ranzato, A. Beygelzimer, Y. N. Dauphin, P. Liang, and J. W. Vaughan, Eds., pp. 15084–15097.
- [8] DEVLIN, J., CHANG, M., LEE, K., AND TOUTANOVA, K. BERT: pre-training of deep bidirectional transformers for language understanding. In Proceedings of the 2019 Conference of the NAACL-HLT, Volume 1 (Long and Short Papers) (2019), J. Burstein, C. Doran, and T. Solorio, Eds., Association for Computational Linguistics, pp. 4171–4186.
- [9] BROWN, T. B., MANN, B., RYDER, N., SUBBIAH, M., KAPLAN, J., DHARIWAL, P., NEELAKANTAN, A., SHYAM, P., SASTRY, G., ASKELL, A., AGARWAL, S., HERBERT-VOSS, A., KRUEGER, G., HENIGHAN, T., CHILD, R., RAMESH, A., ZIEGLER, D. M., WU, J., WINTER, C., HESSE, C., CHEN, M., SIGLER, E., LITWIN, M., GRAY, S., CHESS, B., CLARK, J., BERNER, C., MCCANDLISH, S., RADFORD, A., SUTSKEVER, I., AND AMODEI, D. Language models are few-shot learners. In Advances in Neural Information Processing Systems 33: NeurIPS 2020 (2020), H. Larochelle, M. Ranzato, R. Hadsell, M. Balcan, and H. Lin, Eds.